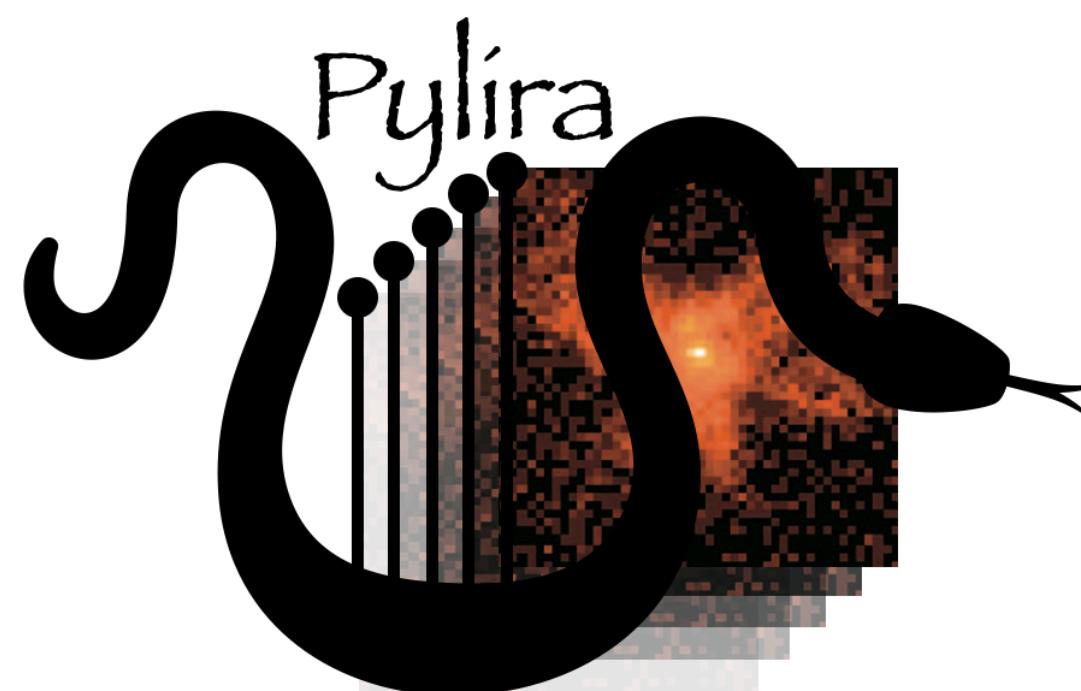


# Pylira: deconvolution of images in the presence of Poisson noise

July 13th, 2022

Scipy 2022, Austin TX



**Axel Donath**

for the CHASC Astro-Statistics Collaboration

# About me



- Post-Doc @ Center for Astrophysics | Harvard Smithsonian
- Gamma-ray & X-ray astronomer, open source scientific software developer, self-taught “statistician”
- Member of the **CHASC Astro-Statistics Collaboration**
  - <https://hea-www.harvard.edu/astrostat/>
  - <https://github.com/astrostat/>
- Also one of the **lead developers of Gammapy** (<https://gammapy.org>) a Python package for analysis of astronomical gamma-ray data and “Science Tool” library for the Cherenkov Telescope Array (<https://www.cta-observatory.org>)
- My GitHub profile: <https://github.com/adonath>

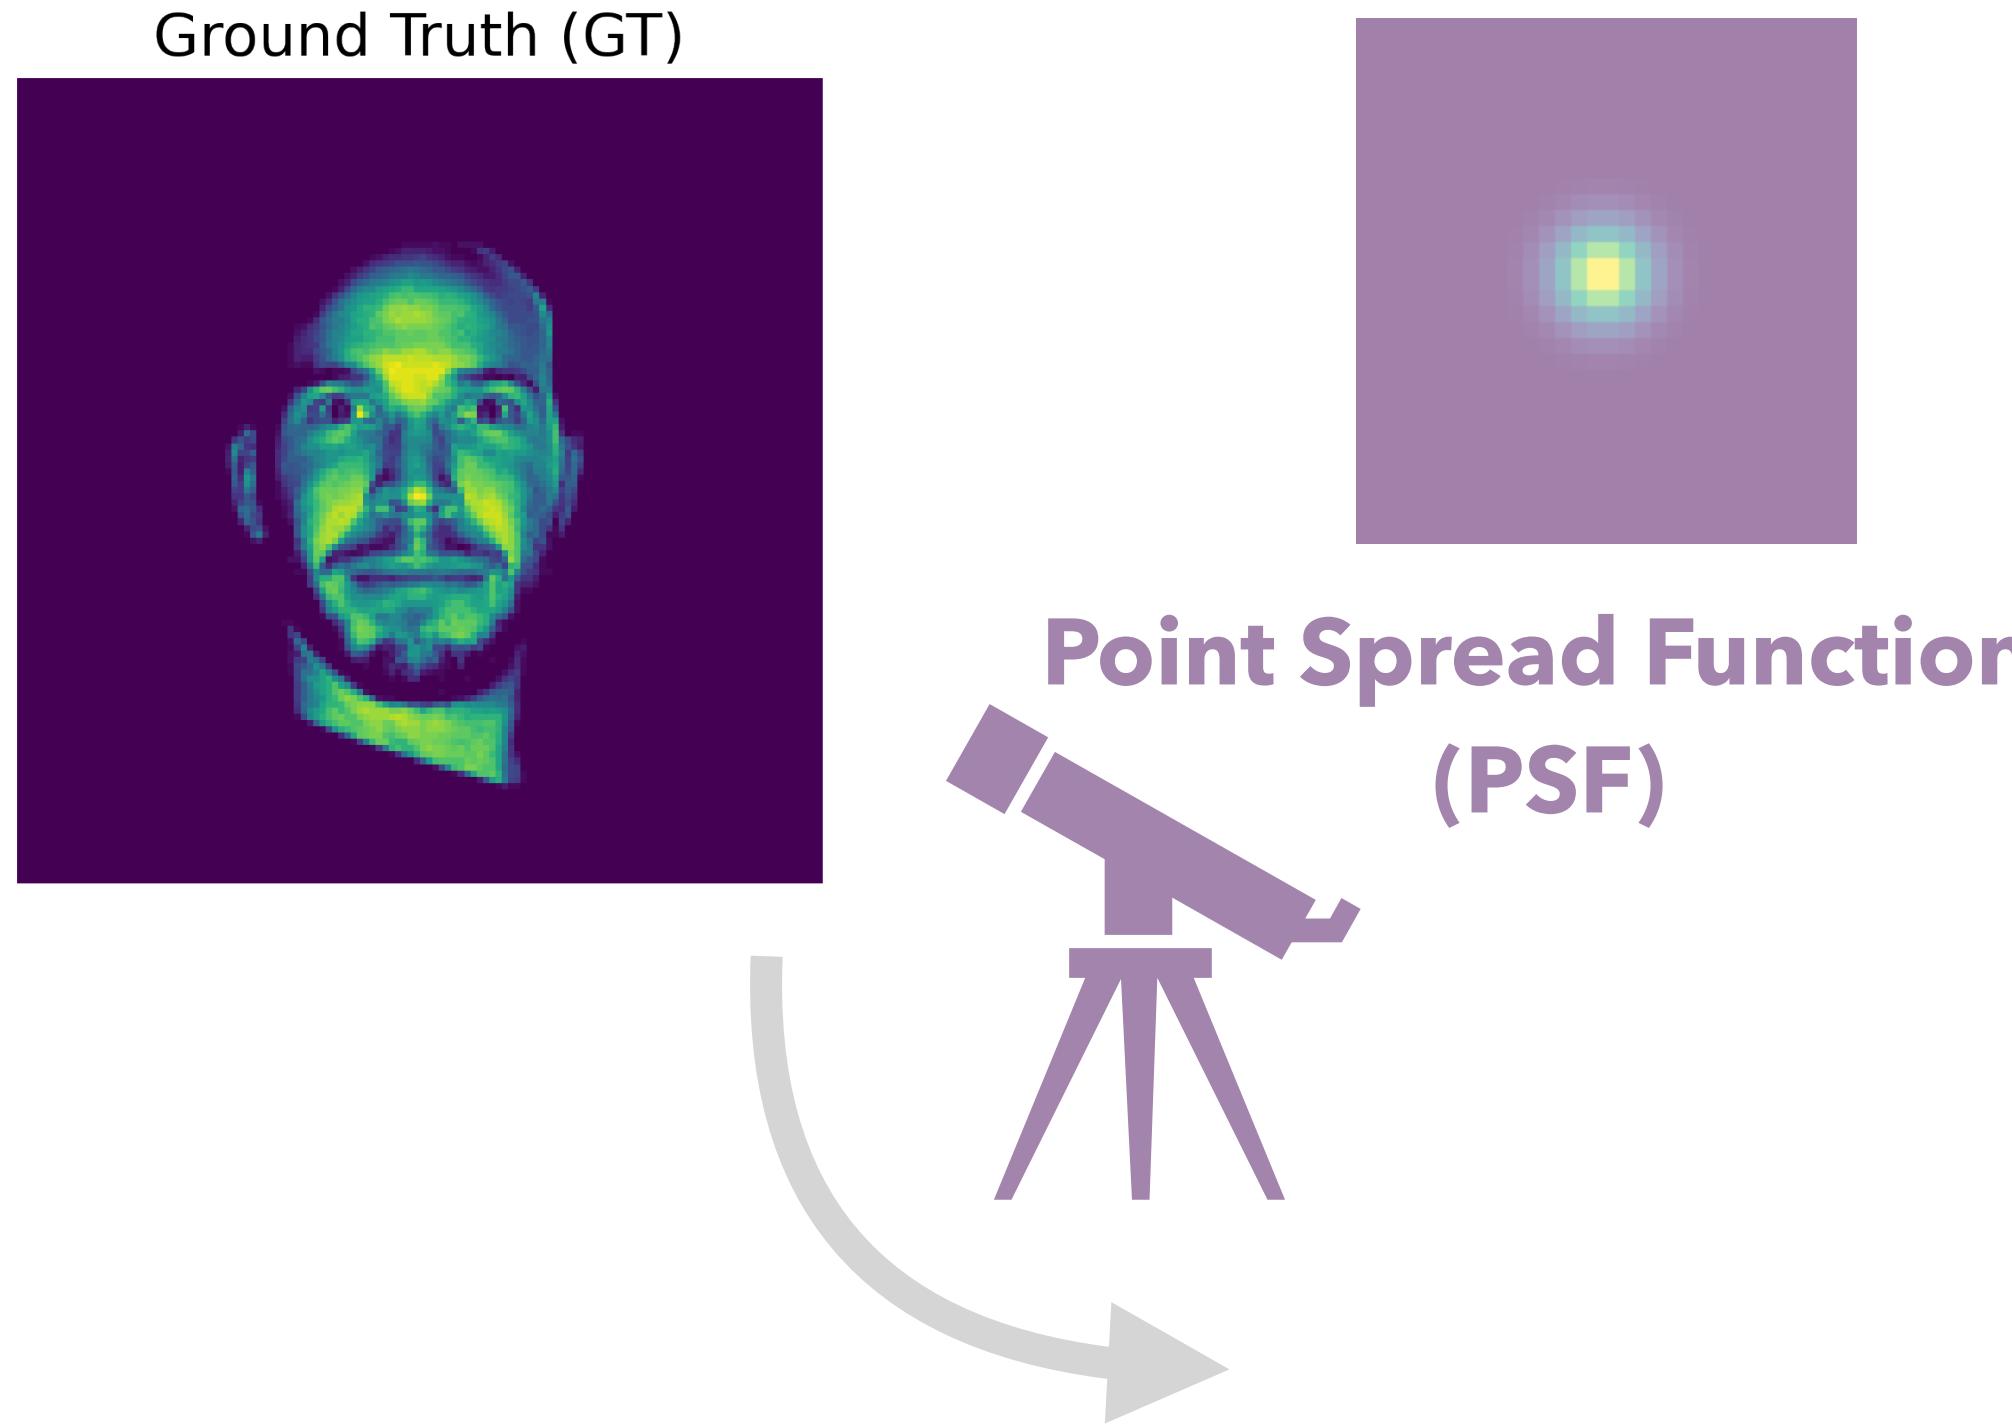
# Introduction

# The “low counts” imaging process

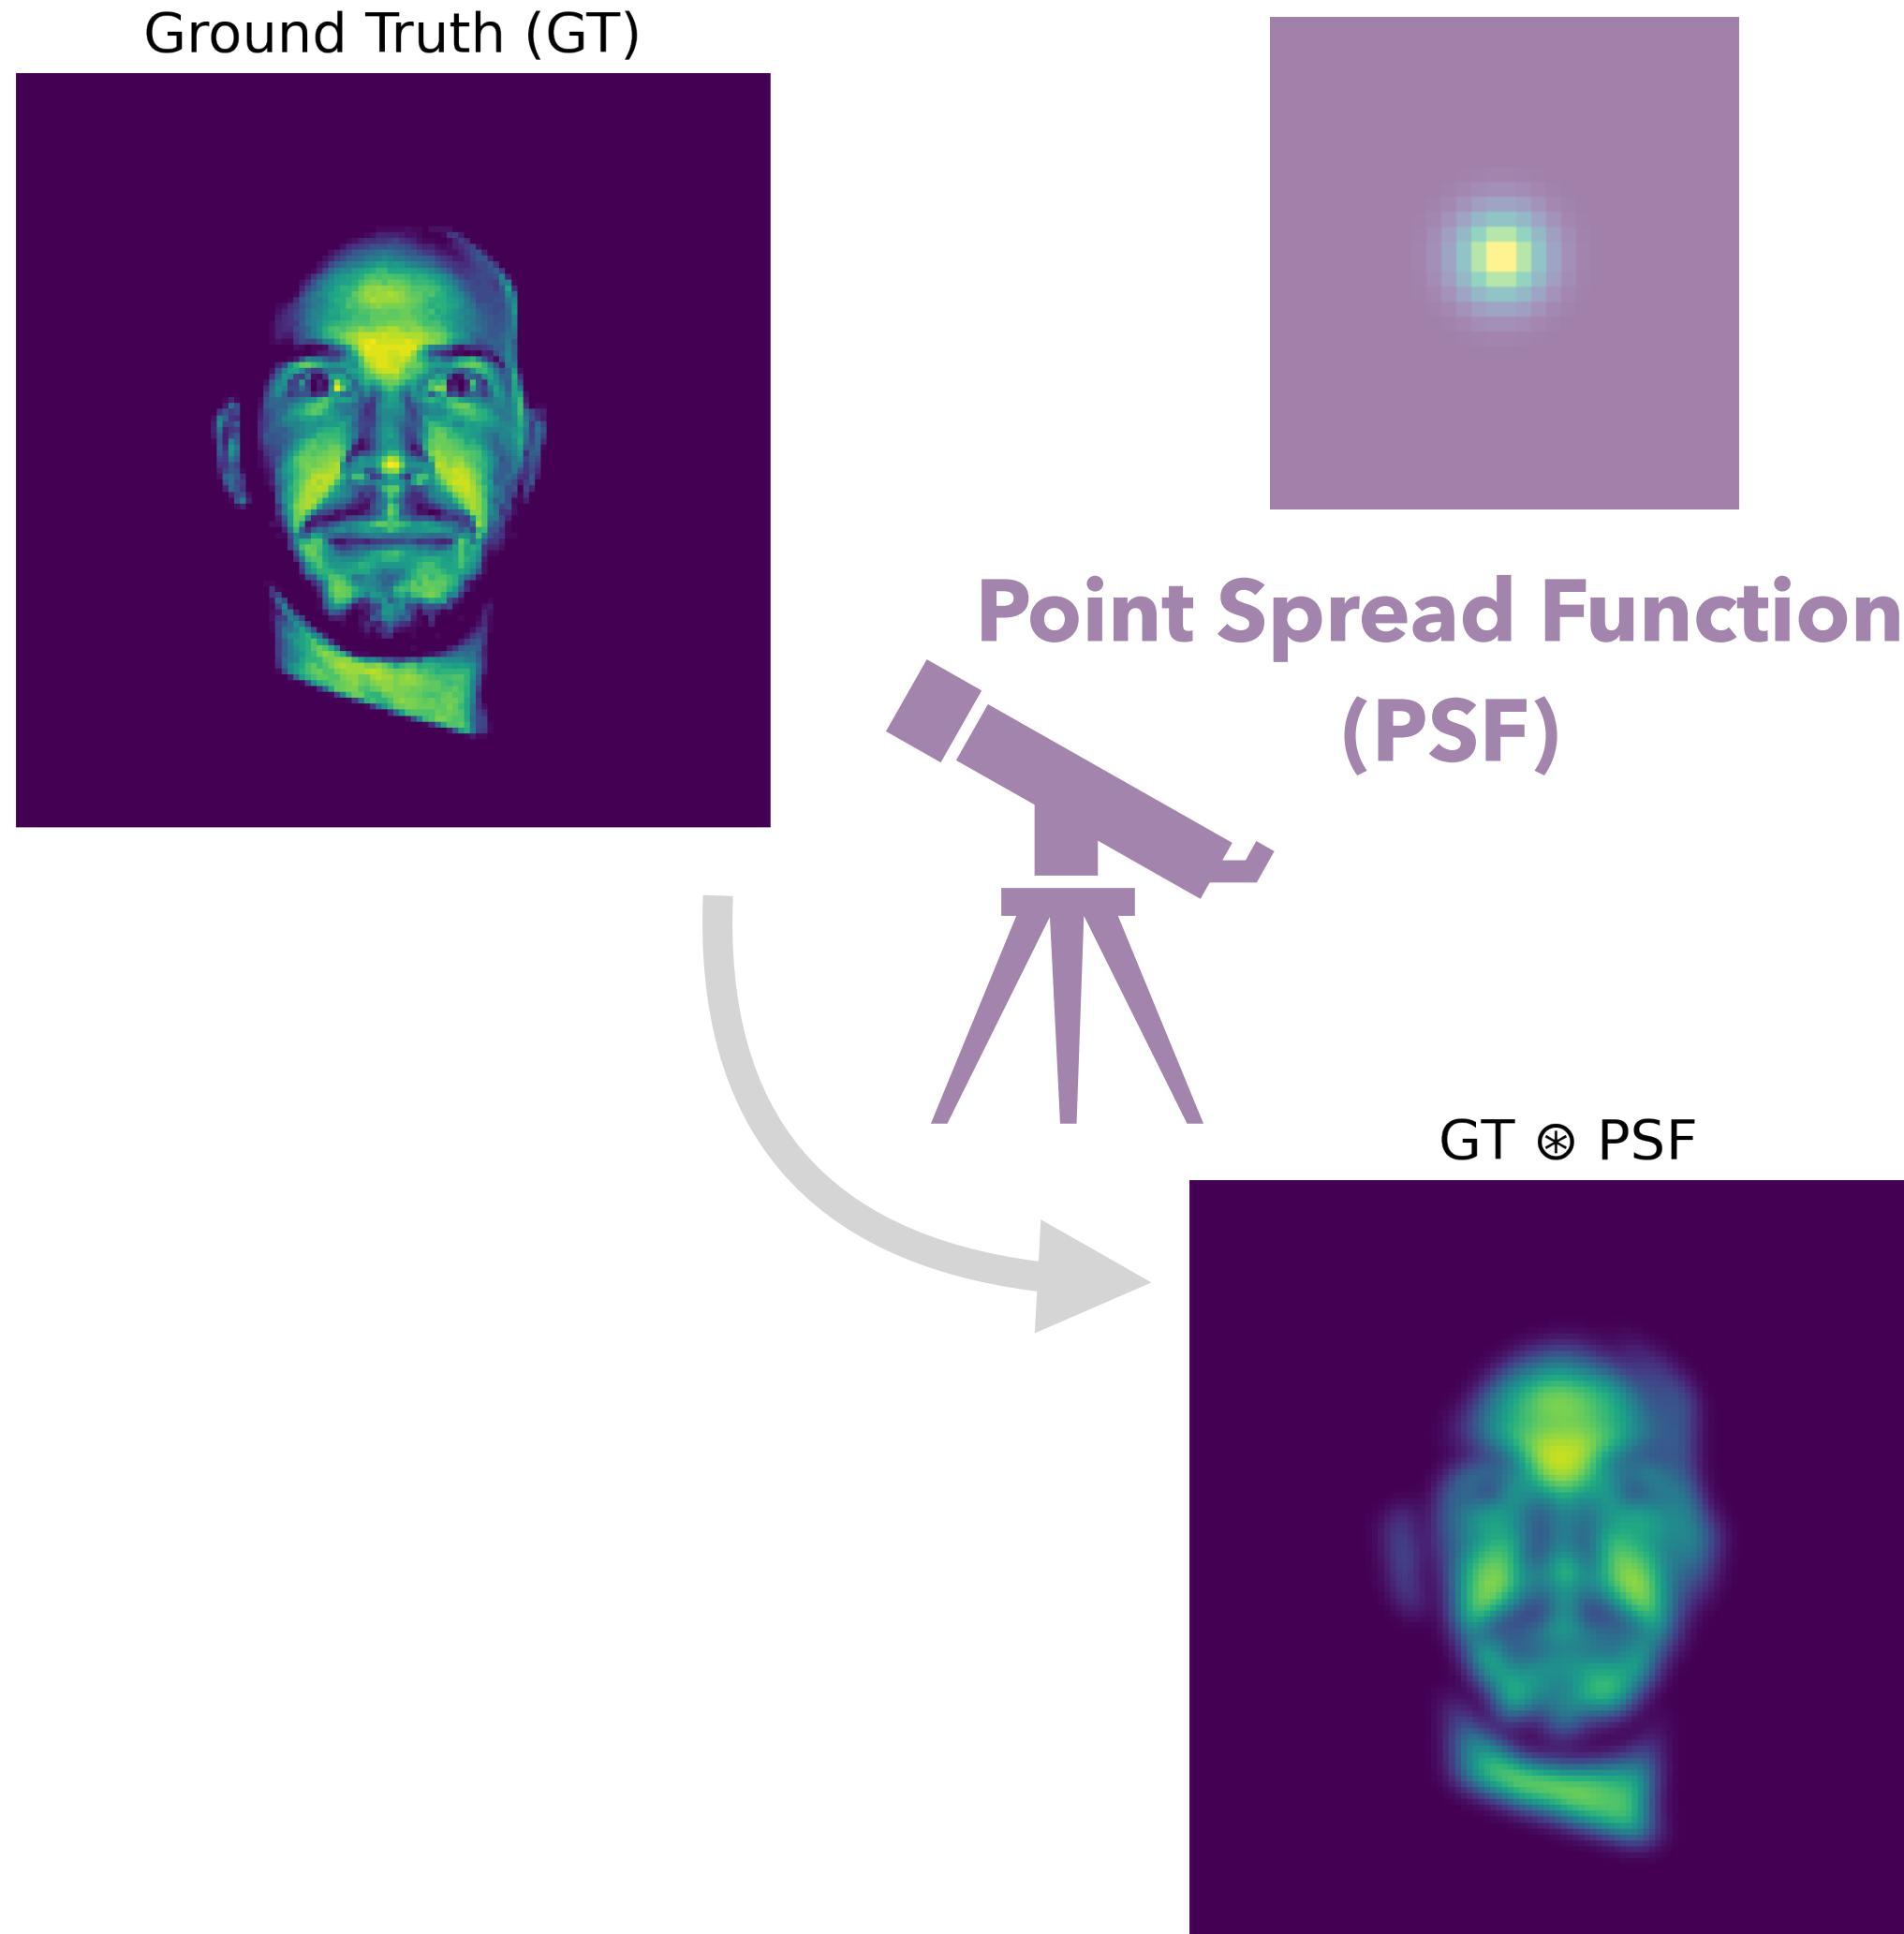
Ground Truth (GT)



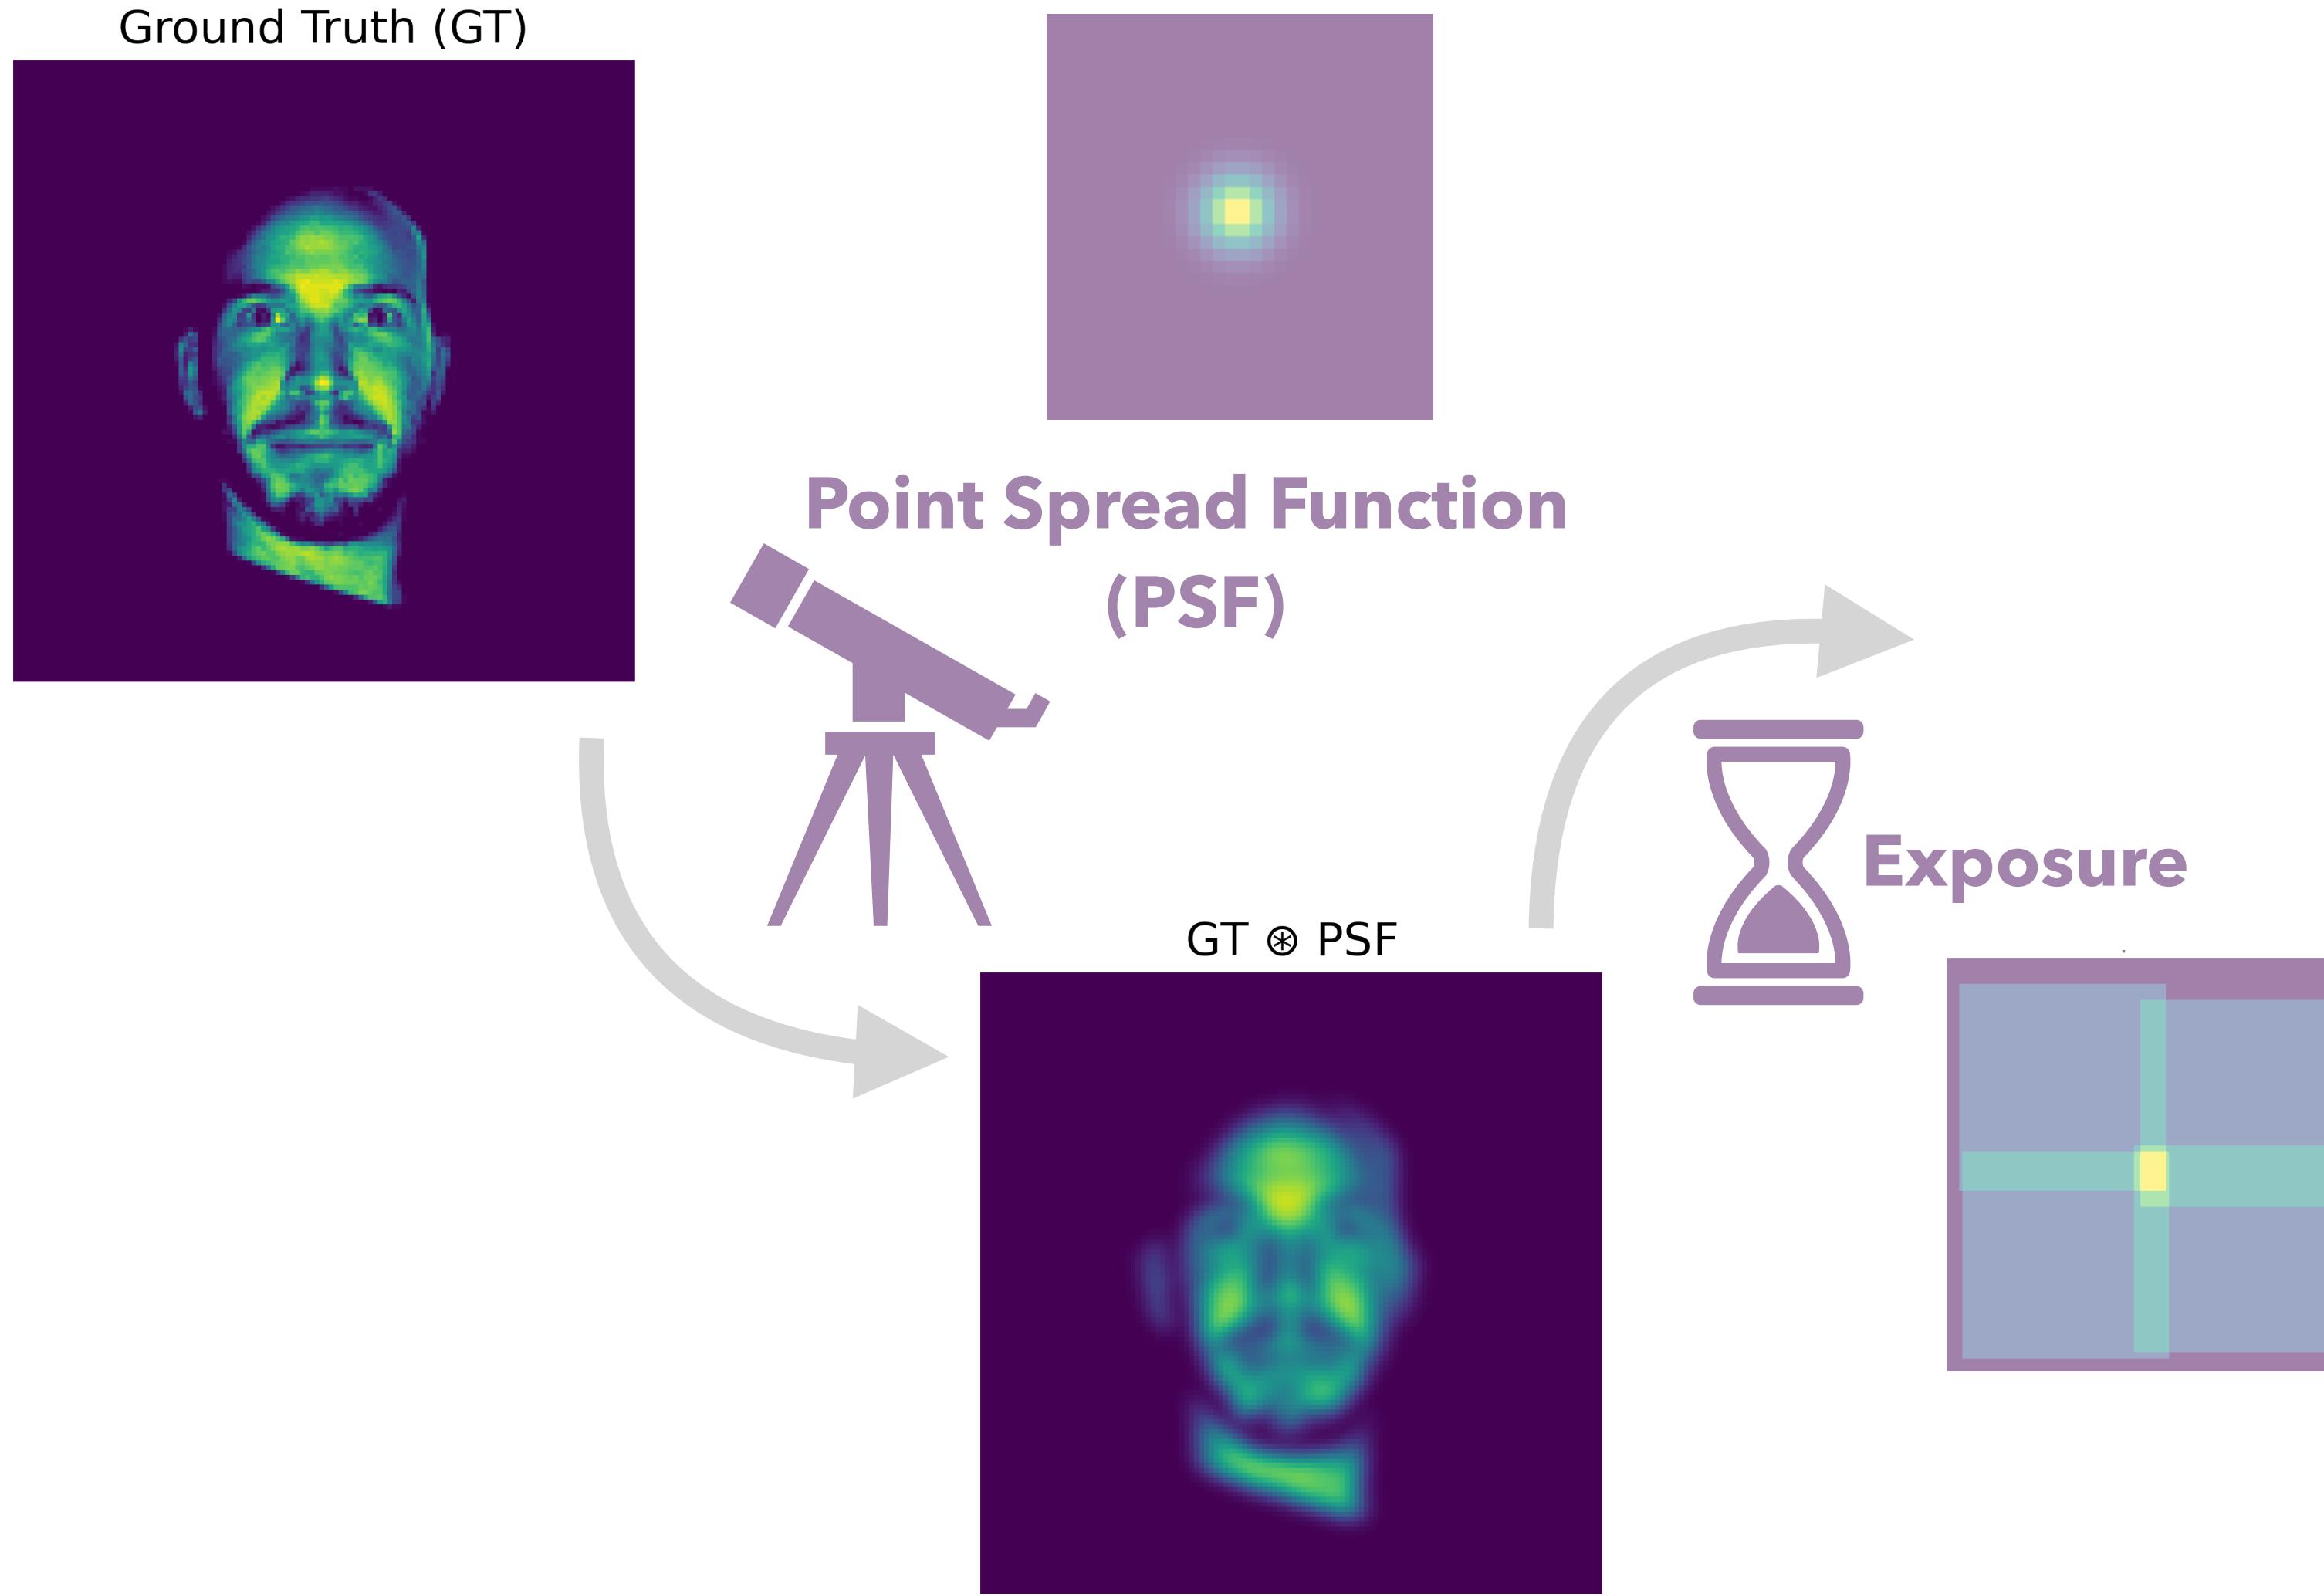
# The “low counts” imaging process



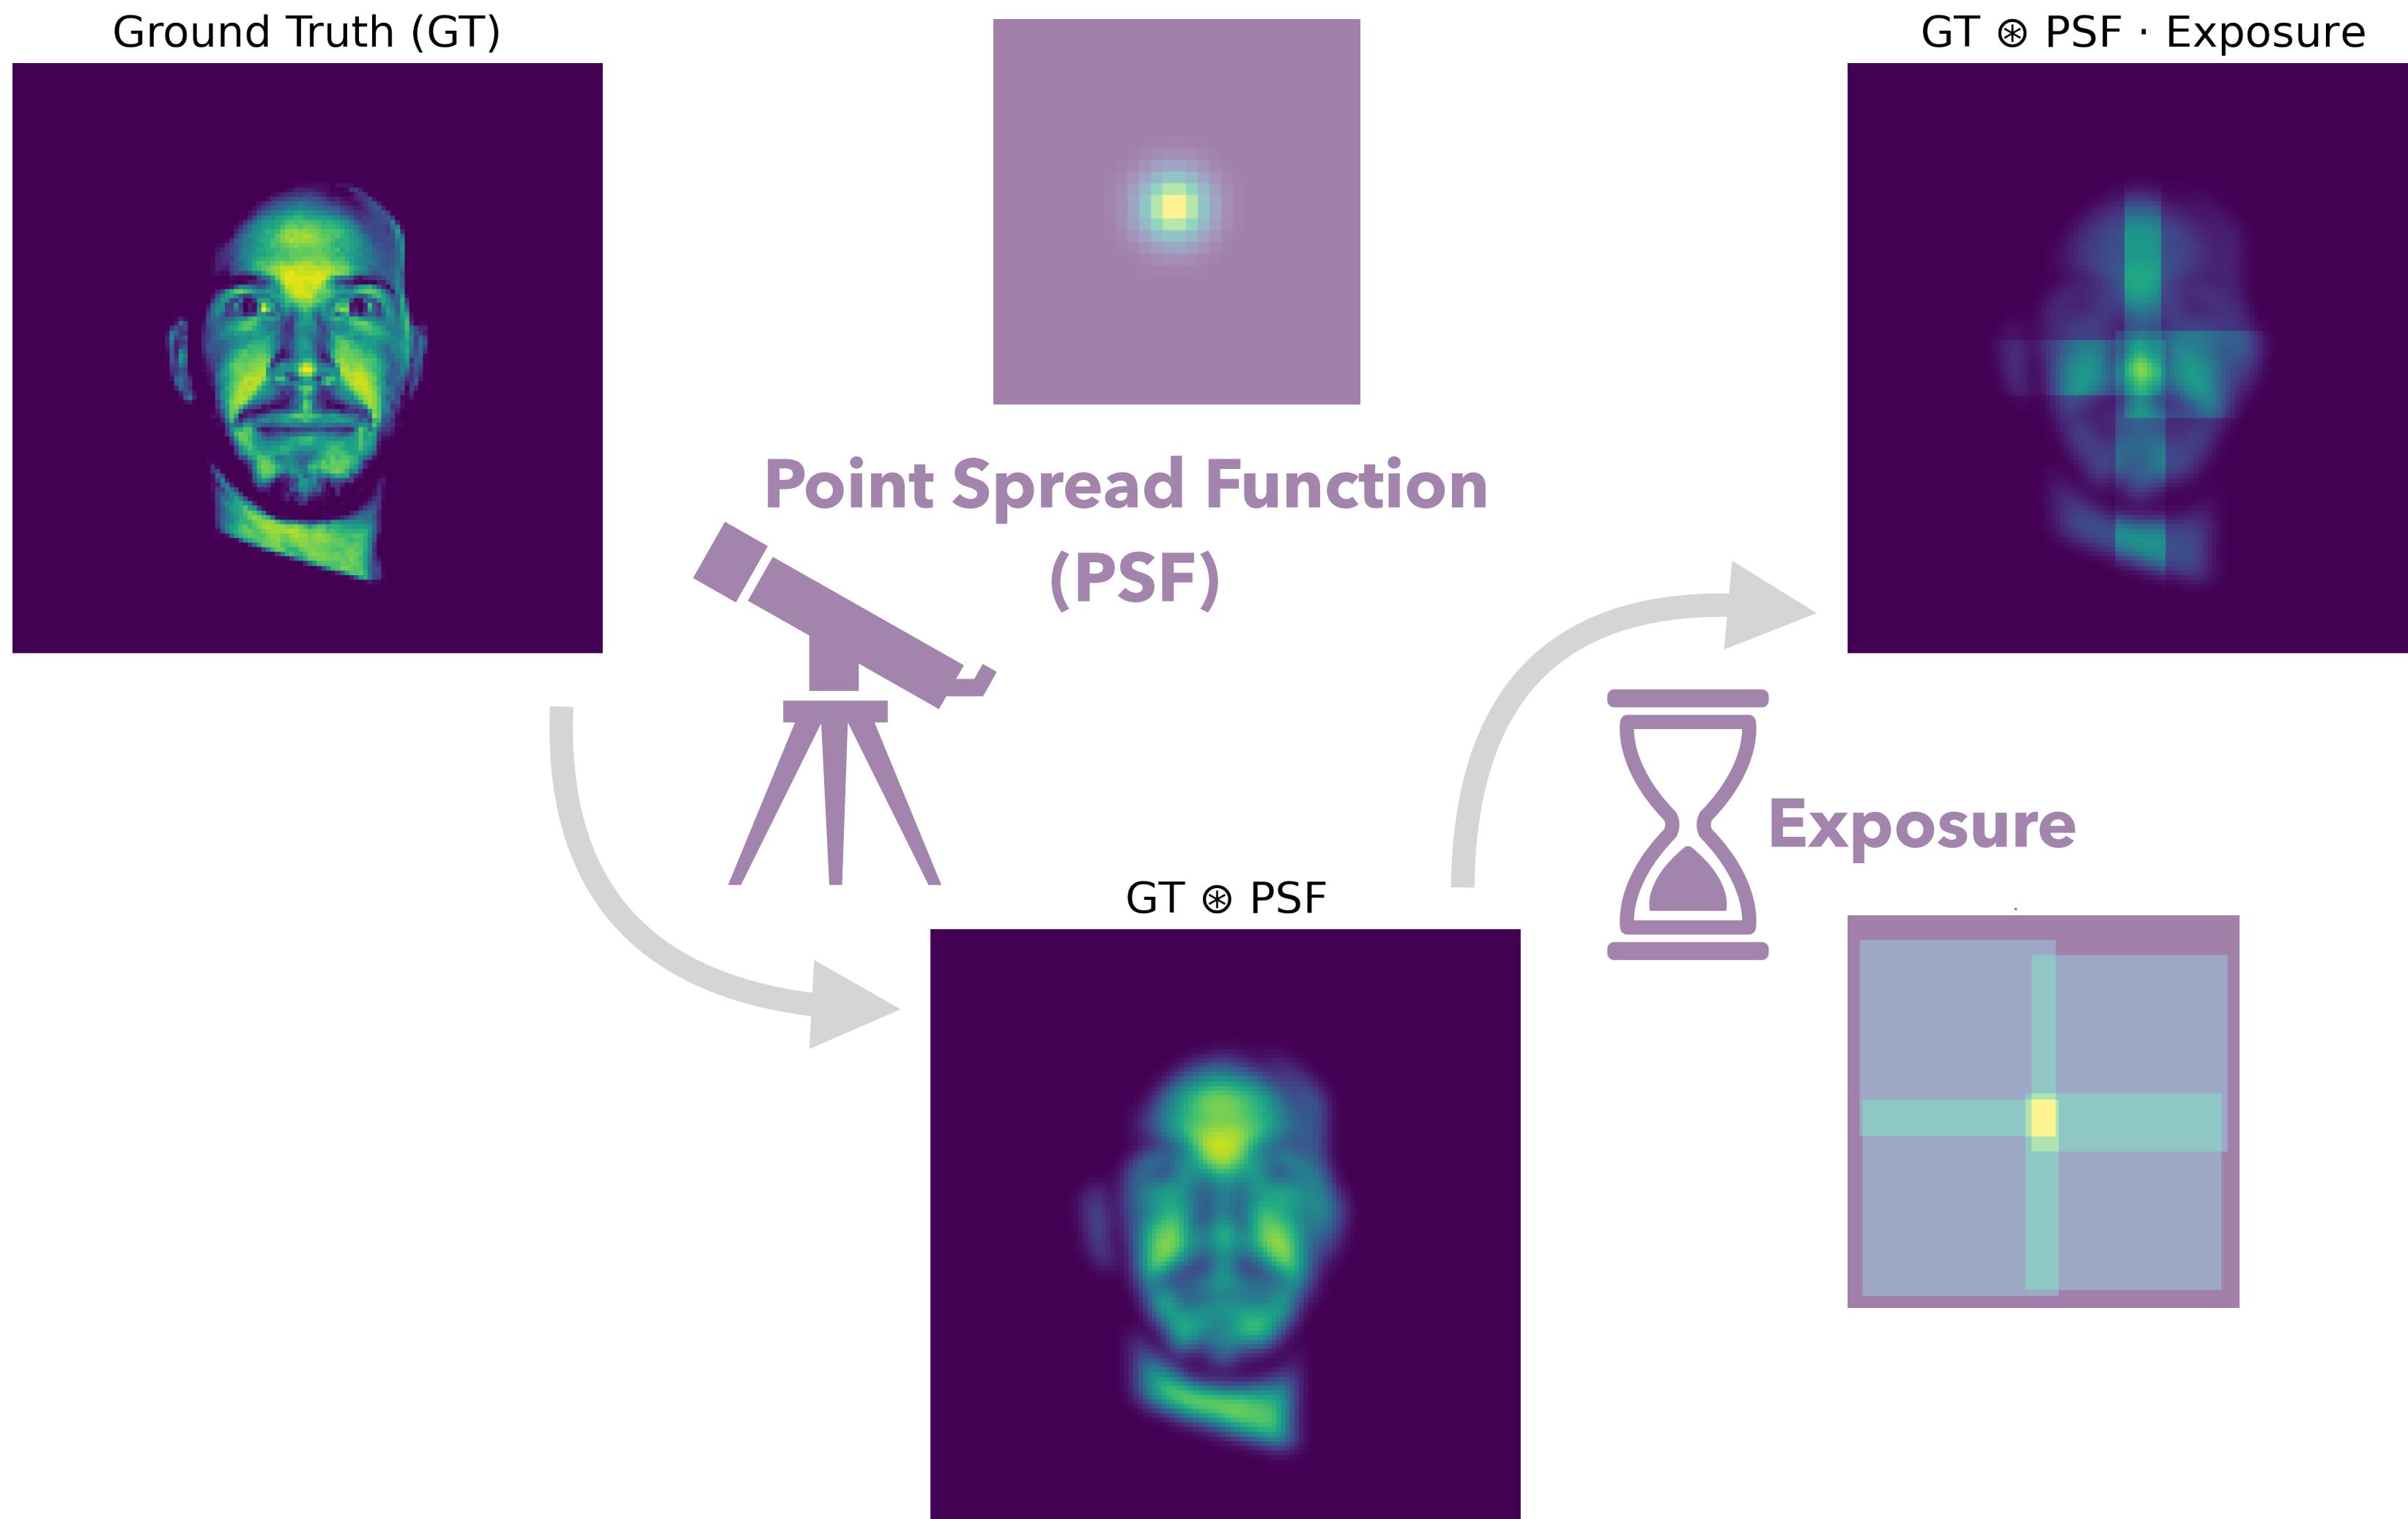
# The “low counts” imaging process



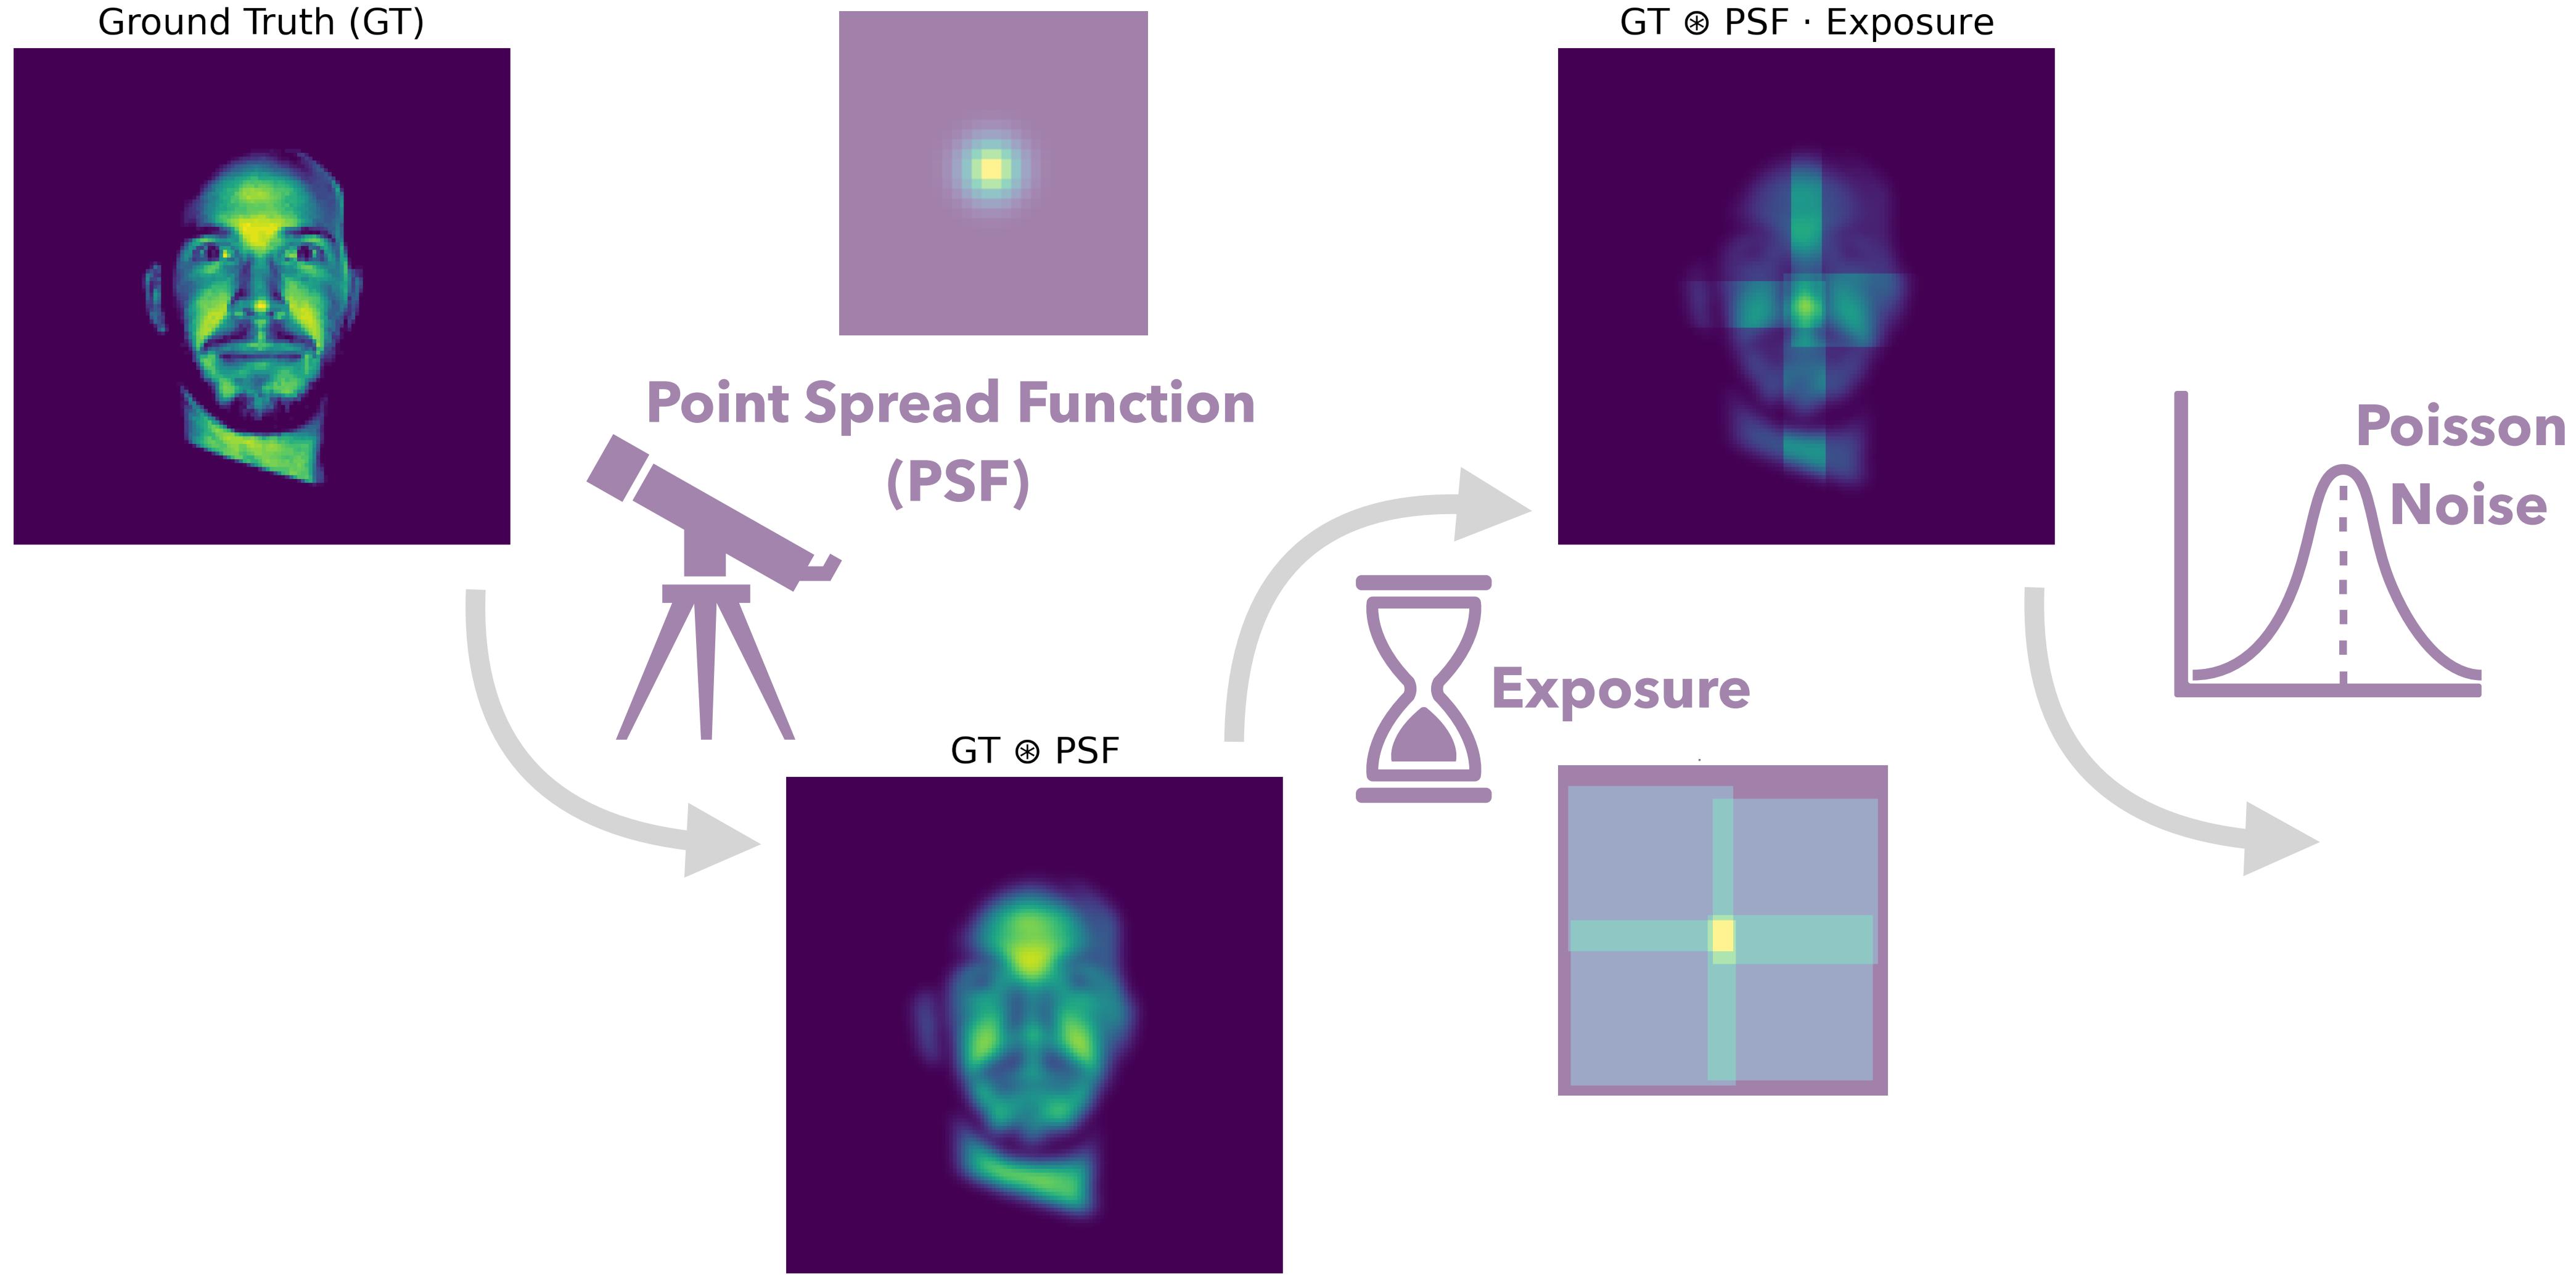
# The “low counts” imaging process



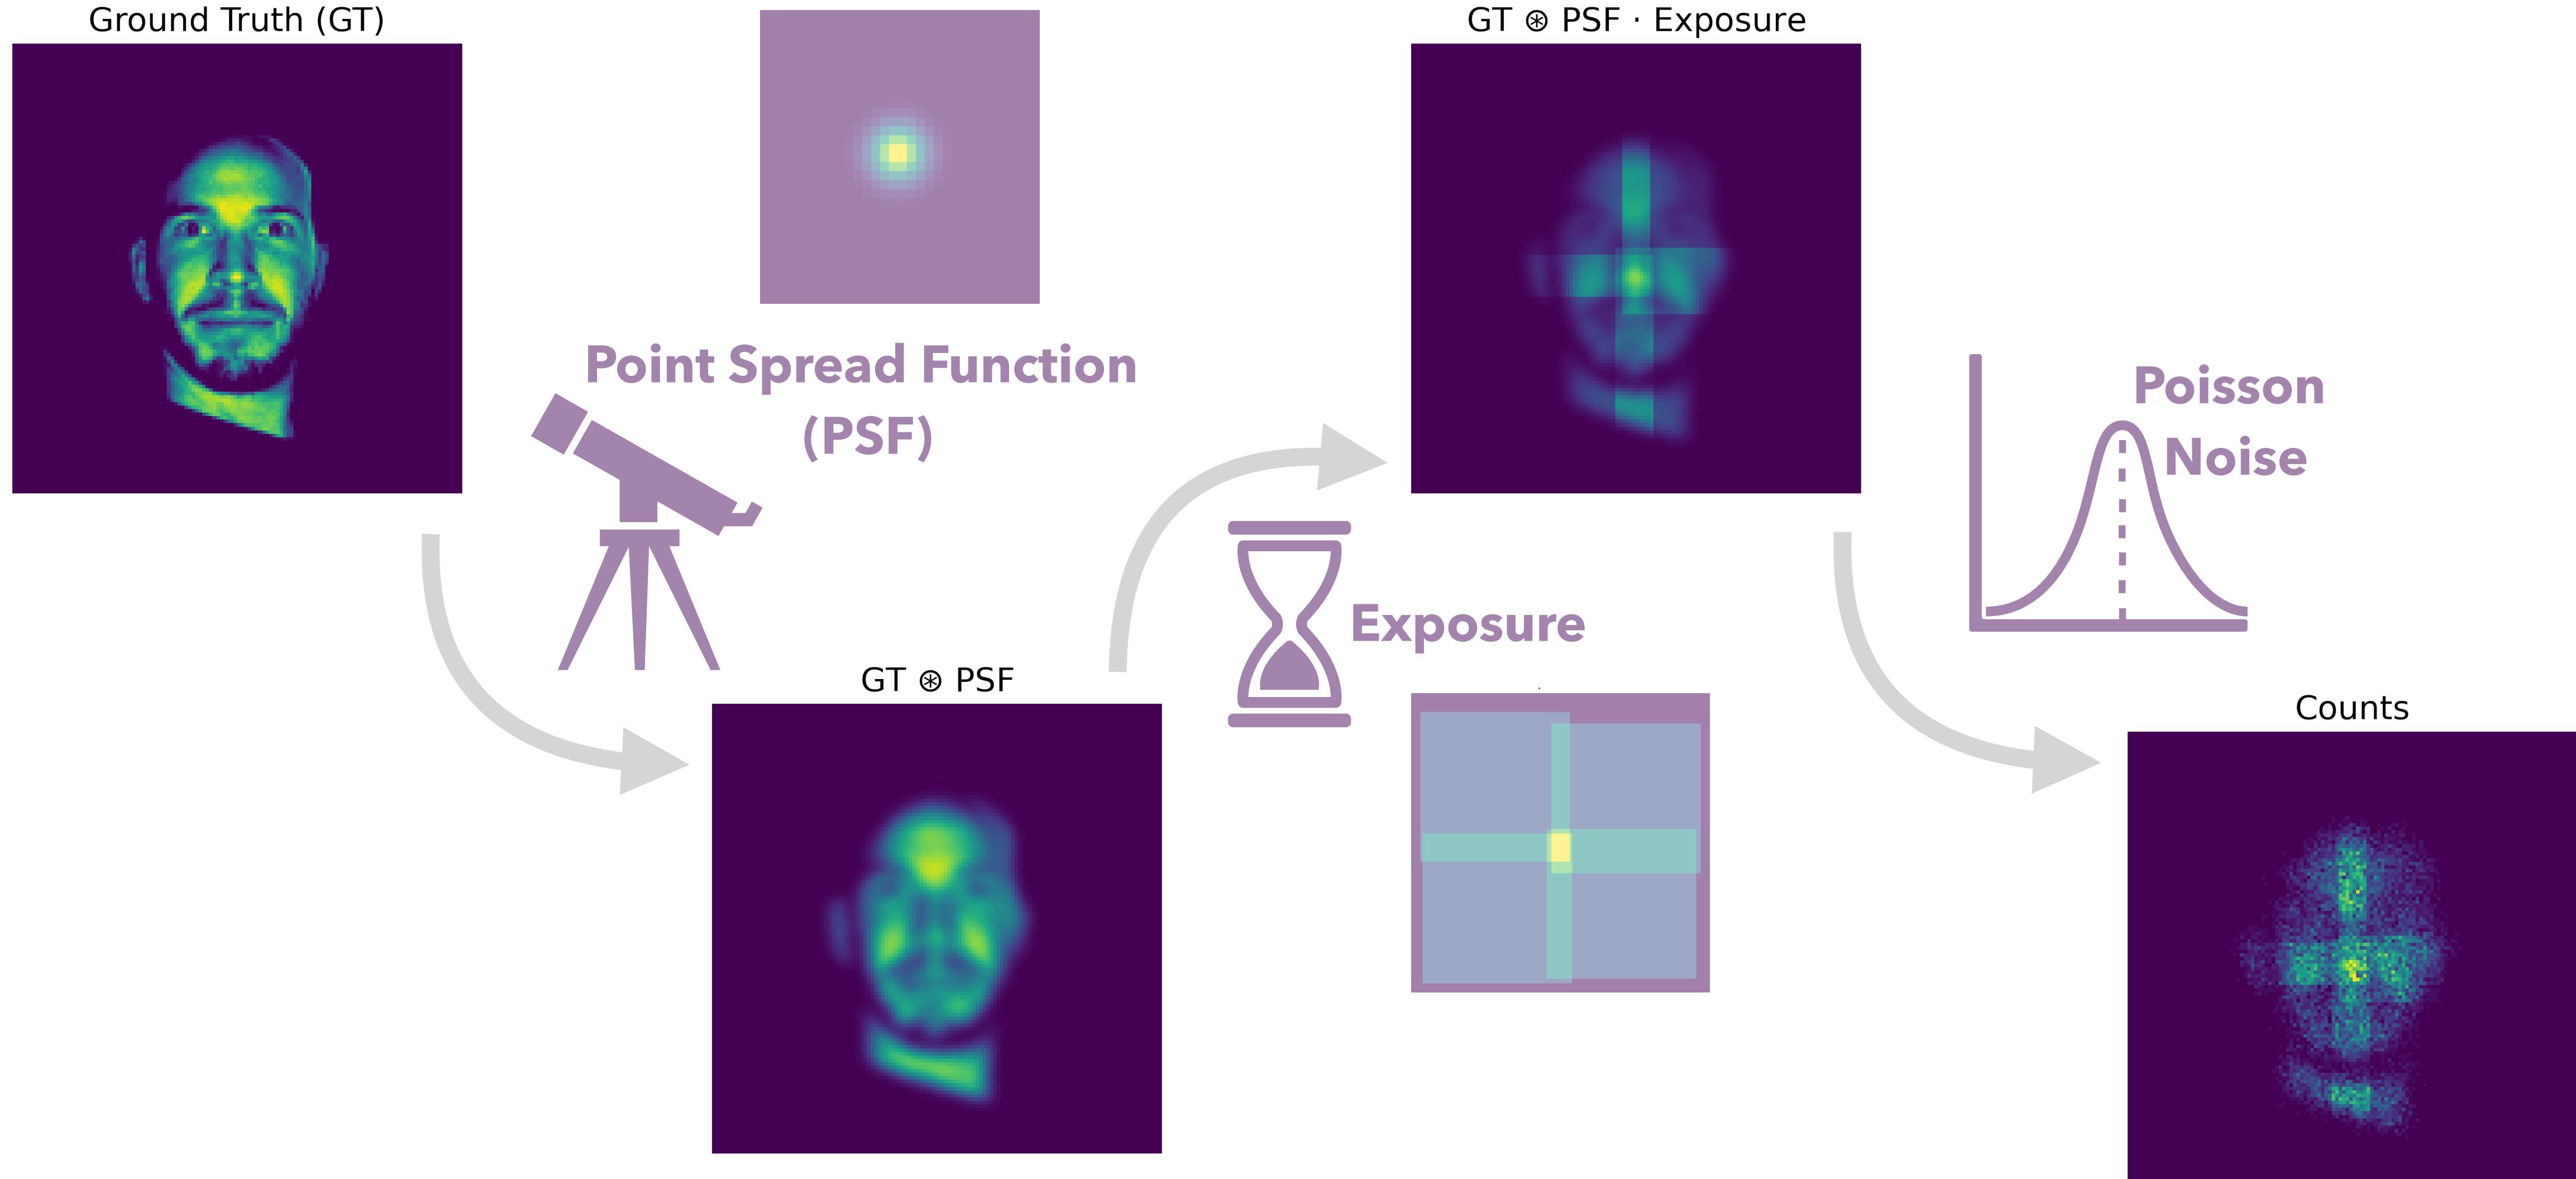
# The “low counts” imaging process



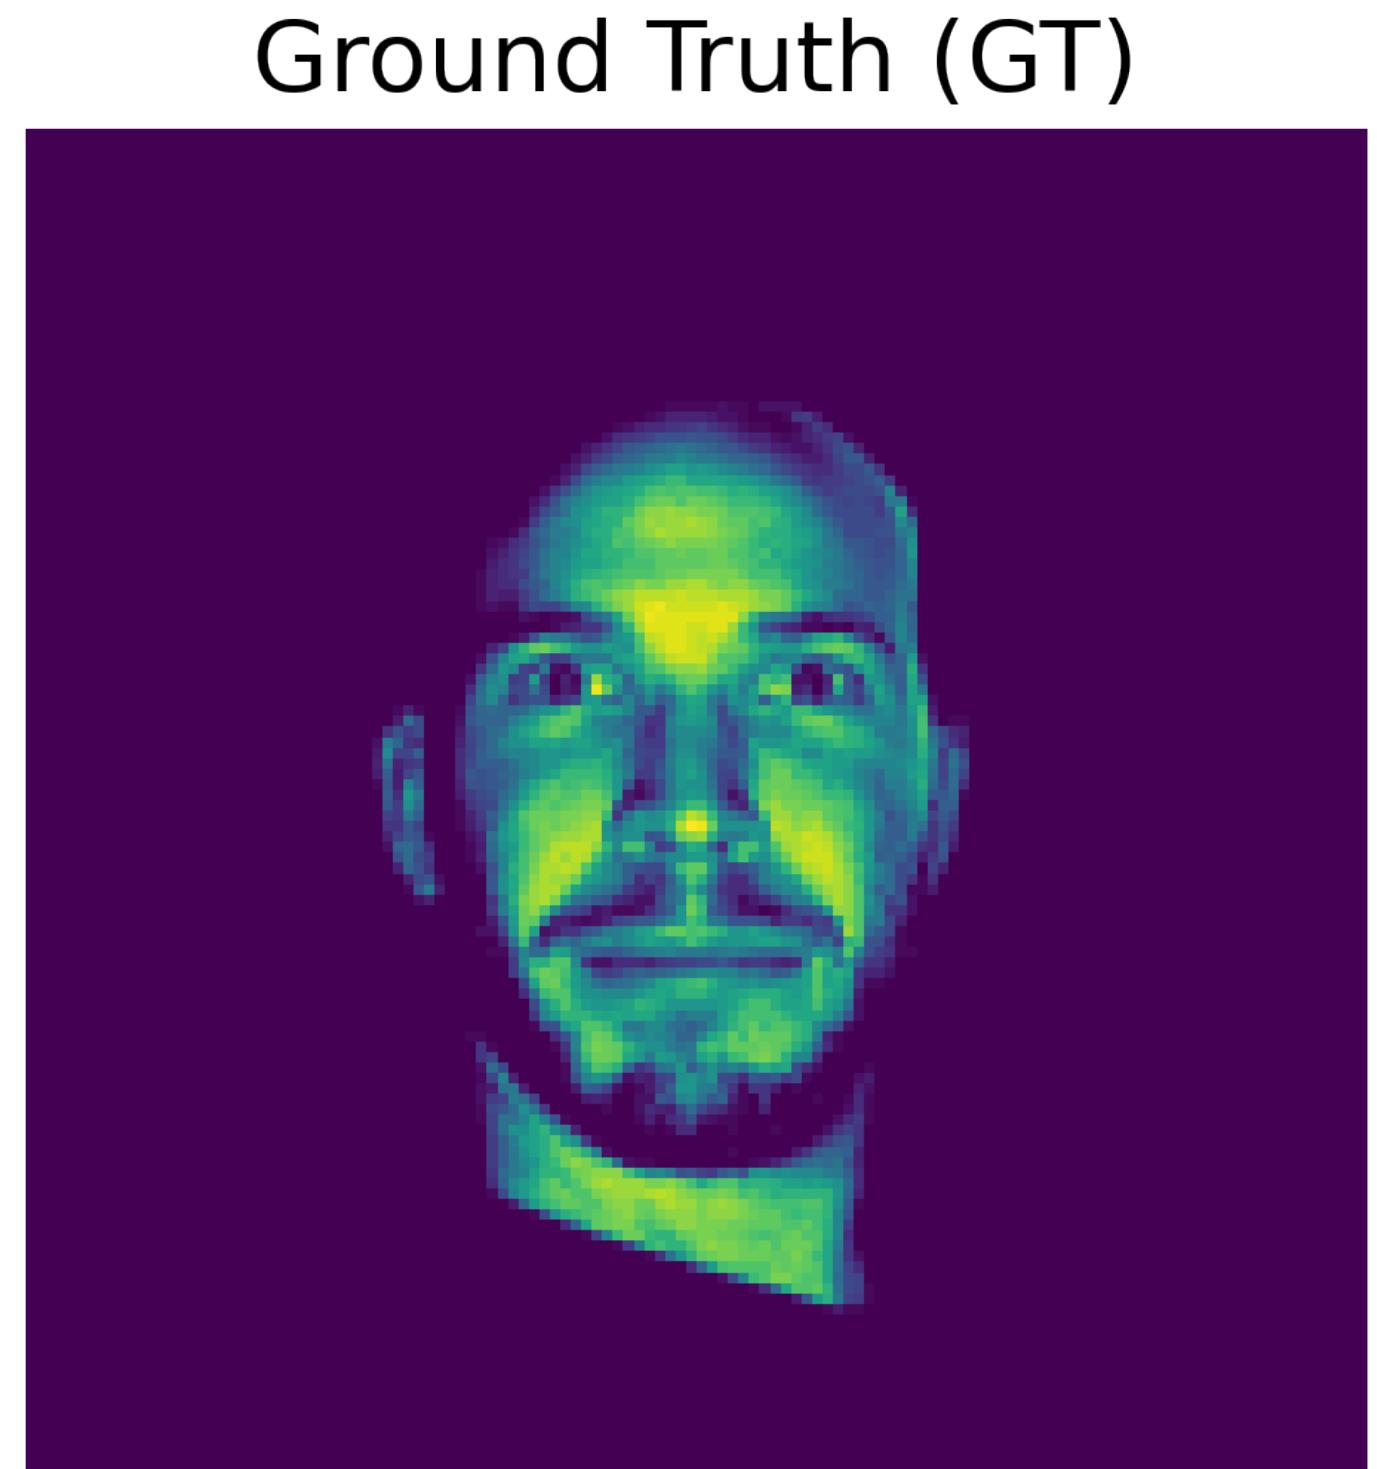
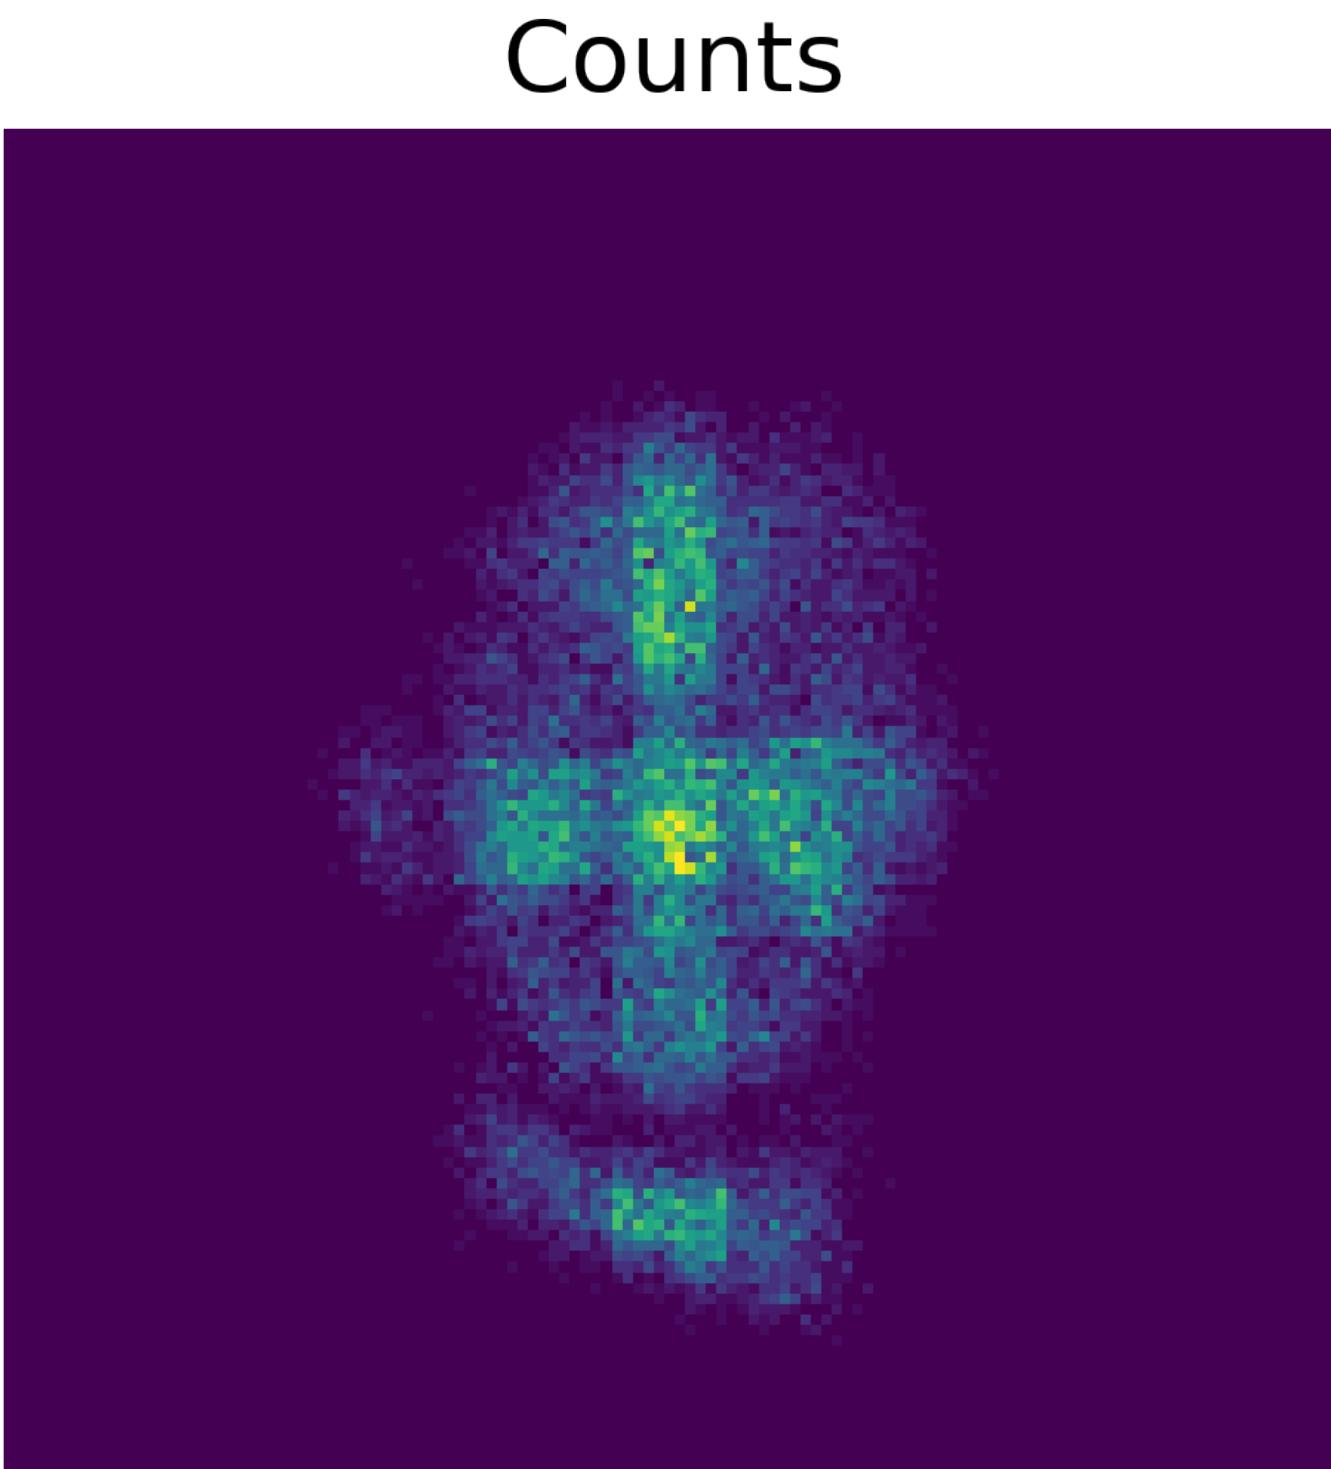
# The “low counts” imaging process



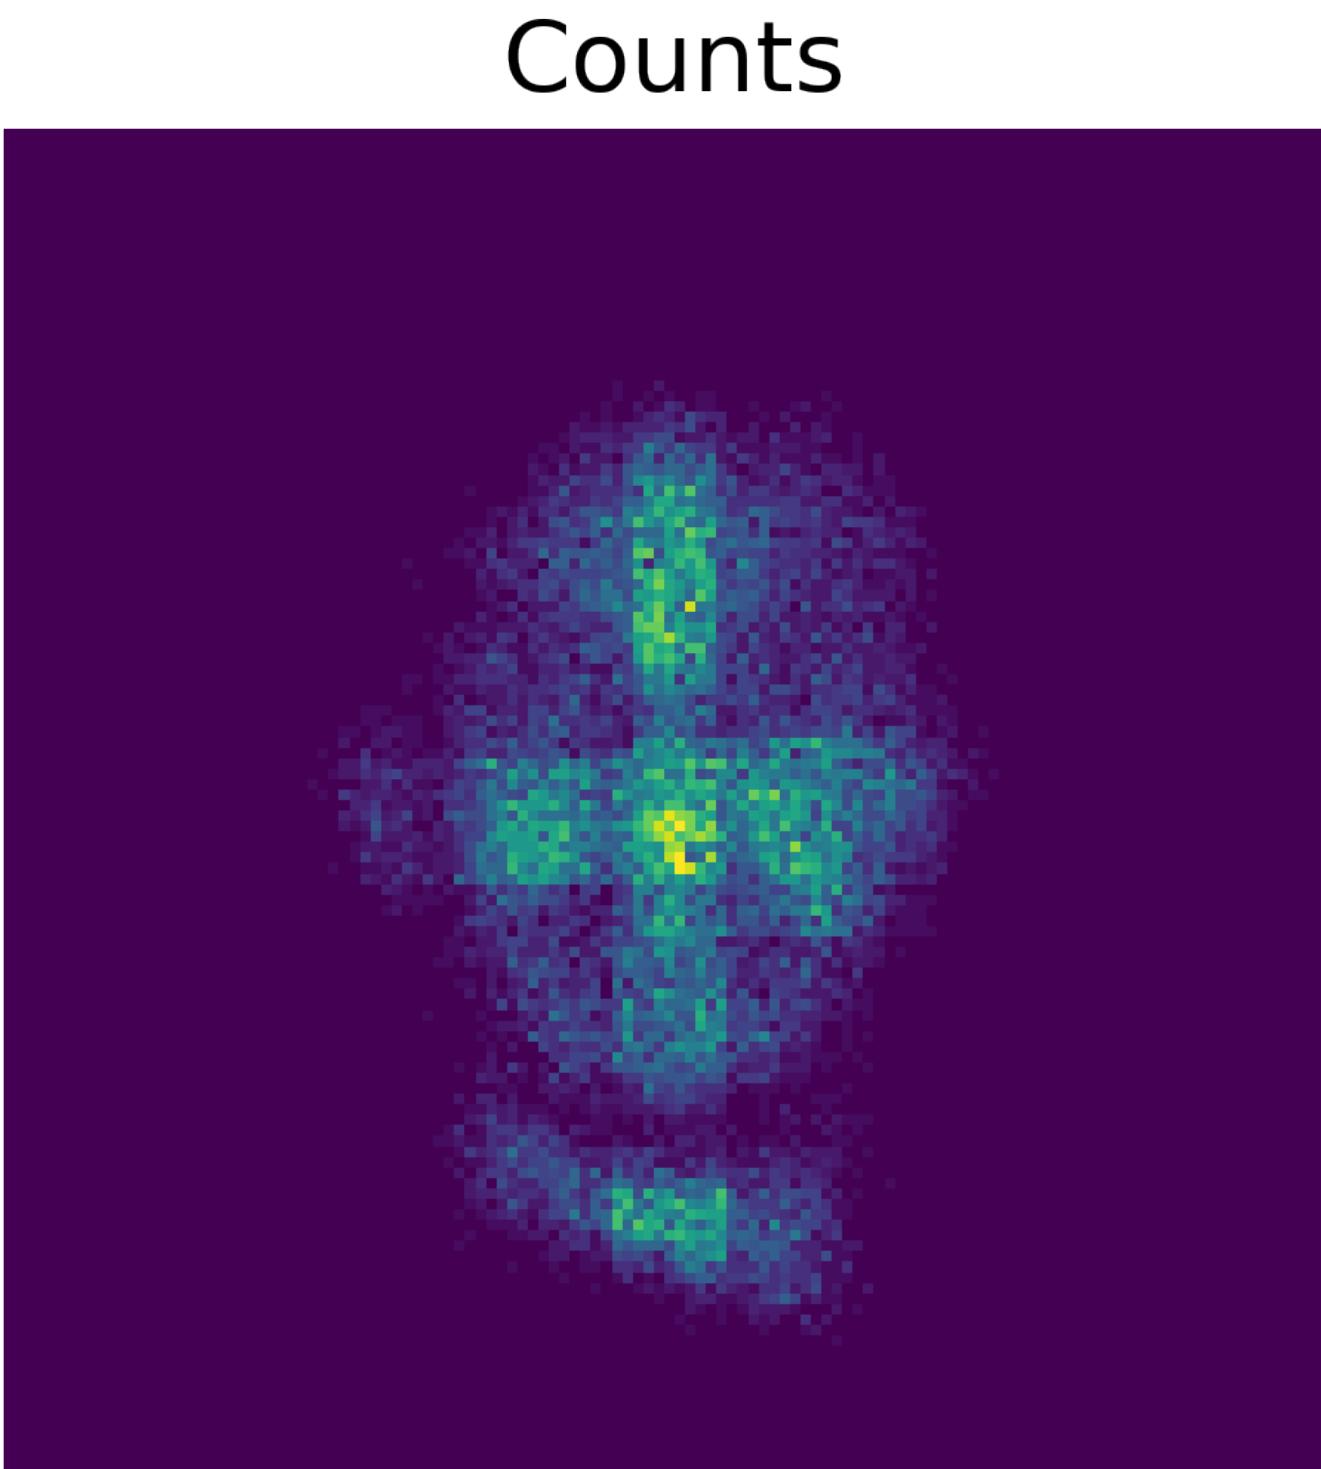
# The “low counts” imaging process



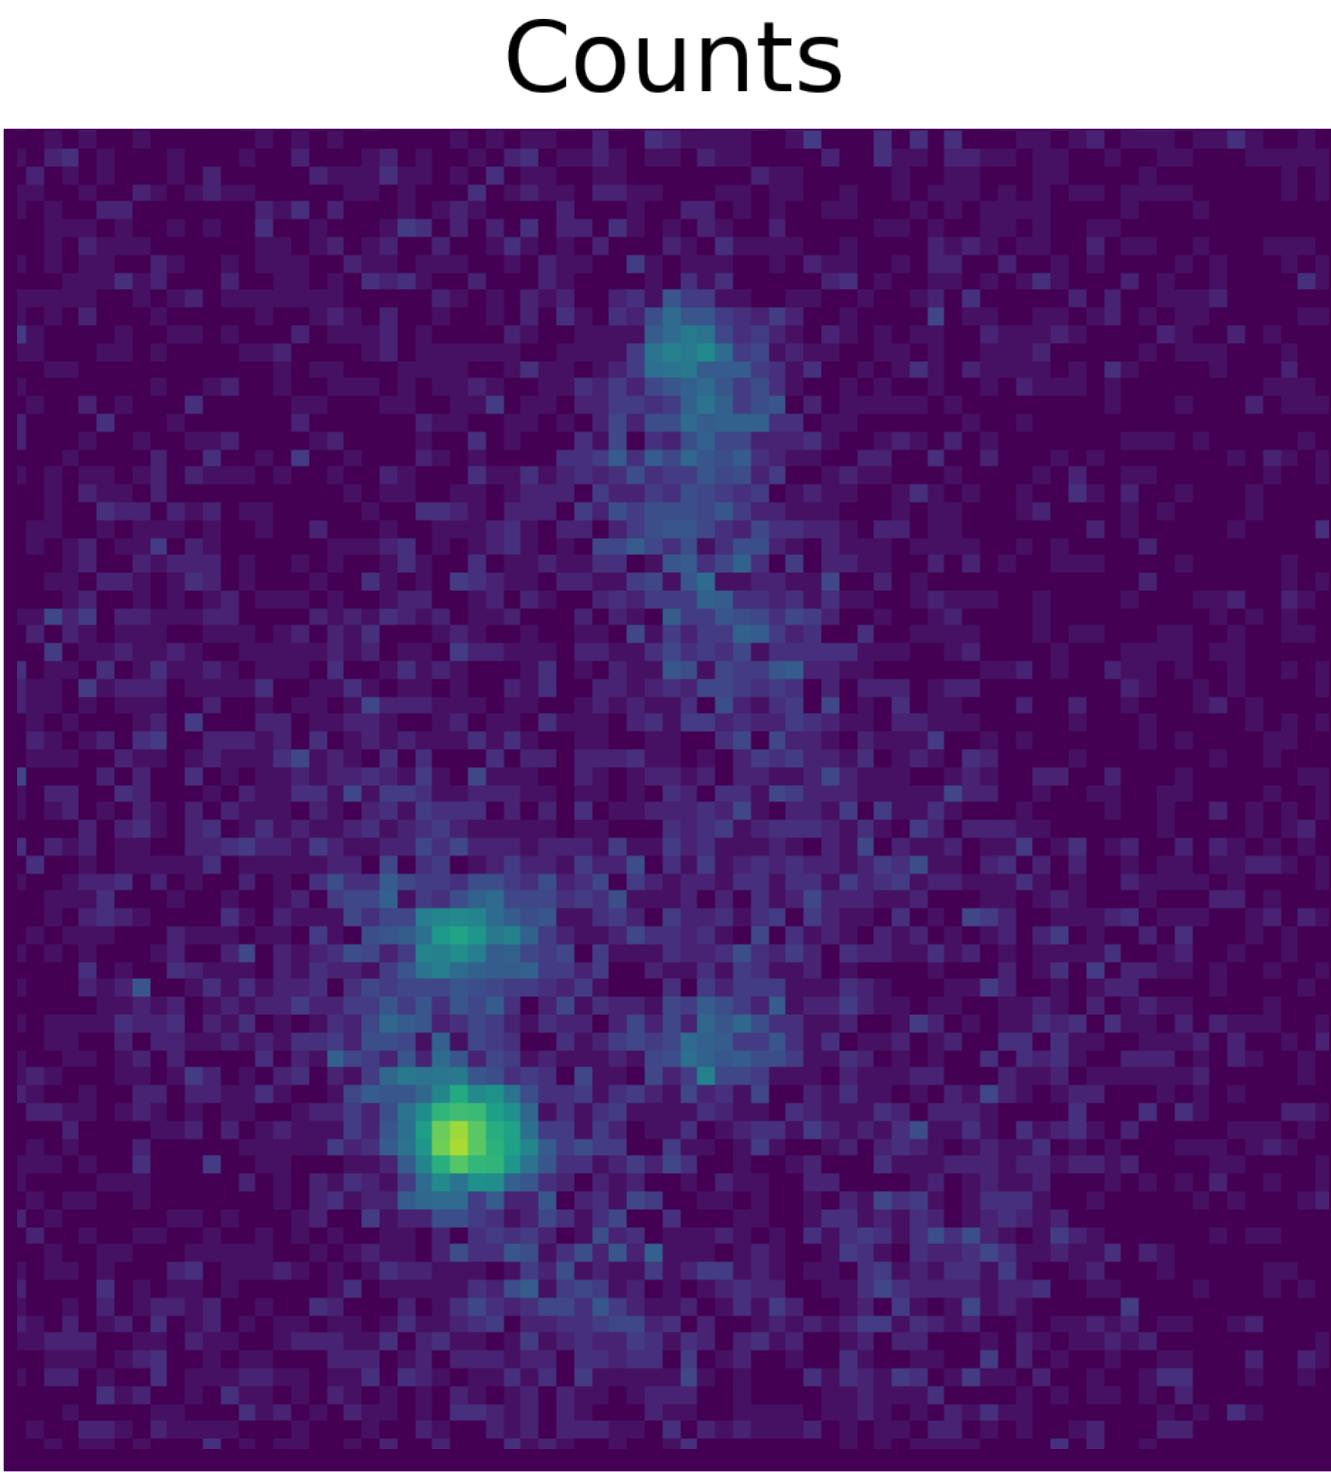
# The problem of “unblind” deconvolution



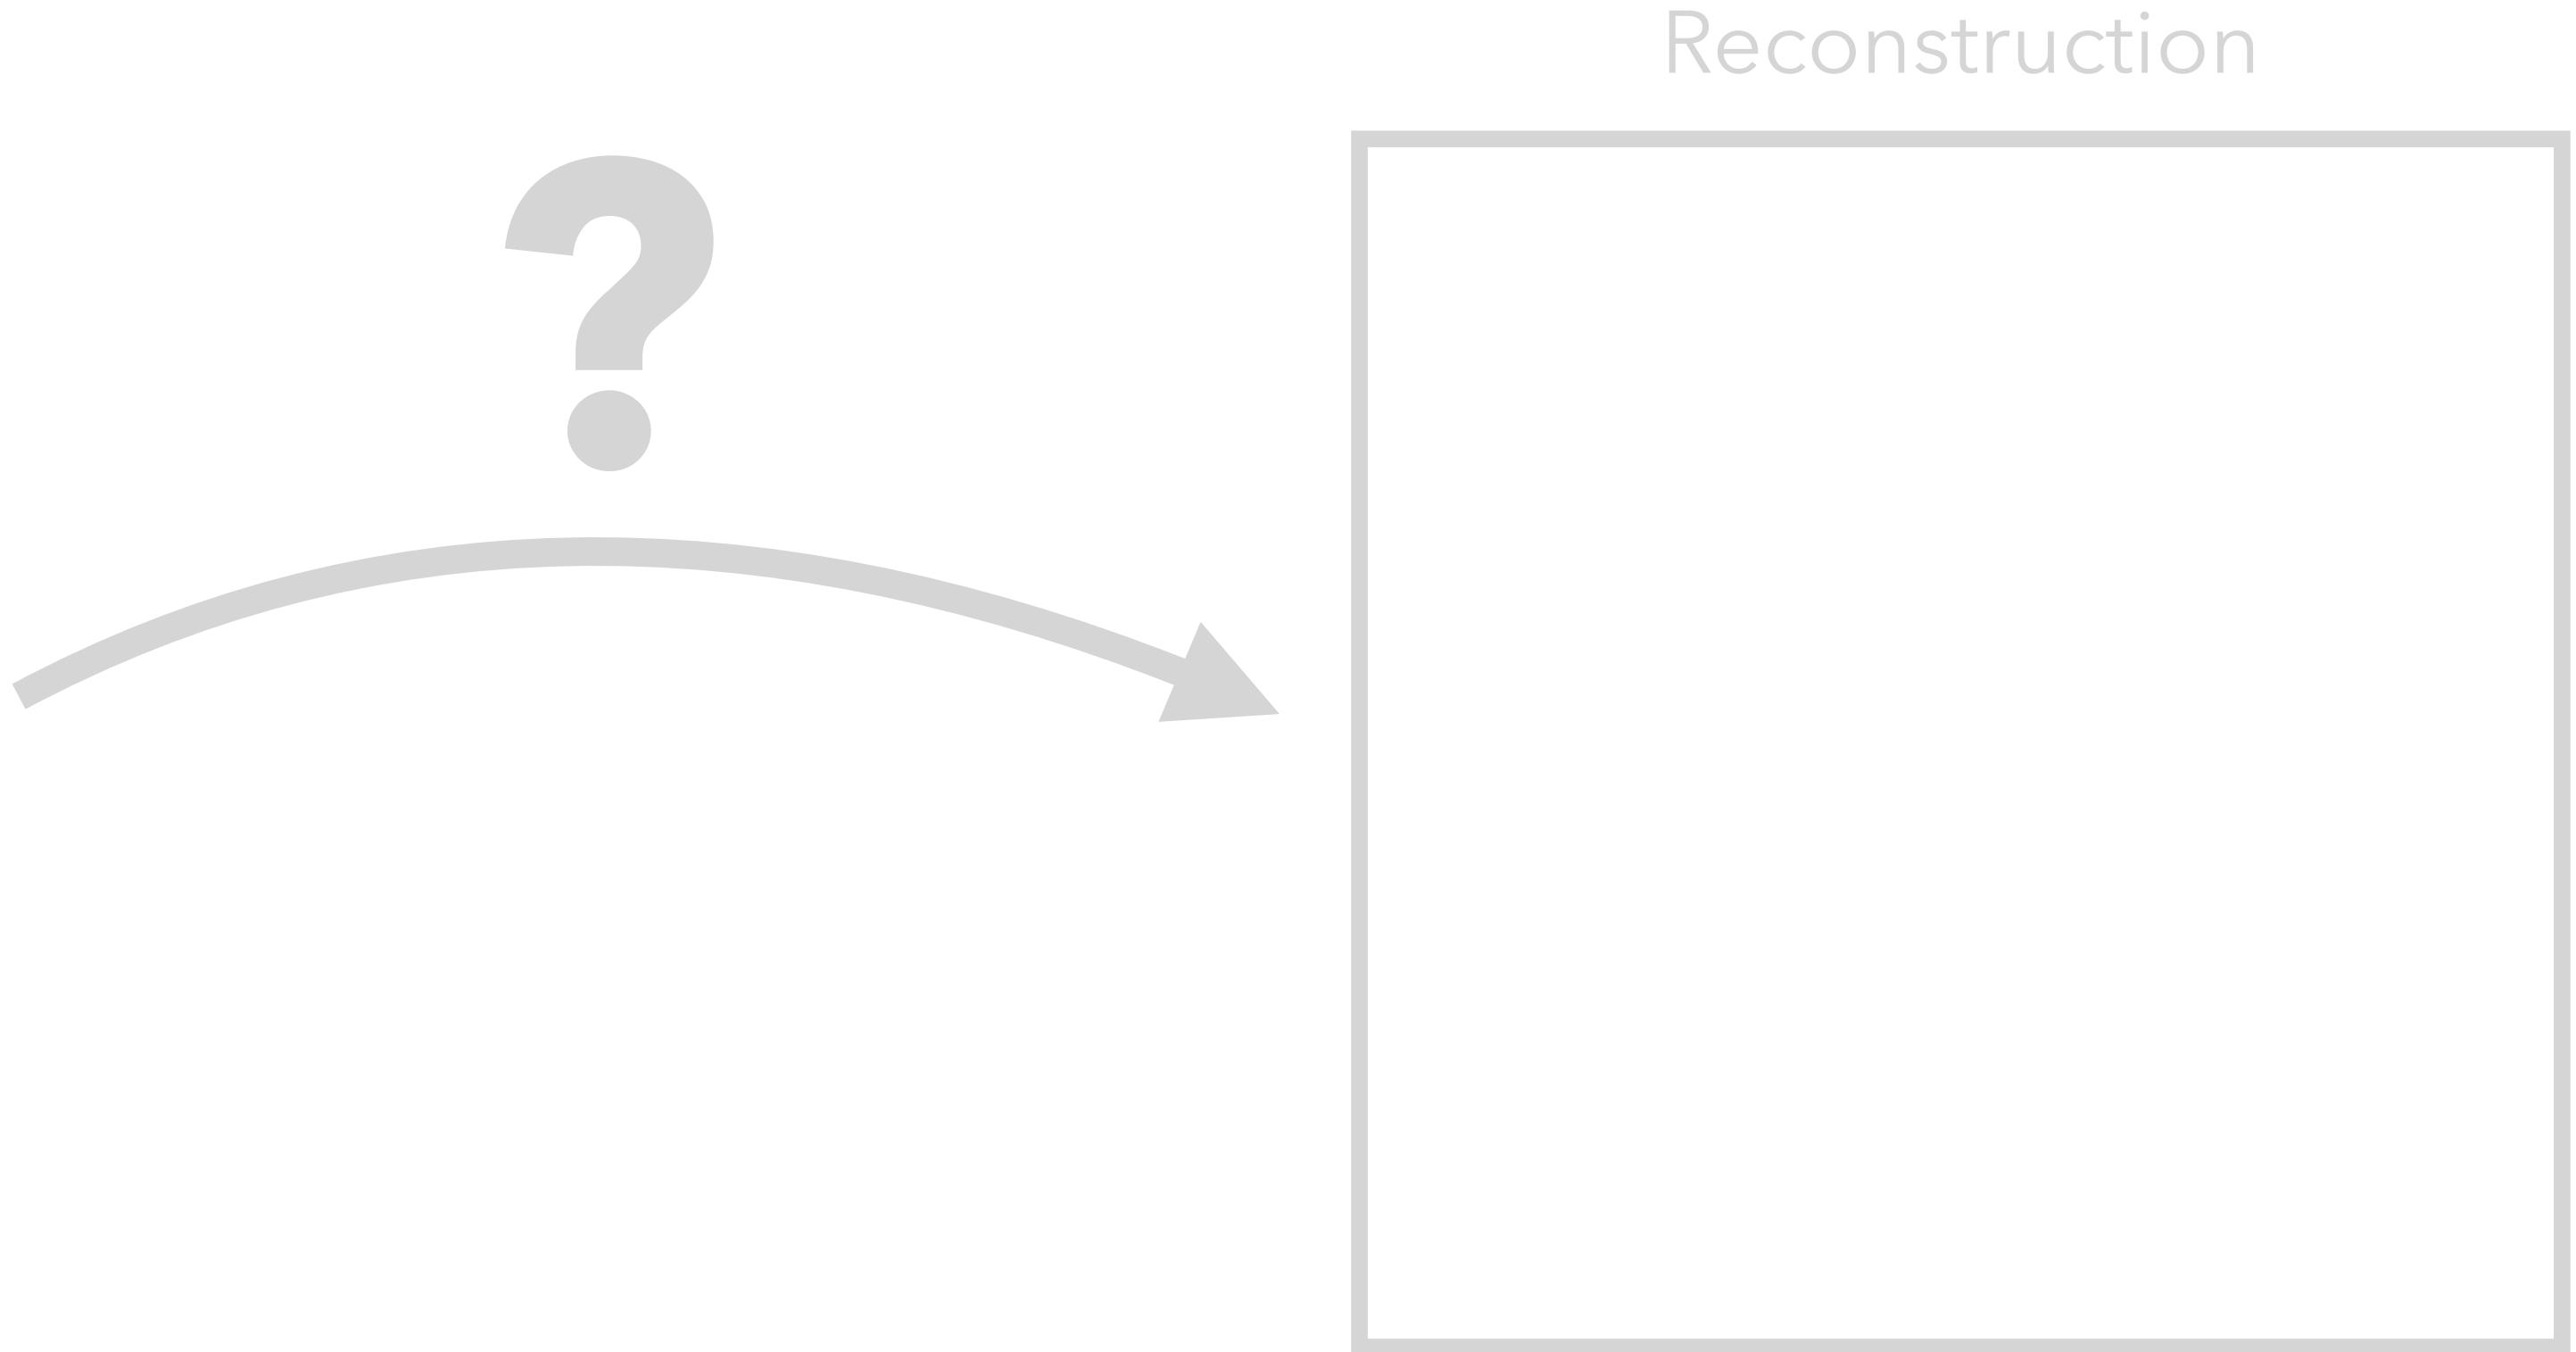
# The problem of “unblind” deconvolution



# The problem of “unblind” deconvolution

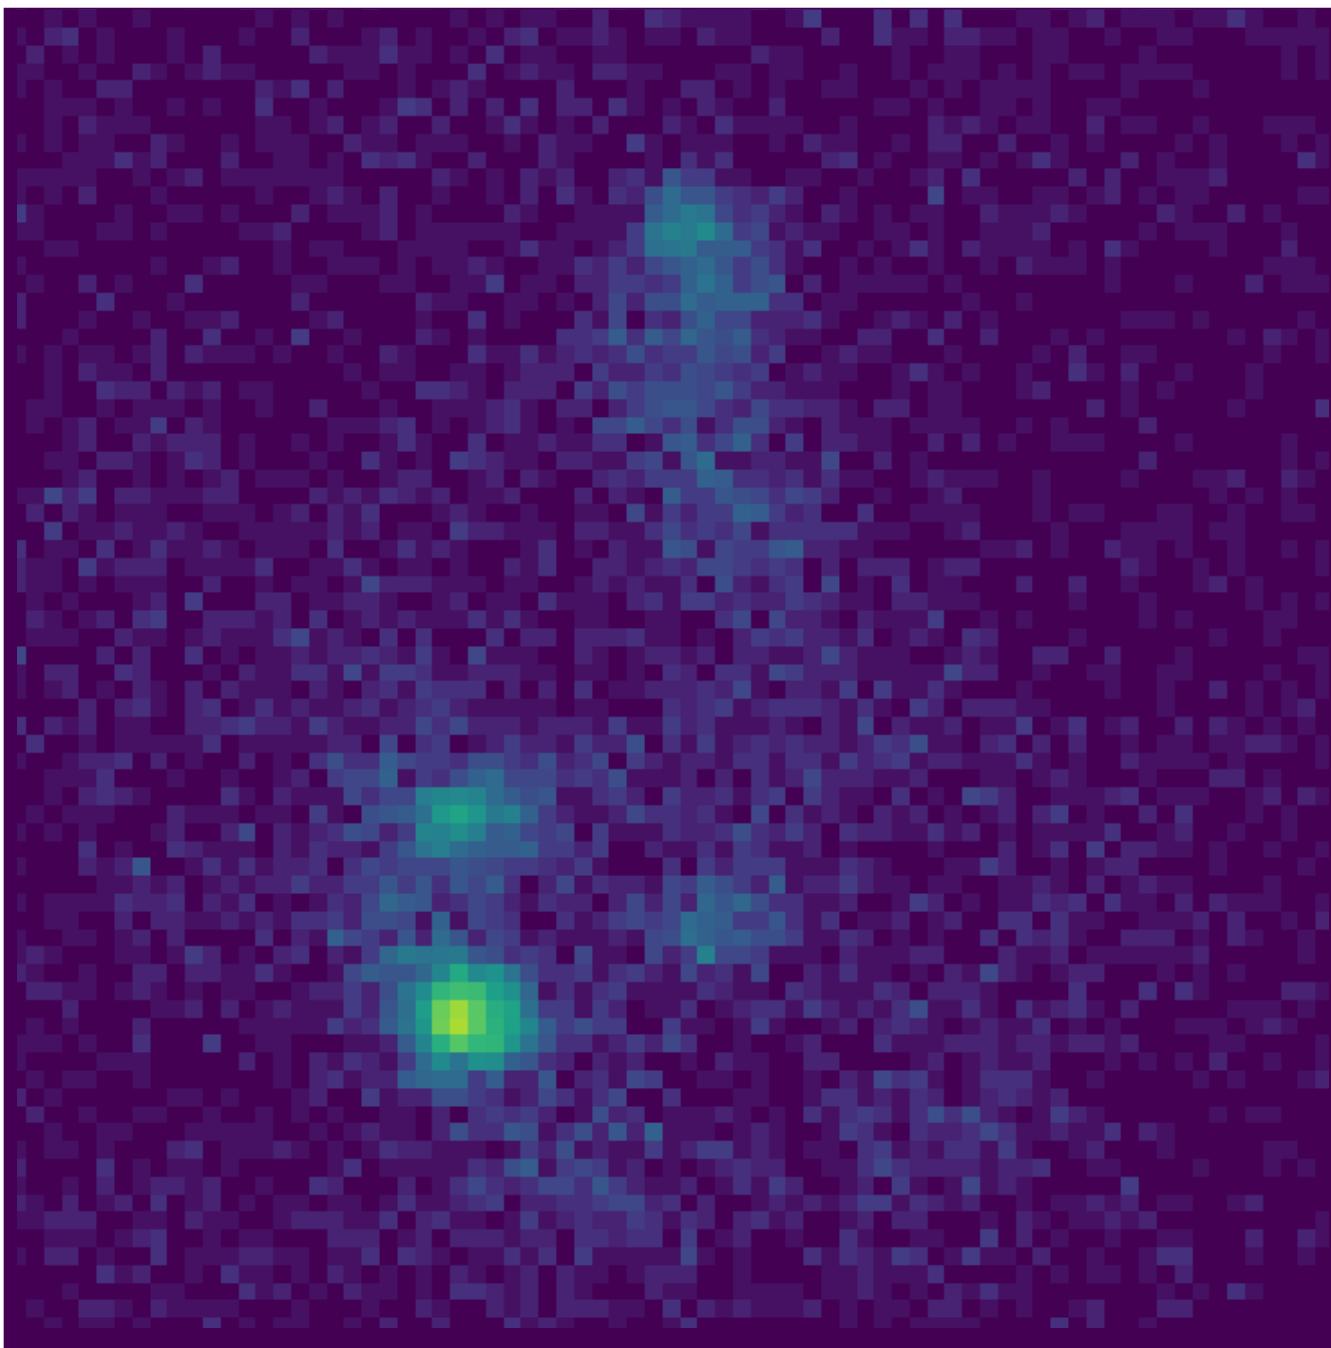


Low counts astronomical images...



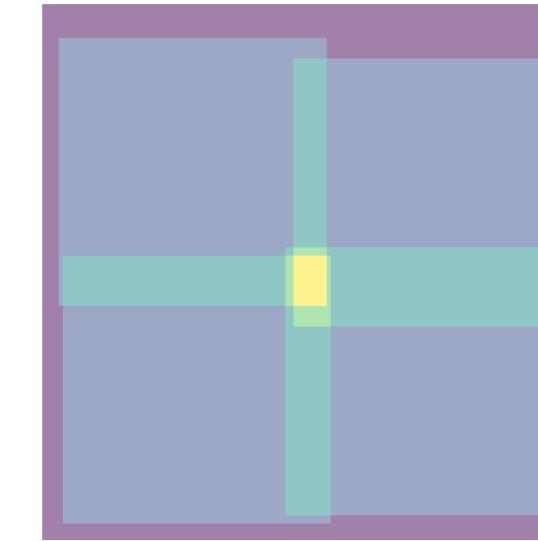
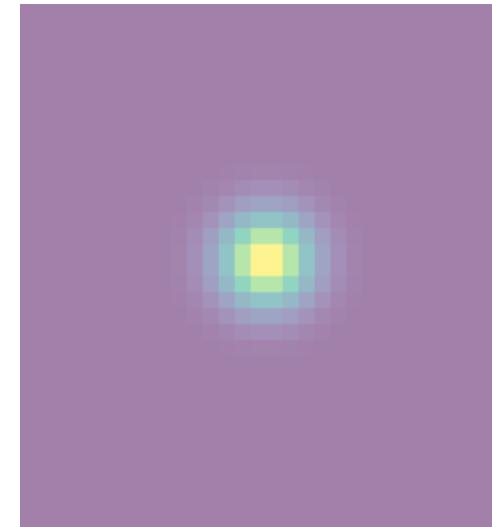
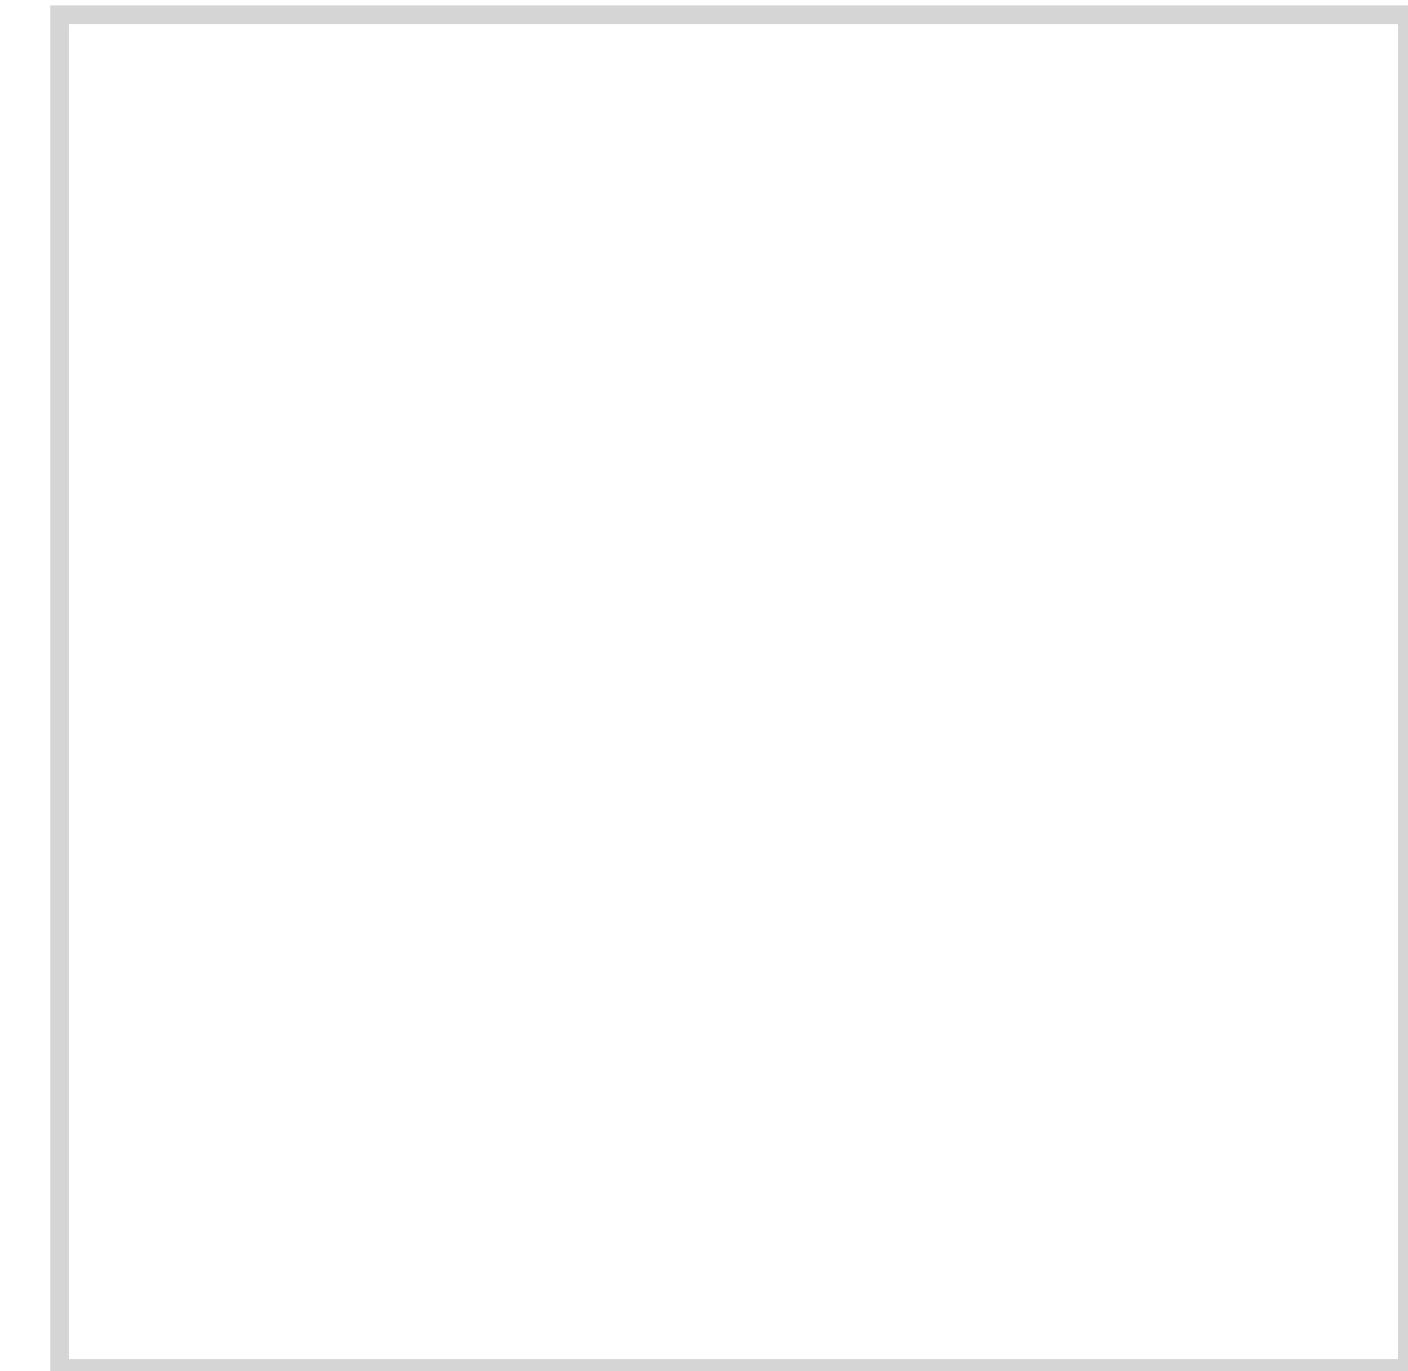
# The problem of “unblind” deconvolution

Counts



Low counts astronomical images...

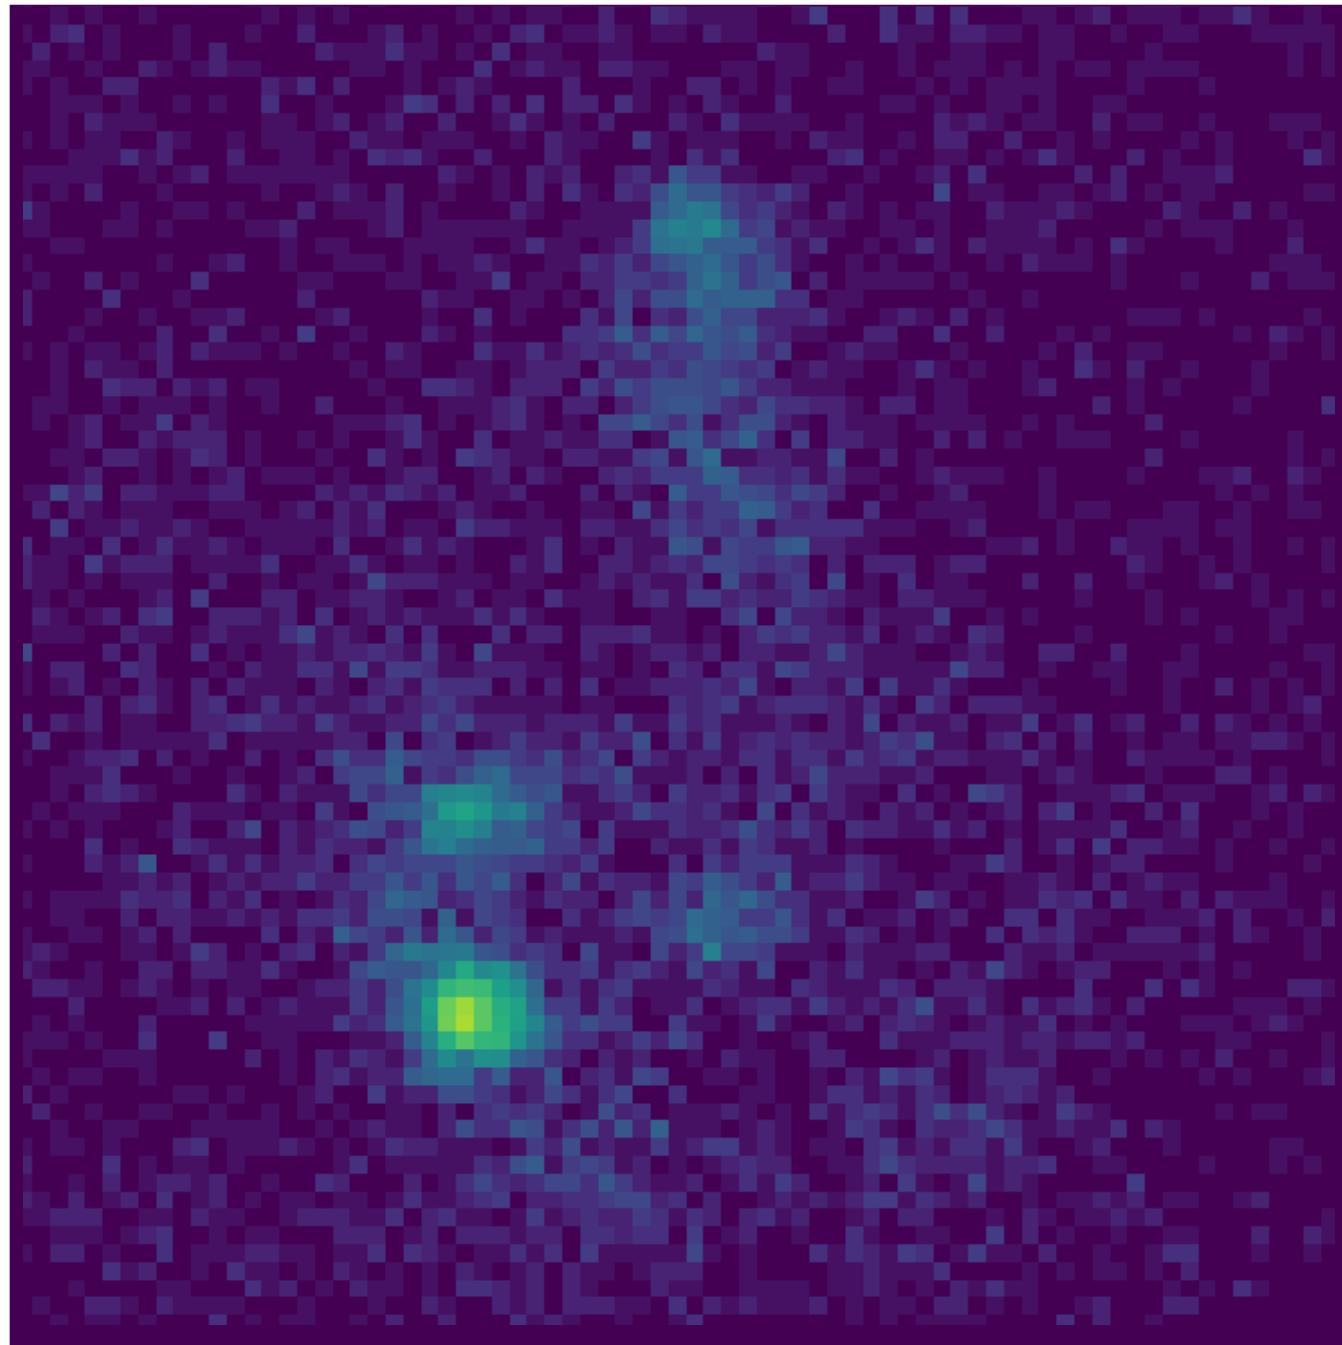
Reconstruction



“Unblind”: PSF and exposure are known or can be simulated

# The problem of “unblind” deconvolution

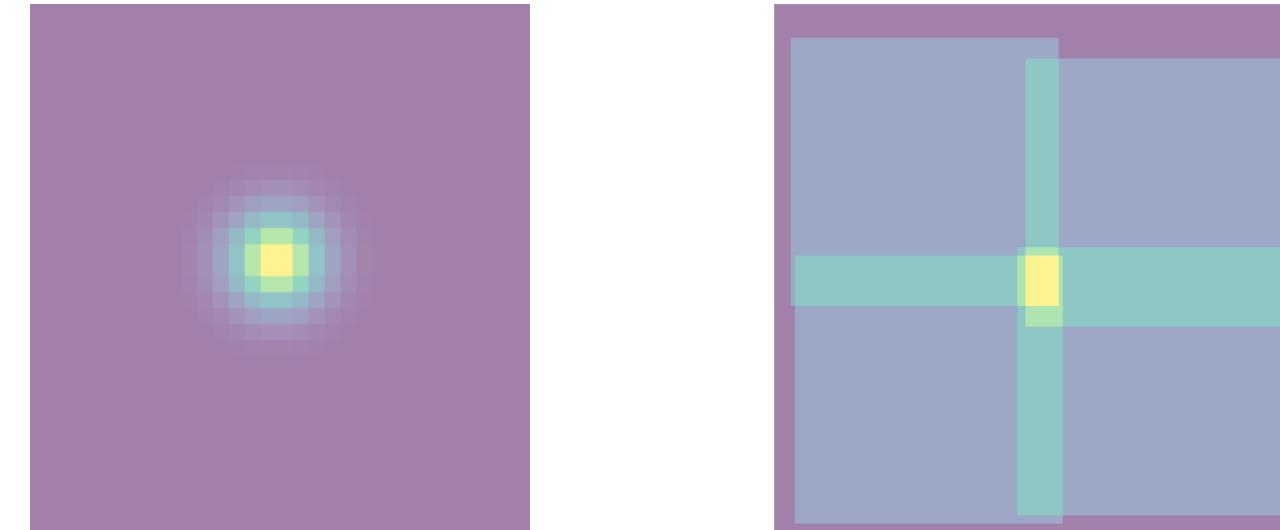
Counts



Low counts astronomical images...

$$\mathcal{L}(\mathbf{d} | \lambda) = \prod_i^N \frac{\lambda_i^{d_i} e^{-\lambda_i}}{\Gamma(d_i + 1)} \quad \lambda = \mathbf{x} \otimes \text{PSF}$$

$$\lambda_i = \sum_k (e_i \cdot (x_i \odot b_i)) p_{i-k}$$



“Unblind”: PSF and exposure are known or can be simulated

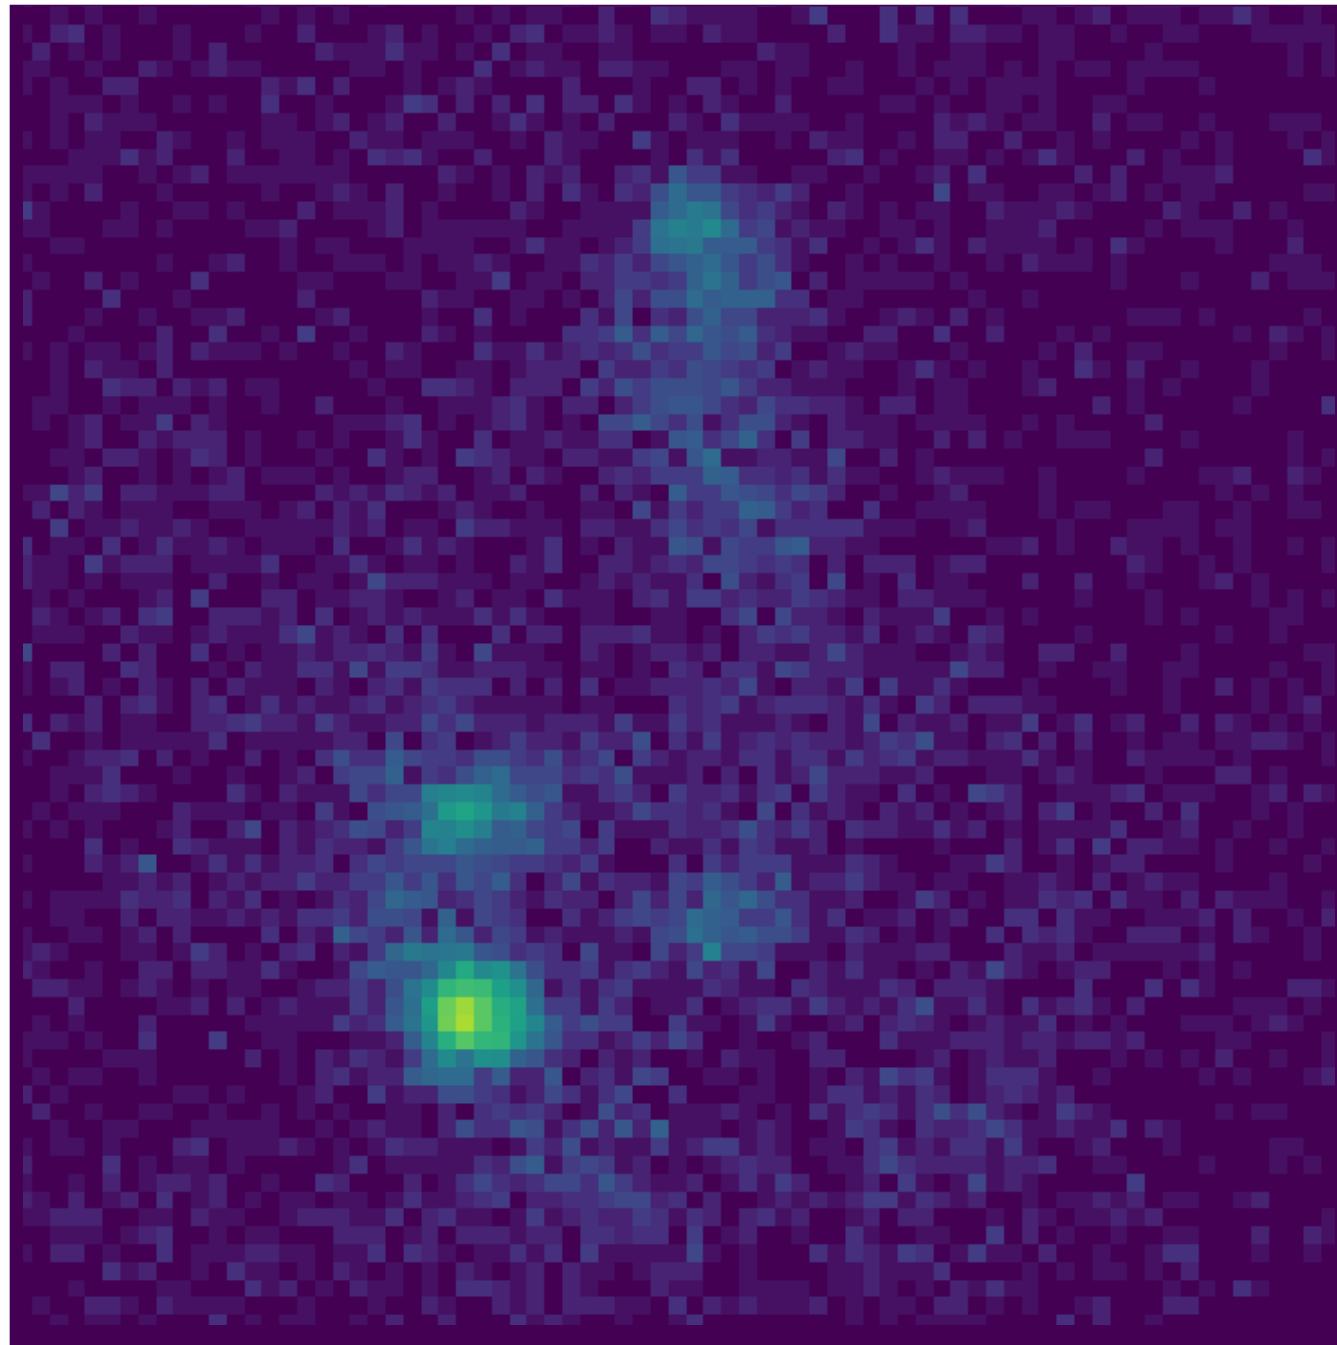
Reconstruction



Needs to be reconstructed using statistical methods

# The problem of “unblind” deconvolution

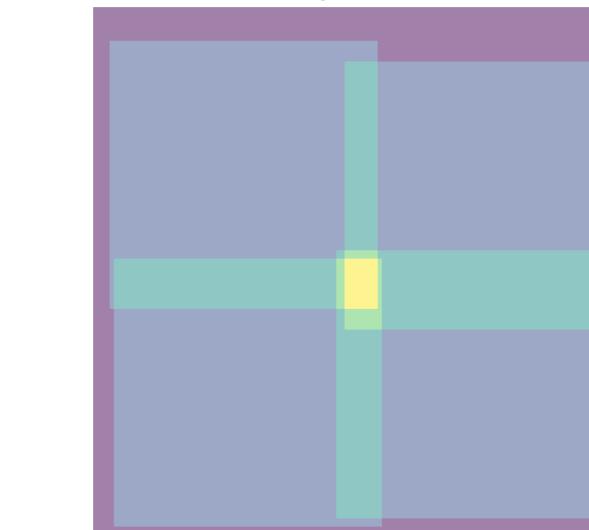
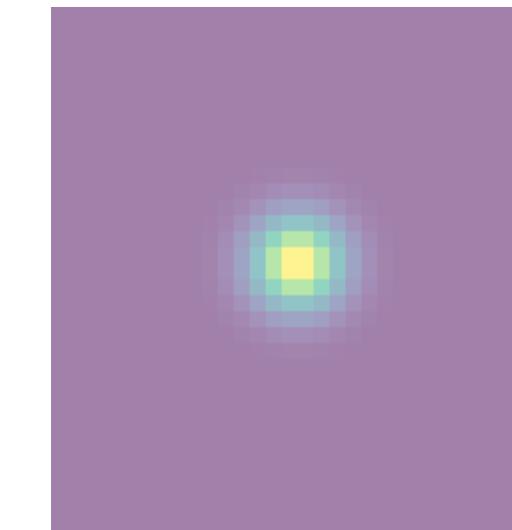
Counts



Low counts astronomical images...

$$\mathcal{L}(\mathbf{d}|\lambda) = \prod_i^N \frac{\lambda_i^{d_i} e^{-\lambda_i}}{\Gamma(d_i+1)} \quad \lambda = \mathbf{x} \otimes \text{PSF}$$

$$\lambda_i = \sum_k (e_i \cdot (x_i \odot b_i)) p_{i-k}$$



“Unblind”: PSF and exposure are known or can be simulated

Reconstruction

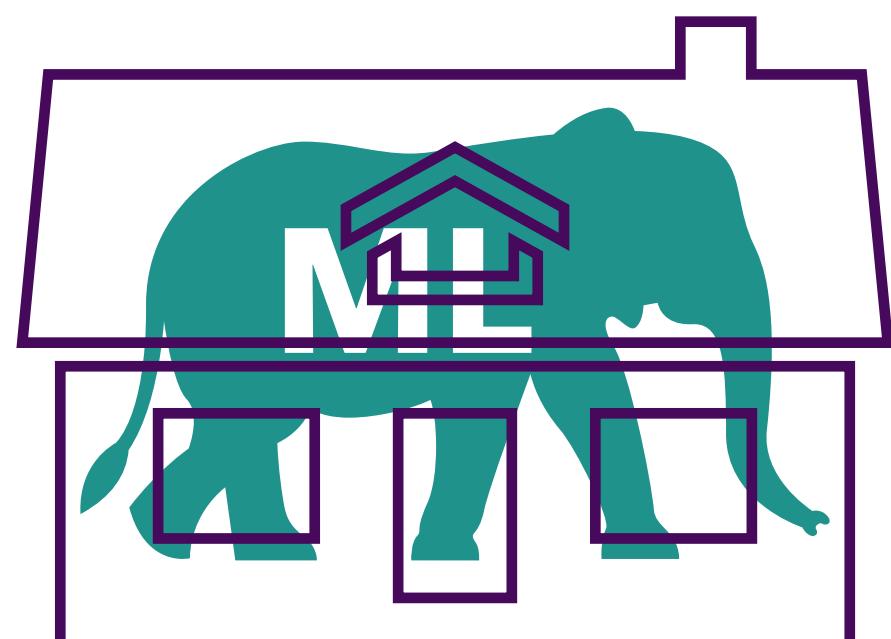


Needs to be reconstructed using statistical methods

“Ill-posed inference problem”

# Why not just use ML?

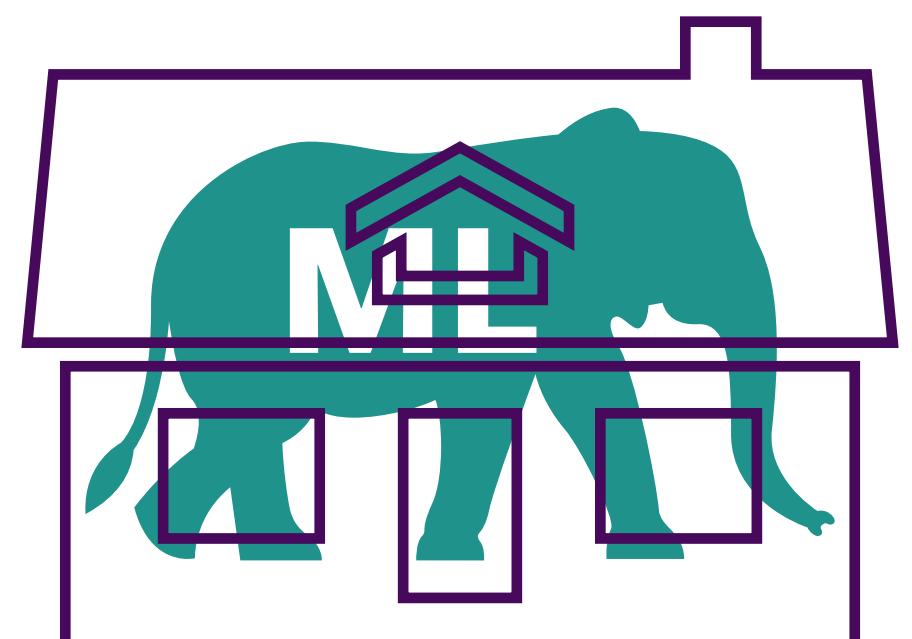
- **Lack of training data:** astronomical images from X-ray and gamma-ray observatories are expensive.
- **There is no ground truth:** we don't know what the sky "really" looks like. We only have the data that is observed and distorted through our instruments.
- **Standard deep CNN architectures are not suited** for the deconvolution task. They can learn a "pseudo inverse PSF kernel" in the spatial domain, but this requires large kernels and adapted architectures [[Xu et al. 2014](#)].
- For Poisson data one is always in the low signal to noise regime. So **specifying "Type I" errors (false detections) is important**. Ideally one can keep track of the full likelihood.
- Needs to learn information that is available a priory, e.g. Poisson statistics of the data.
- However "**transfer learning**" and **simulated training data could maybe work**, or hybrid approaches...



The "MLephant in the room"  
question...

# Why not just use ML?

- **Lack of training data:** astronomical images from X-ray and gamma-ray observatories are expensive.
- **There is no ground truth:** we don't know what the sky "really" looks like. We only have the data that is observed and distorted through our instruments.
- **Standard deep CNN architectures are not suited** for the deconvolution task. They can learn a "pseudo inverse PSF kernel" in the spatial domain, but this requires large kernels and adapted architectures [[Xu et al. 2014](#)].
- For Poisson data one is always in the low signal to noise regime. So **specifying "Type I" errors (false detections) is important**. Ideally one can keep track of the full likelihood.
- Needs to learn information that is available a priori, e.g. Poisson statistics of the data.
- However "**transfer learning**" and **simulated training data could maybe work**, or hybrid approaches...



The "MLephant in the room"  
question...

# Methods

# Some math...

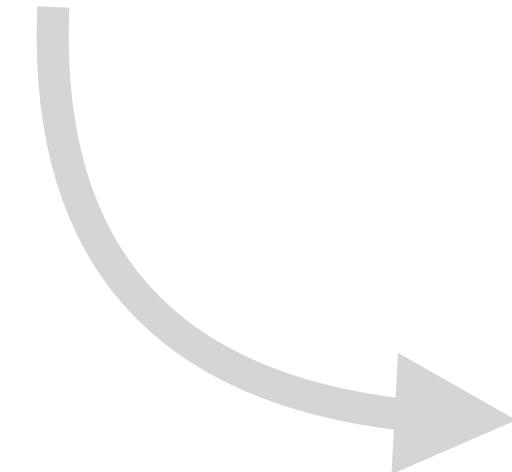
For the counts image assume per  
pixel **Poisson** likelihood:

$$\mathcal{L}(\mathbf{d} | \lambda) = \prod_i^N \frac{\lambda_i^{d_i} e^{-\lambda_i}}{d_i!}$$

# Some math...

For the counts image assume per pixel **Poisson** likelihood:

$$\mathcal{L}(\mathbf{d} | \lambda) = \prod_i^N \frac{\lambda_i^{d_i} e^{-\lambda_i}}{d_i!}$$



As usual: take the **negative log-likelihood**, in Astronomy often call "Cash" statistics

$$\mathcal{C}(\mathbf{d} | \lambda) = \sum_i^N (\lambda_i - d_i \log \lambda_i)$$

# Some math...

For the counts image assume per pixel **Poisson** likelihood:

$$\mathcal{L}(\mathbf{d} | \lambda) = \prod_i^N \frac{\lambda_i^{d_i} e^{-\lambda_i}}{d_i!}$$

$\lambda$  are the "model counts"

$$\lambda = \mathbf{x} \circledast \text{PSF}$$

$\mathbf{x}$  is the reconstructed image we are looking for. Consider each pixel  $x_i$  as independent parameter in the model...

As usual: take the **negative log-likelihood**, in Astronomy often call "Cash" statistics

$$\mathcal{C}(\mathbf{d} | \lambda) = \sum_i^N (\lambda_i - d_i \log \lambda_i)$$

# Some math...

For the counts image assume per pixel **Poisson** likelihood:

$$\mathcal{L}(\mathbf{d} | \lambda) = \prod_i^N \frac{\lambda_i^{d_i} e^{-\lambda_i}}{d_i!}$$

$\lambda$  are the "model counts"

$$\lambda = \mathbf{x} \circledast \text{PSF}$$

$\mathbf{x}$  is the reconstructed image we are looking for. Consider each pixel  $x_i$  as independent parameter in the model...

As usual: take the **negative log-likelihood**, in Astronomy often call "Cash" statistics

$$\mathcal{C}(\mathbf{d} | \lambda) = \sum_i^N (\lambda_i - d_i \log \lambda_i)$$

E.g. could be solved by "**Gradient Descent**"...

$$\mathbf{x}_{n+1} = \mathbf{x}_n - \alpha \cdot \frac{\partial \mathcal{C}(\mathbf{d} | \mathbf{x})}{\partial x_i}$$

# Some math...

For the counts image assume per pixel **Poisson** likelihood:

$$\mathcal{L}(\mathbf{d} | \lambda) = \prod_i^N \frac{\lambda_i^{d_i} e^{-\lambda_i}}{d_i!}$$

$\lambda$  are the "model counts"

$$\lambda = \mathbf{x} \circledast \text{PSF}$$

$\mathbf{x}$  is the reconstructed image we are looking for. Consider each pixel  $x_i$  as independent parameter in the model...

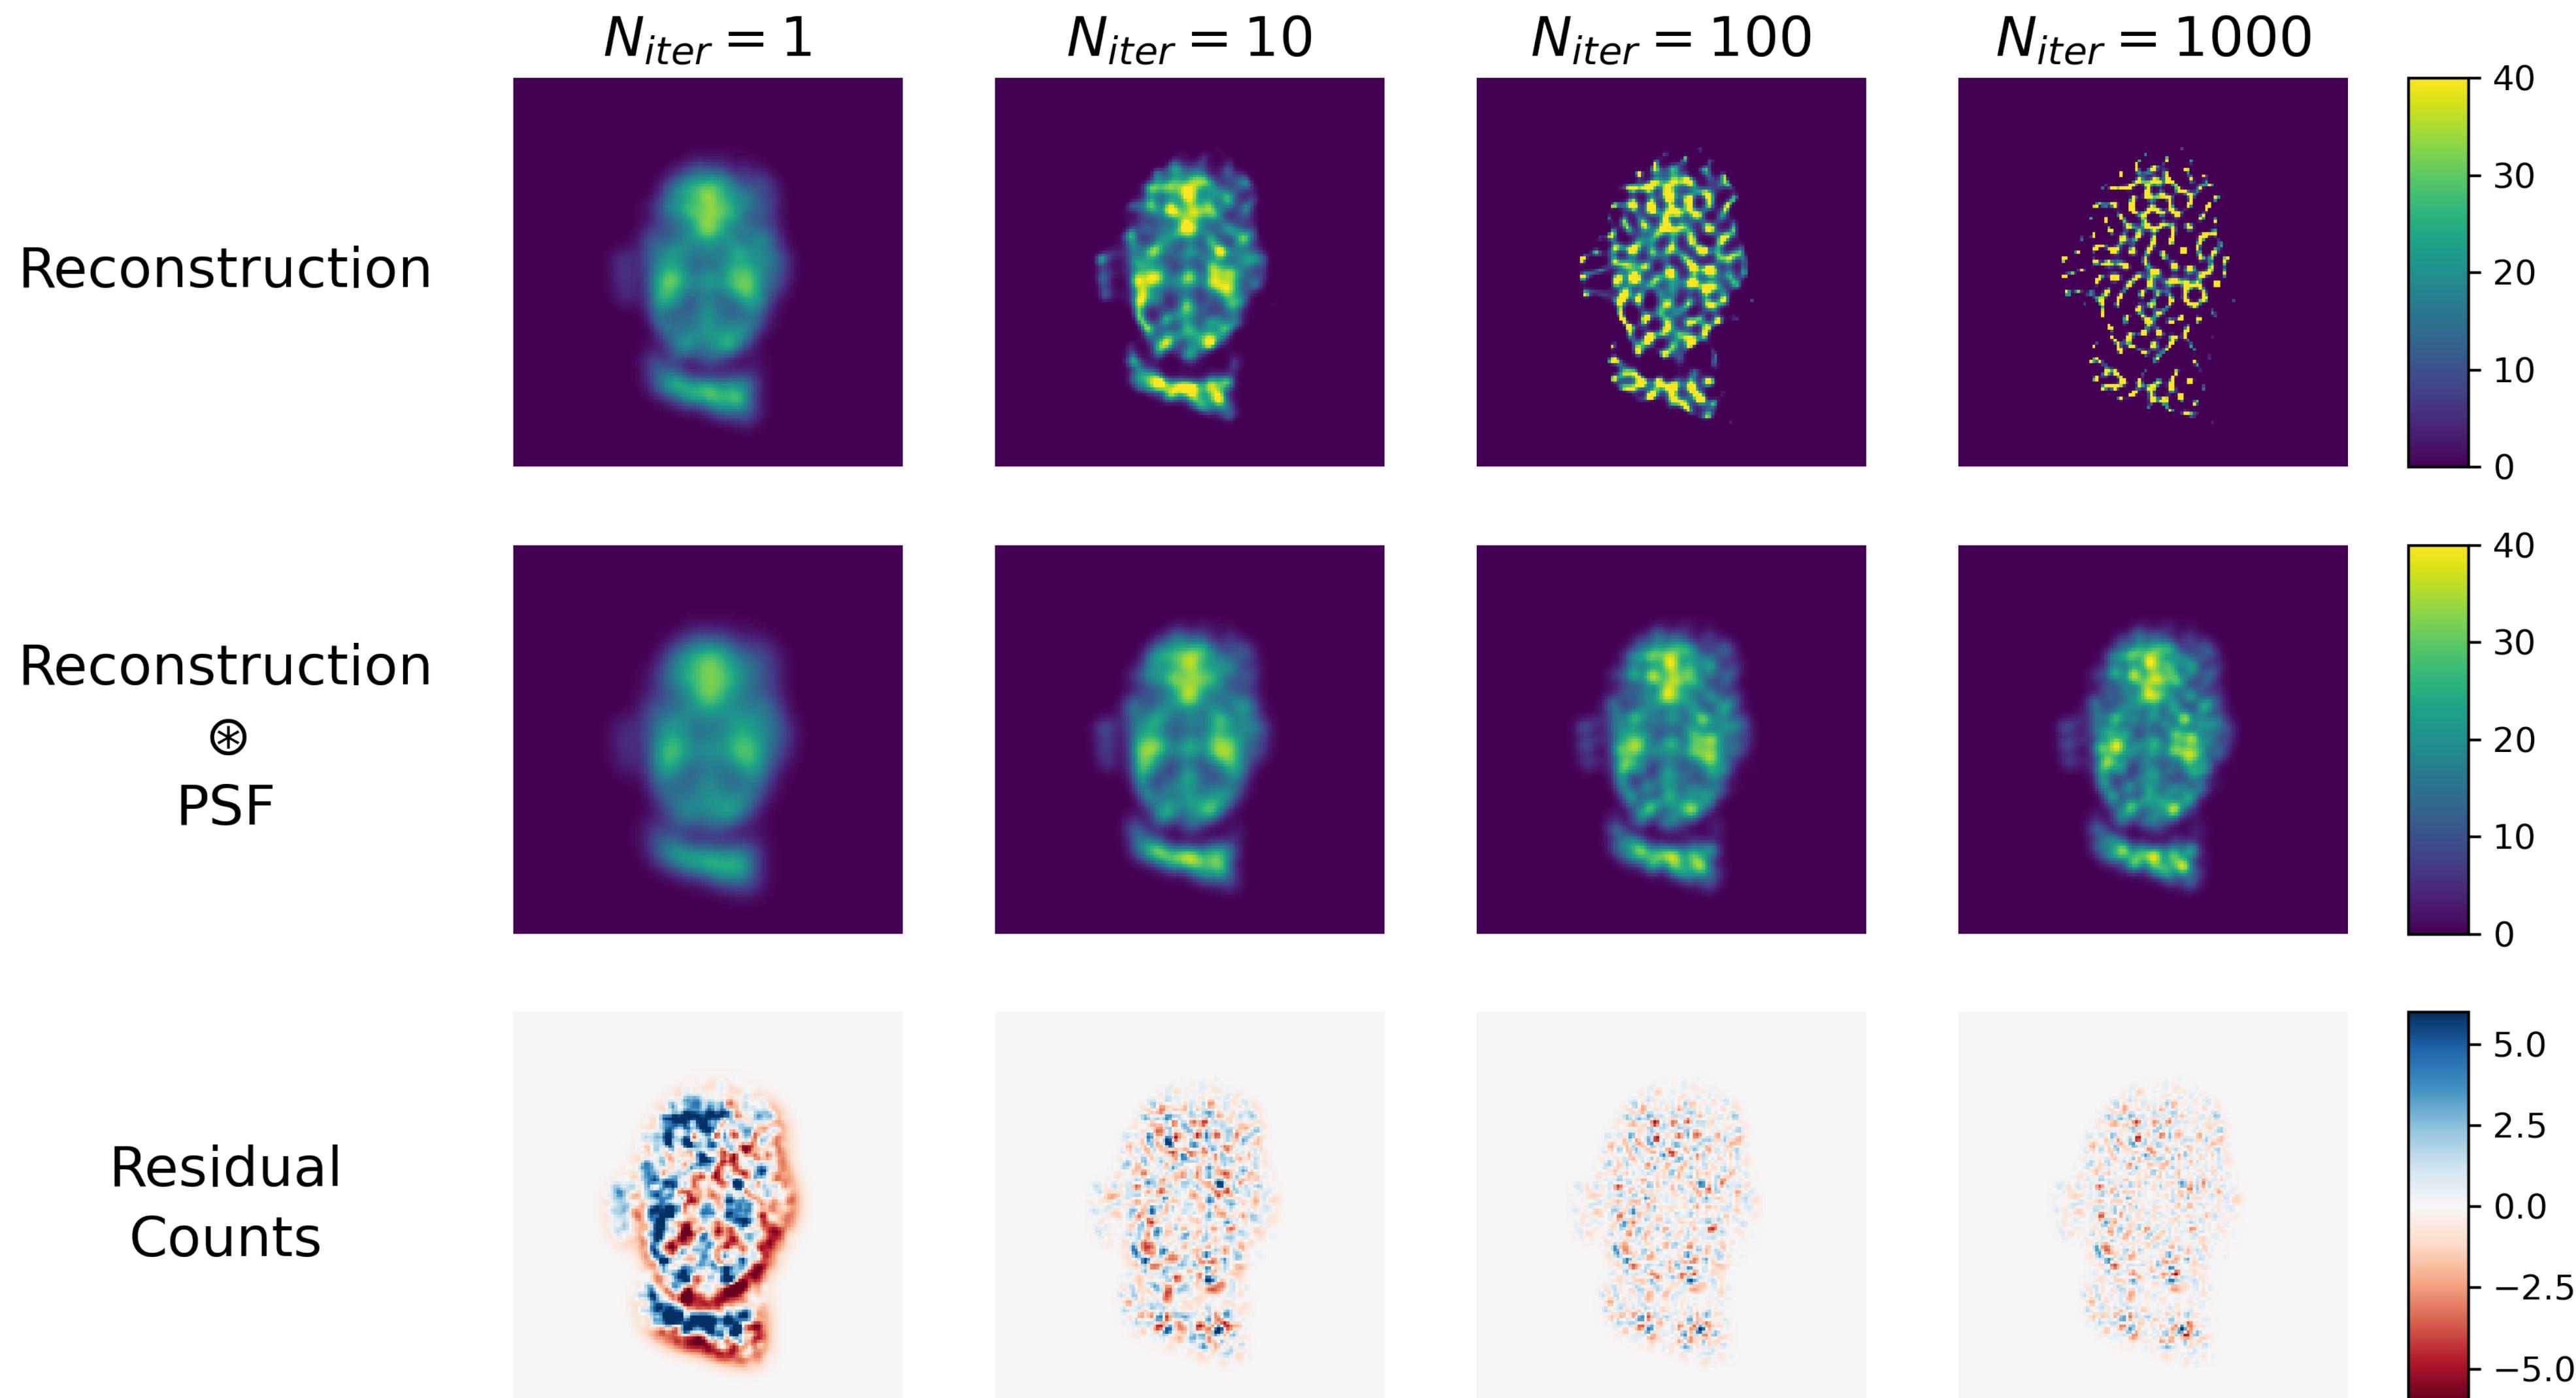
As usual: take the **negative log-likelihood**, in Astronomy often call "Cash" statistics

$$\mathcal{C}(\mathbf{d} | \lambda) = \sum_i^N (\lambda_i - d_i \log \lambda_i)$$

...or using **Expectation Maximisation (EM)**  
Proposed by [[Richardson 1972](#)] & [[Lucy 1974](#)]

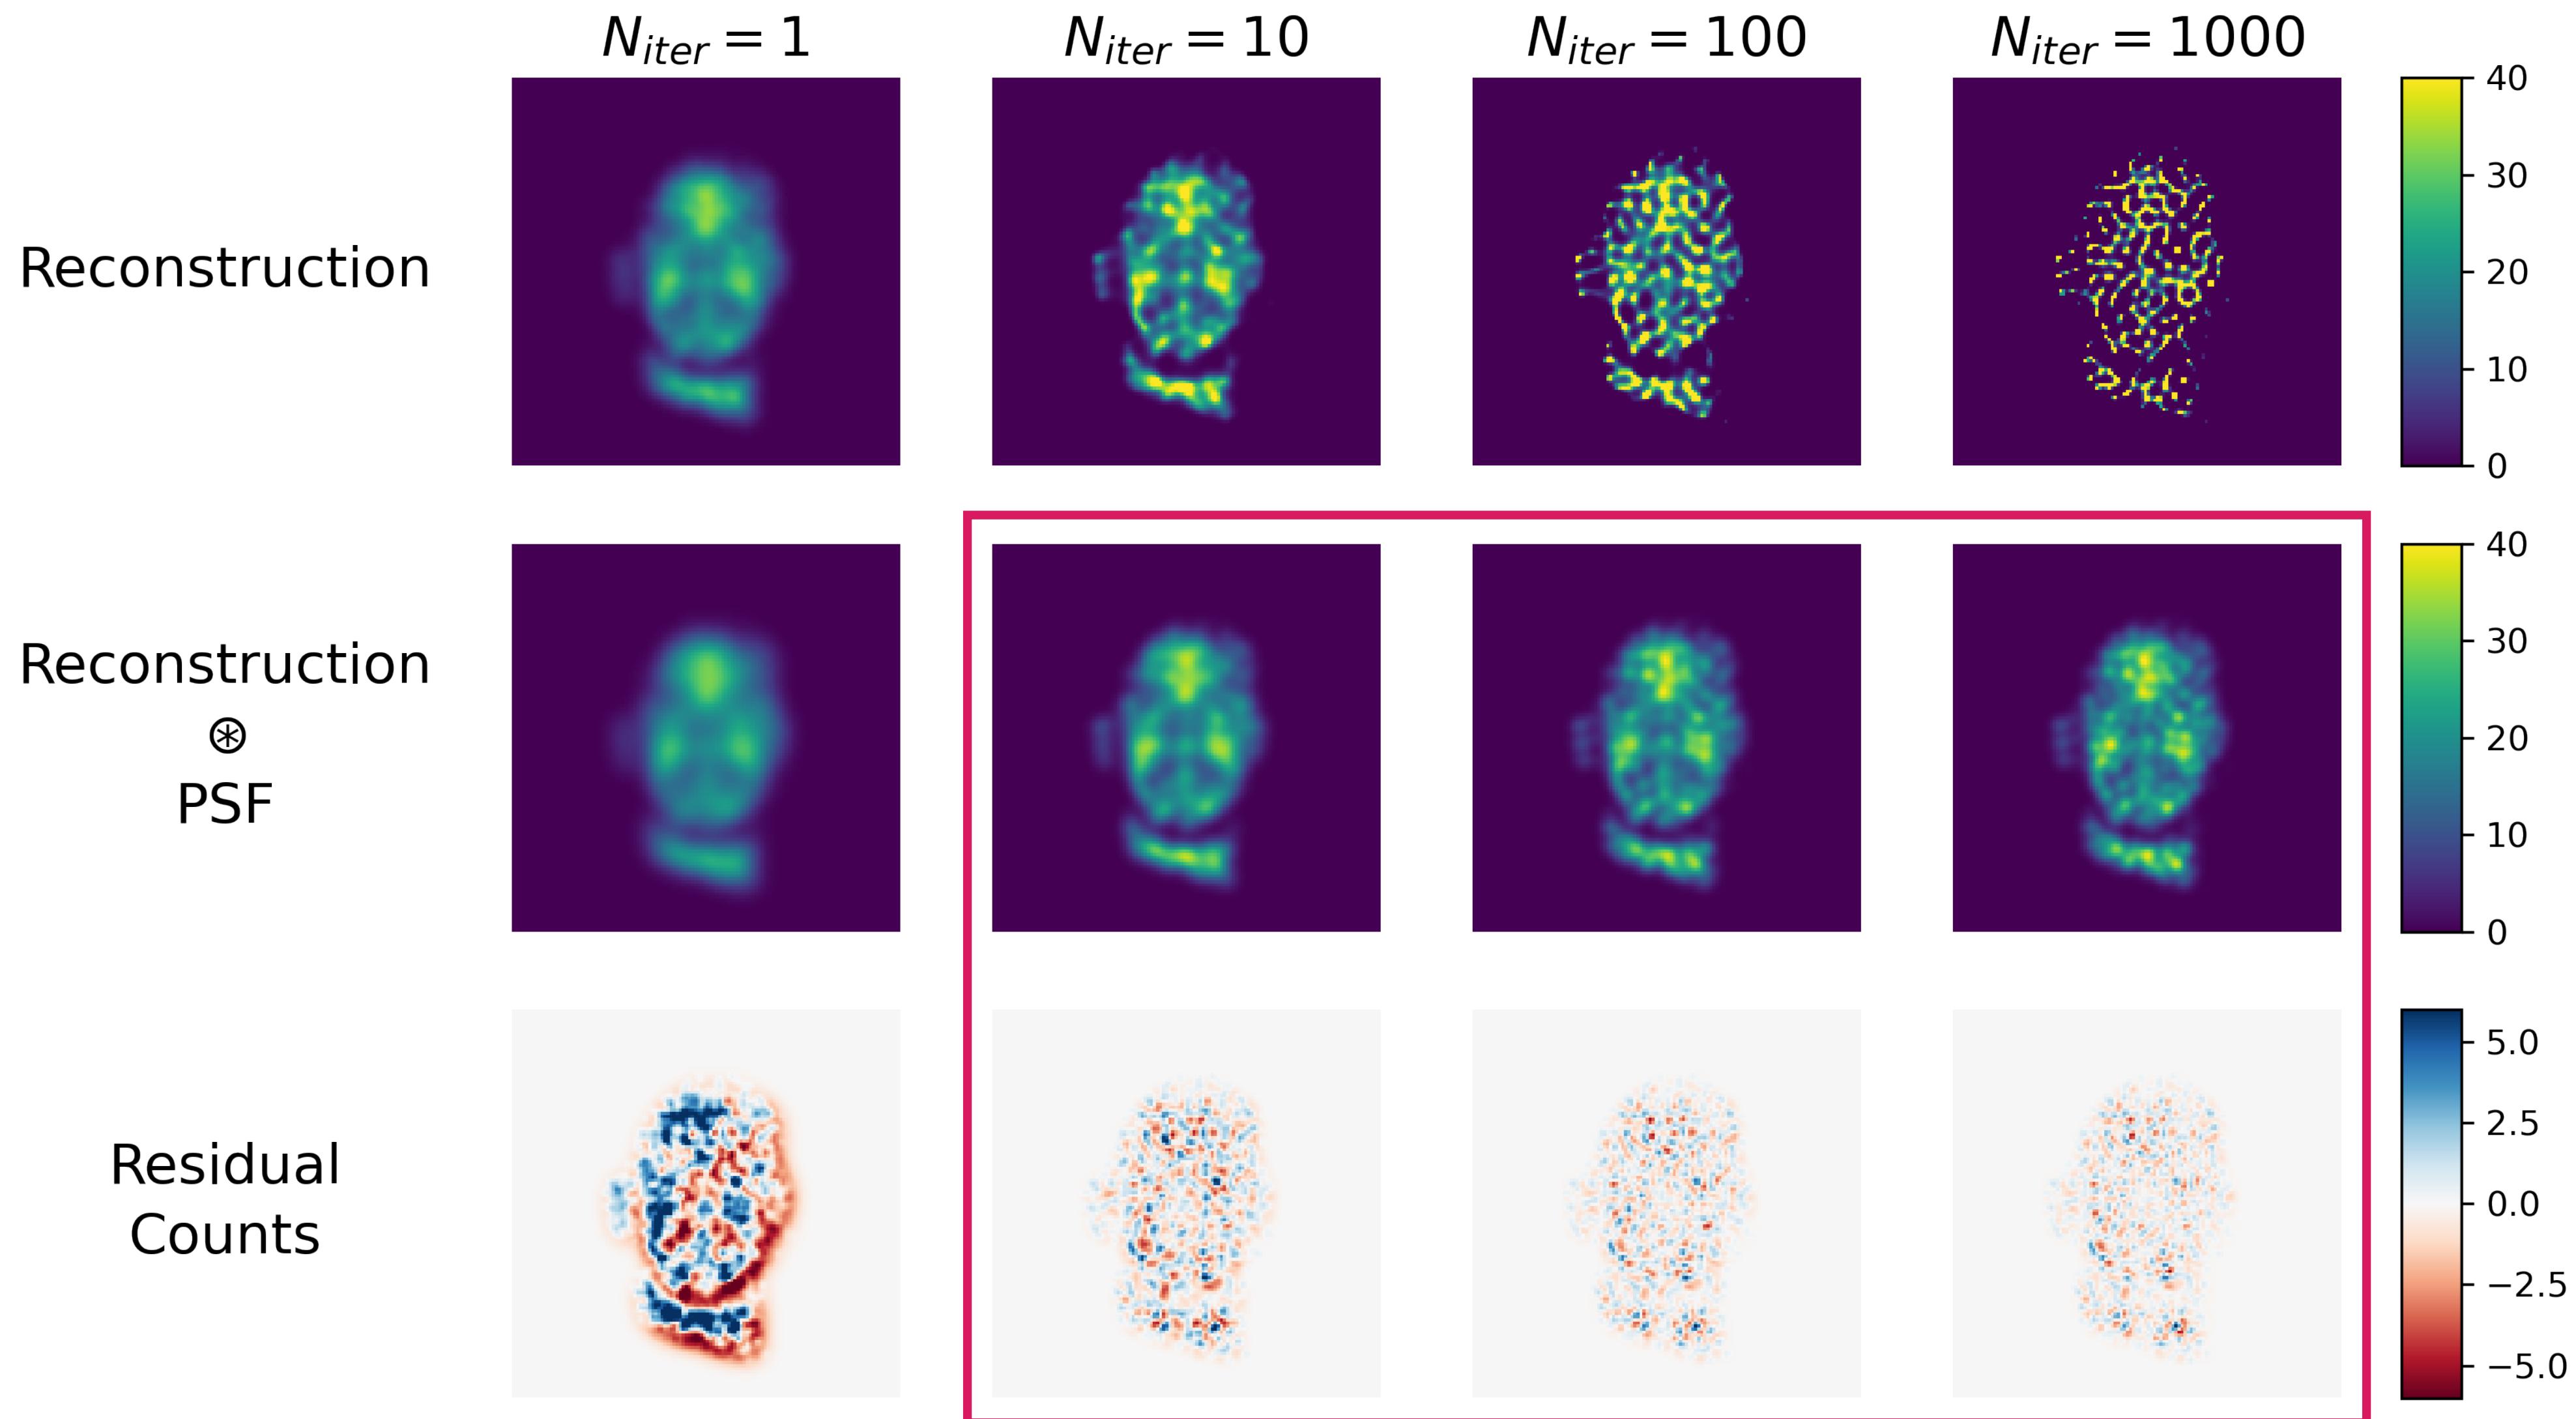
$$\mathbf{x}_{n+1} = \mathbf{x}_n \frac{\mathbf{d}}{\mathbf{x}_n \circledast \text{PSF}} \circledast \text{PSF}^T$$

# RL reconstruction quality



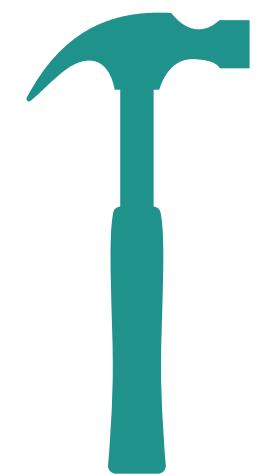
Uses [[skimage.restoration.richardson\\_lucy](#)]

# RL reconstruction quality



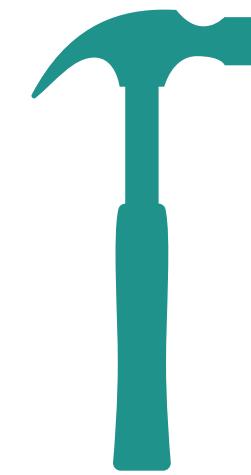
All show good residuals and model counts. But reconstructions very different...

# Possible solutions



**Pragmatic:** early stopping, limit number of iterations until one visually “likes” the result. Often used in literature, good enough for showing nice images...

# Possible solutions

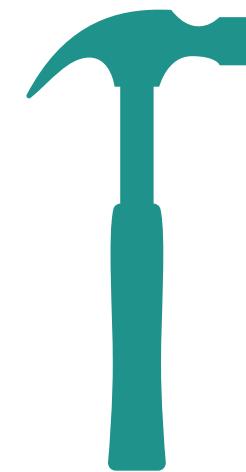


**Pragmatic:** early stopping, limit number of iterations until one visually “likes” the result. Often used in literature, good enough for showing nice images...



“Engineered”: adding additional regularization terms to the likelihood function, such [[total variation](#)], [[Tikhonov regularization](#)] etc. often require additional “hand tuned” parameters...

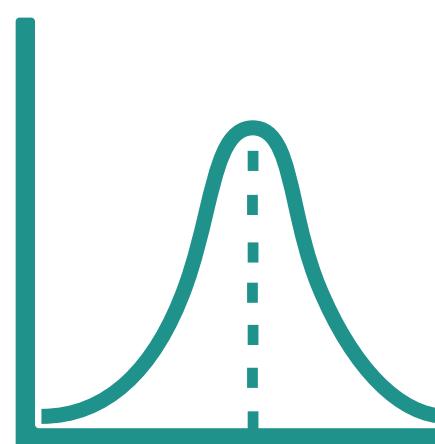
# Possible solutions



**Pragmatic:** early stopping, limit number of iterations until one visually “likes” the result. Often used in literature, good enough for showing nice images...



**“Engineered”:** adding additional regularization terms to the likelihood function, such [[total variation](#)], [[Tikhonov regularization](#)] etc. often require additional “hand tuned” parameters



**“Statistically closed”:** include prior information or assumptions and change to a fully Bayesian treatment of the problem...

# LIRA method

(L)ow counts (I)mage (R)econstruction and (A)analysis

Objective function **extended by a  
log-prior term**, only depending on  $\lambda$  :

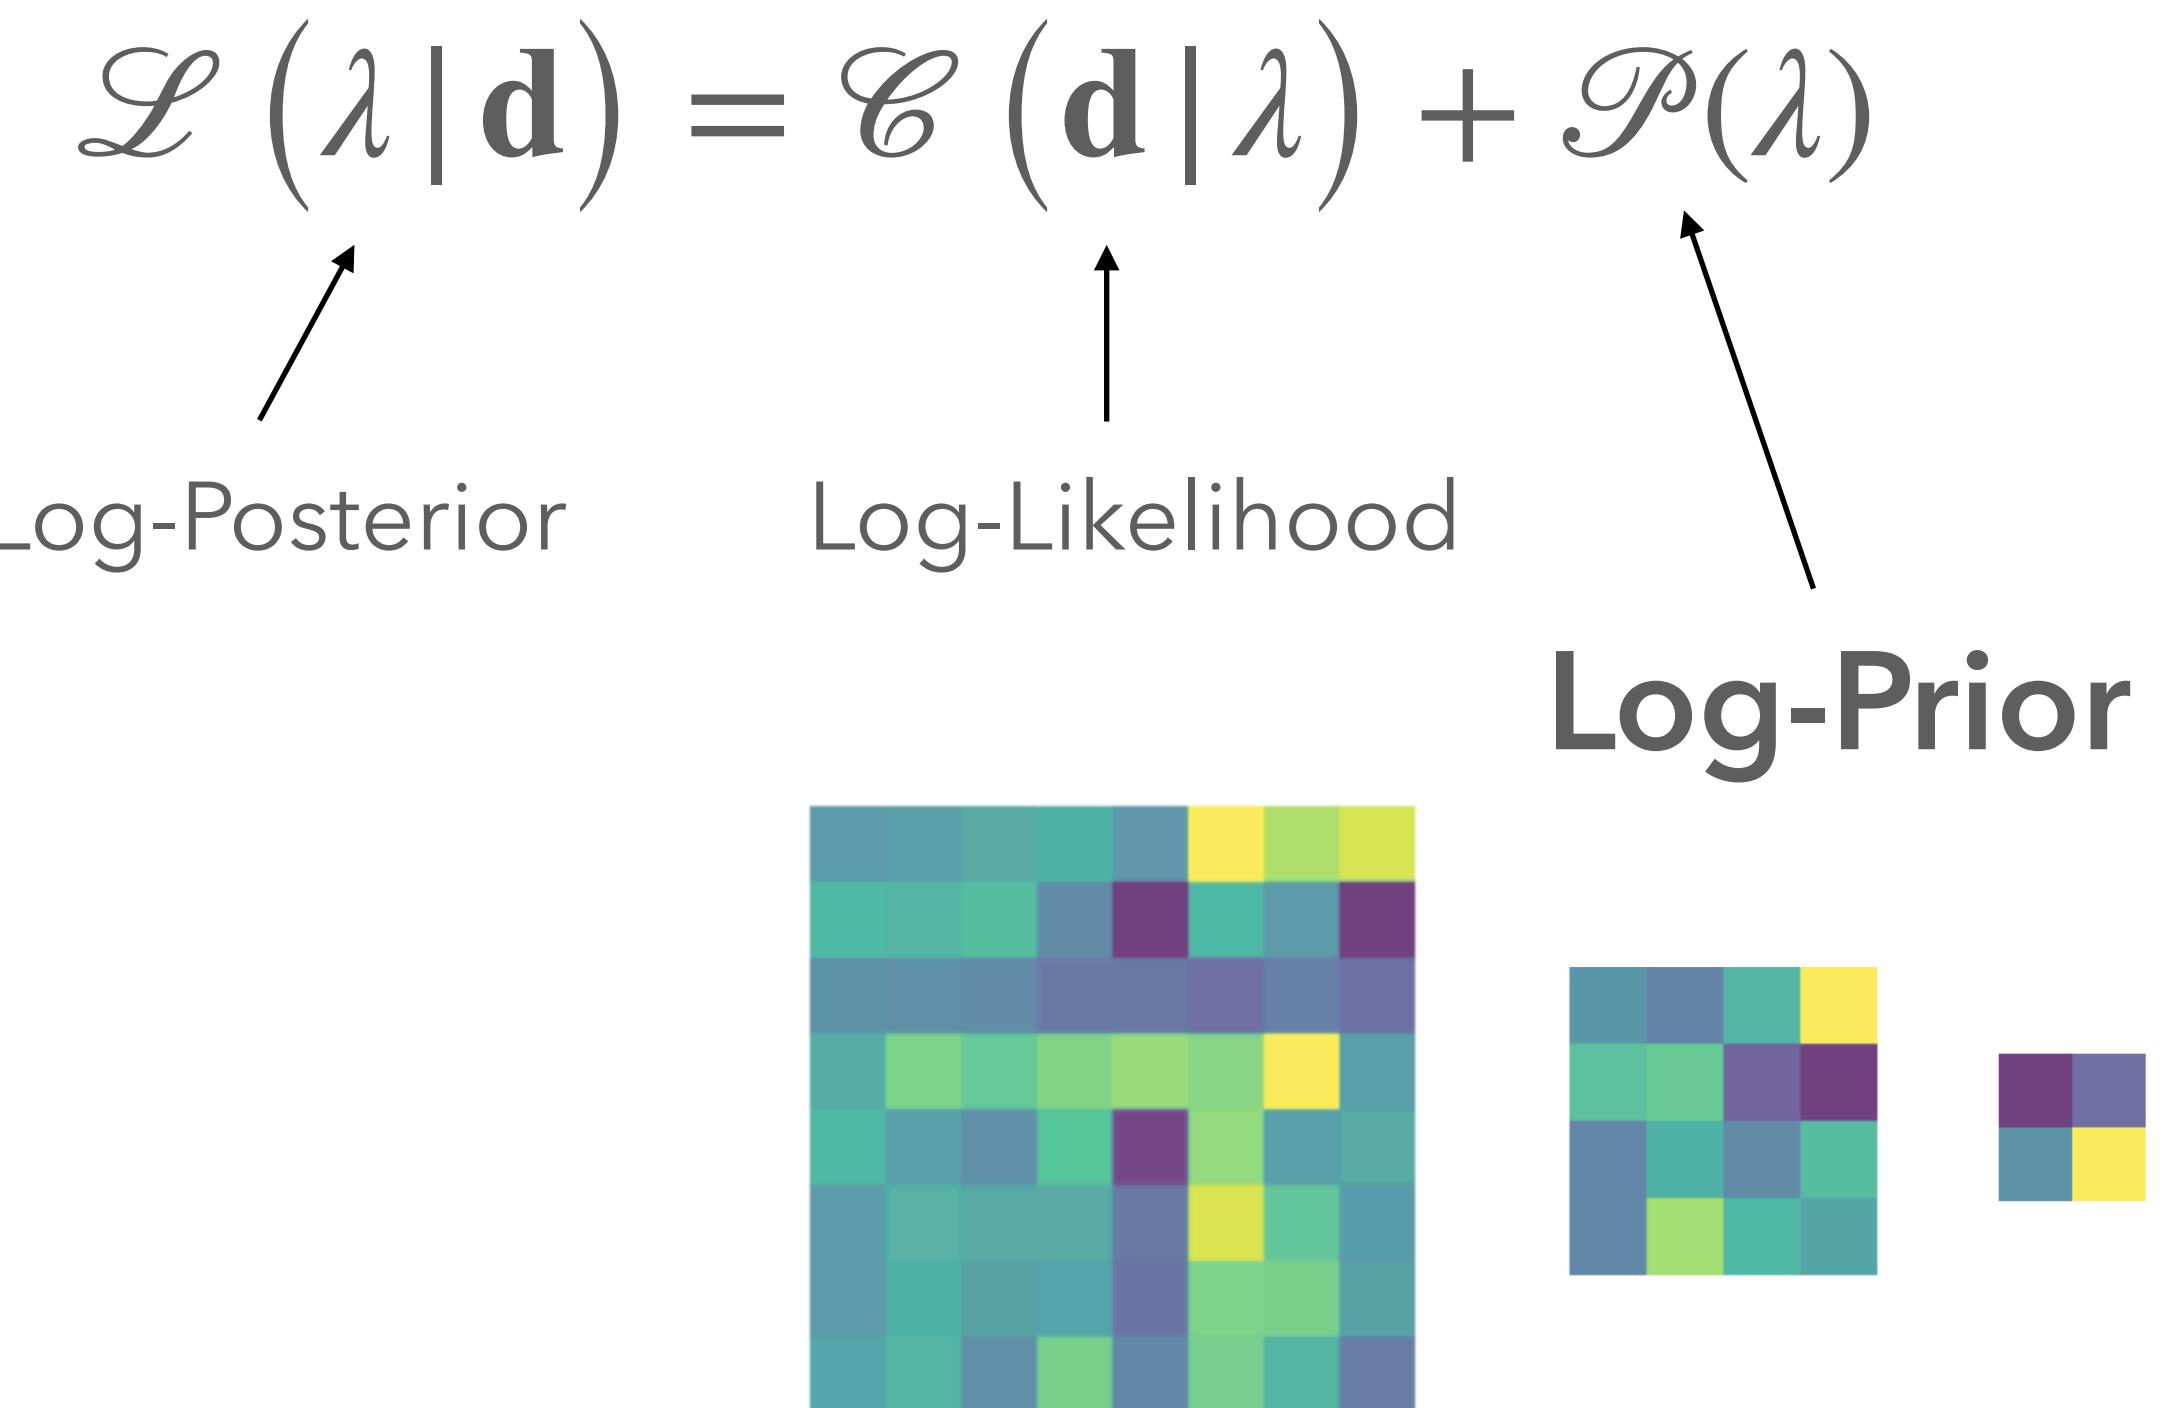
$$\mathcal{L}(\lambda | \mathbf{d}) = \mathcal{C}(\mathbf{d} | \lambda) + \mathcal{P}(\lambda)$$

The equation  $\mathcal{L}(\lambda | \mathbf{d}) = \mathcal{C}(\mathbf{d} | \lambda) + \mathcal{P}(\lambda)$  is displayed. Below the equation, there are two terms: "Log-Posterior" and "Log-Likelihood". A diagonal arrow points from "Log-Posterior" up towards the first term  $\mathcal{C}(\mathbf{d} | \lambda)$ . A vertical arrow points from "Log-Likelihood" up towards the second term  $\mathcal{P}(\lambda)$ .

# LIRA method

(L)ow counts (I)mage (R)econstruction and (A)analysis

Objective function **extended by a log-prior term**, only depending on  $\lambda$  :



A special “multi-scale” prior to achieve **smoothness on multiple scales**.

Initial idea and implementation by [[Esch et al. 2004](#)].

More detailed explanation follows...

# LIRA method

(L)ow counts (I)mage (R)econstruction and (A)analysis

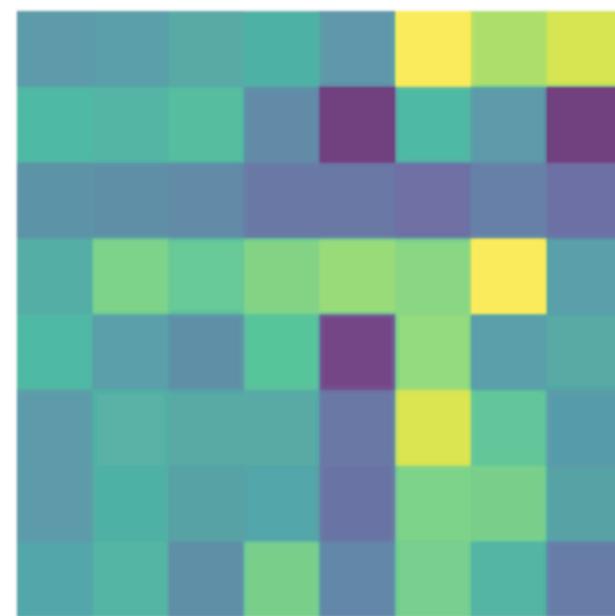
Objective function **extended by a log-prior term**, only depending on  $\lambda$  :

$$\mathcal{L}(\lambda | \mathbf{d}) = \mathcal{C}(\mathbf{d} | \lambda) + \mathcal{P}(\lambda)$$

Log-Posterior

Log-Likelihood

Log-Prior



Model counts extend by **exposure e** and optional **baseline ("background") component b** by [[Connors et al. 2011](#)]

$$\lambda = [(\mathbf{x} + \mathbf{b}) \cdot \mathbf{e}] \circledast \text{PSF}$$

A special "multi-scale" prior to achieve **smoothness on multiple scales**.

Initial idea and implementation by [[Esch et al. 2004](#)].

More detailed explanation follows...

# LIRA method

(L)ow counts (I)mage (R)econstruction and (A)analysis

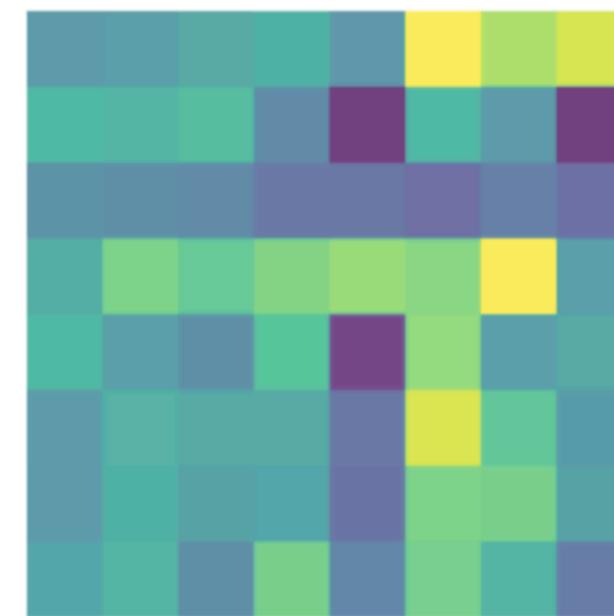
Objective function **extended by a log-prior term**, only depending on  $\lambda$  :

$$\mathcal{L}(\lambda | \mathbf{d}) = \mathcal{C}(\mathbf{d} | \lambda) + \mathcal{P}(\lambda)$$

Log-Posterior

Log-Likelihood

Log-Prior



Model counts extend by **exposure e** and optional **baseline ("background") component b** by [Connors et al. 2011]

$$\lambda = [(\mathbf{x} + \mathbf{b}) \cdot \mathbf{e}] \circledast \text{PSF}$$

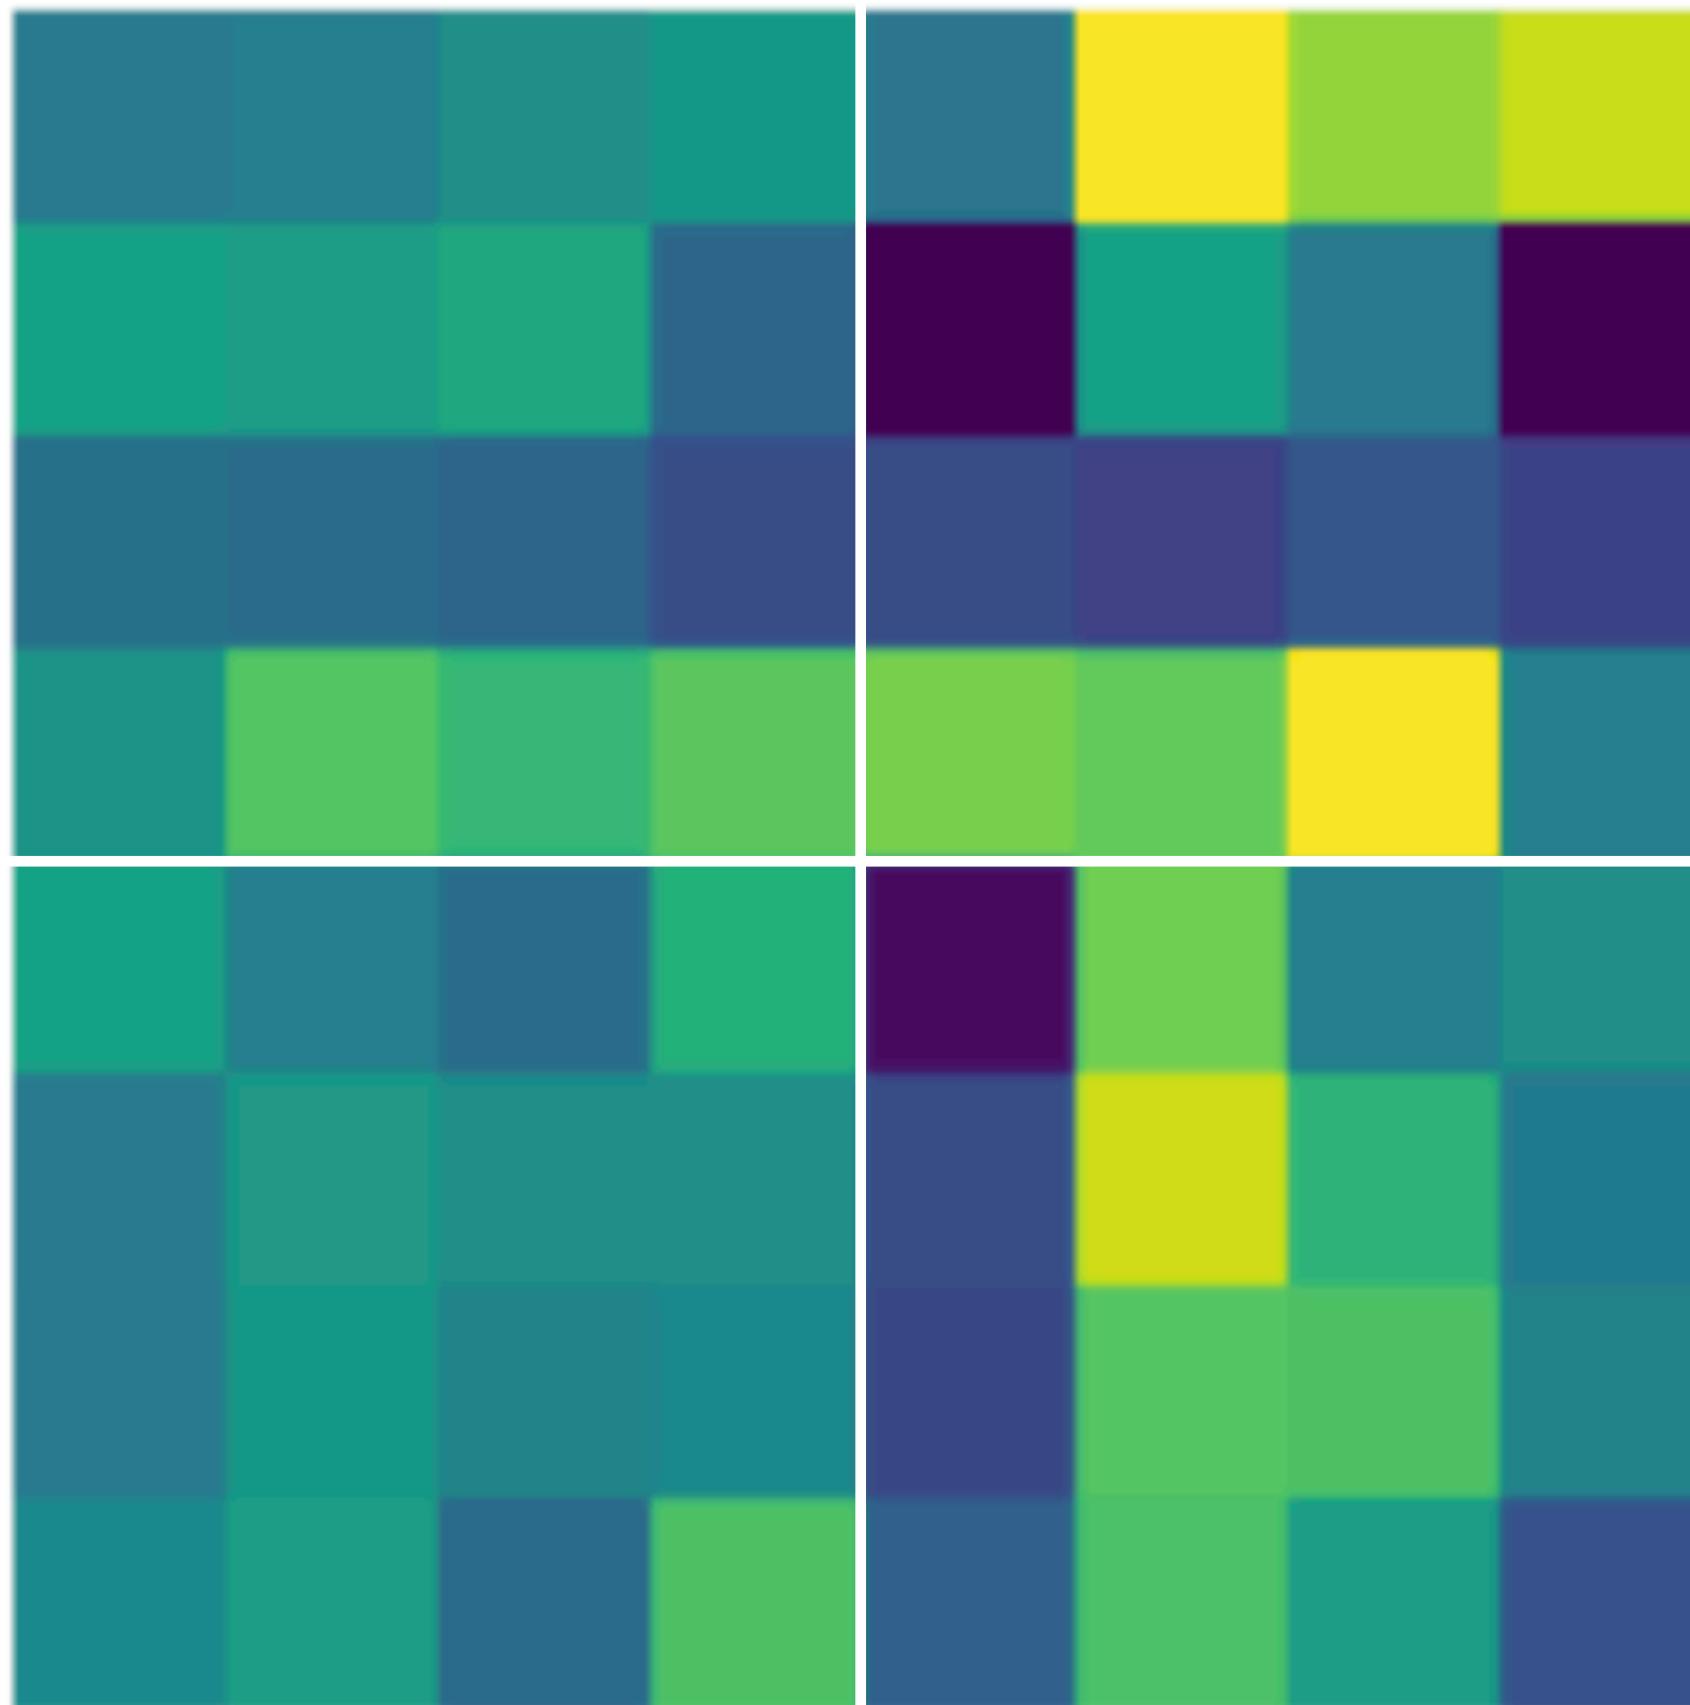
Change from EM to **MCMC sampling** from the log-posterior. Effectively **sampling a series of images...**

A special "multi-scale" prior to achieve **smoothness on multiple scales**.

Initial idea and implementation by [Esch et al. 2004].

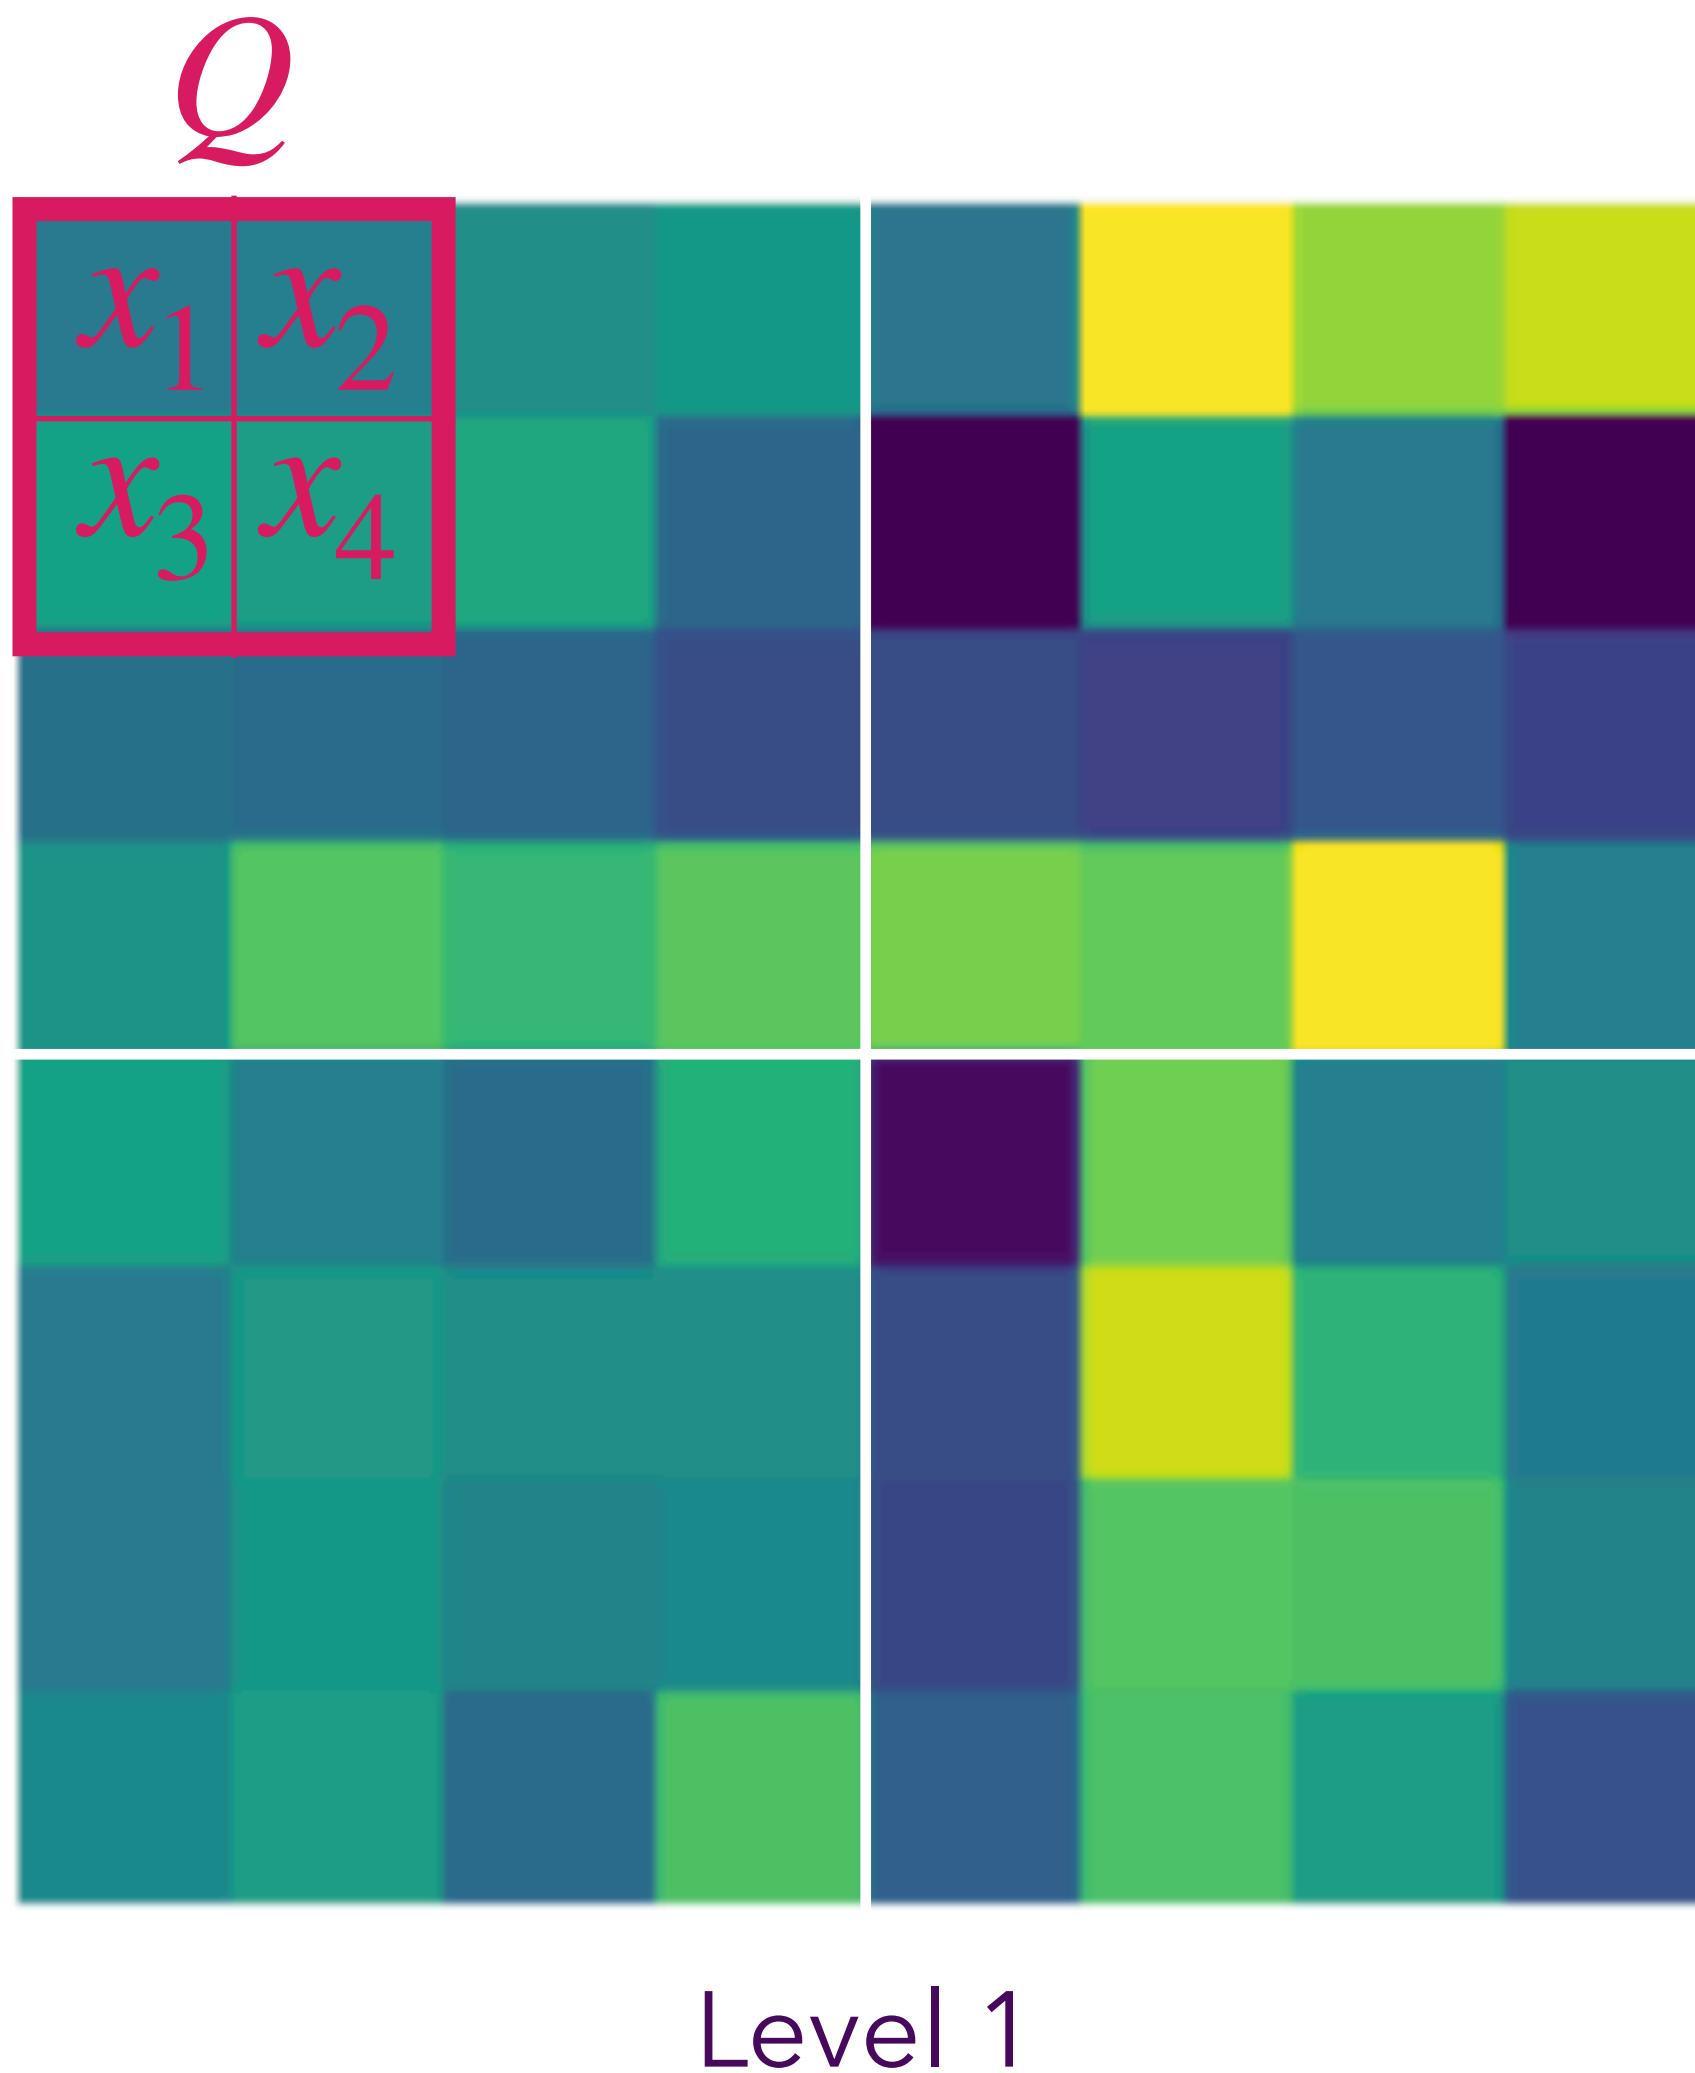
More detailed explanation follows...

# Multiscale decomposition

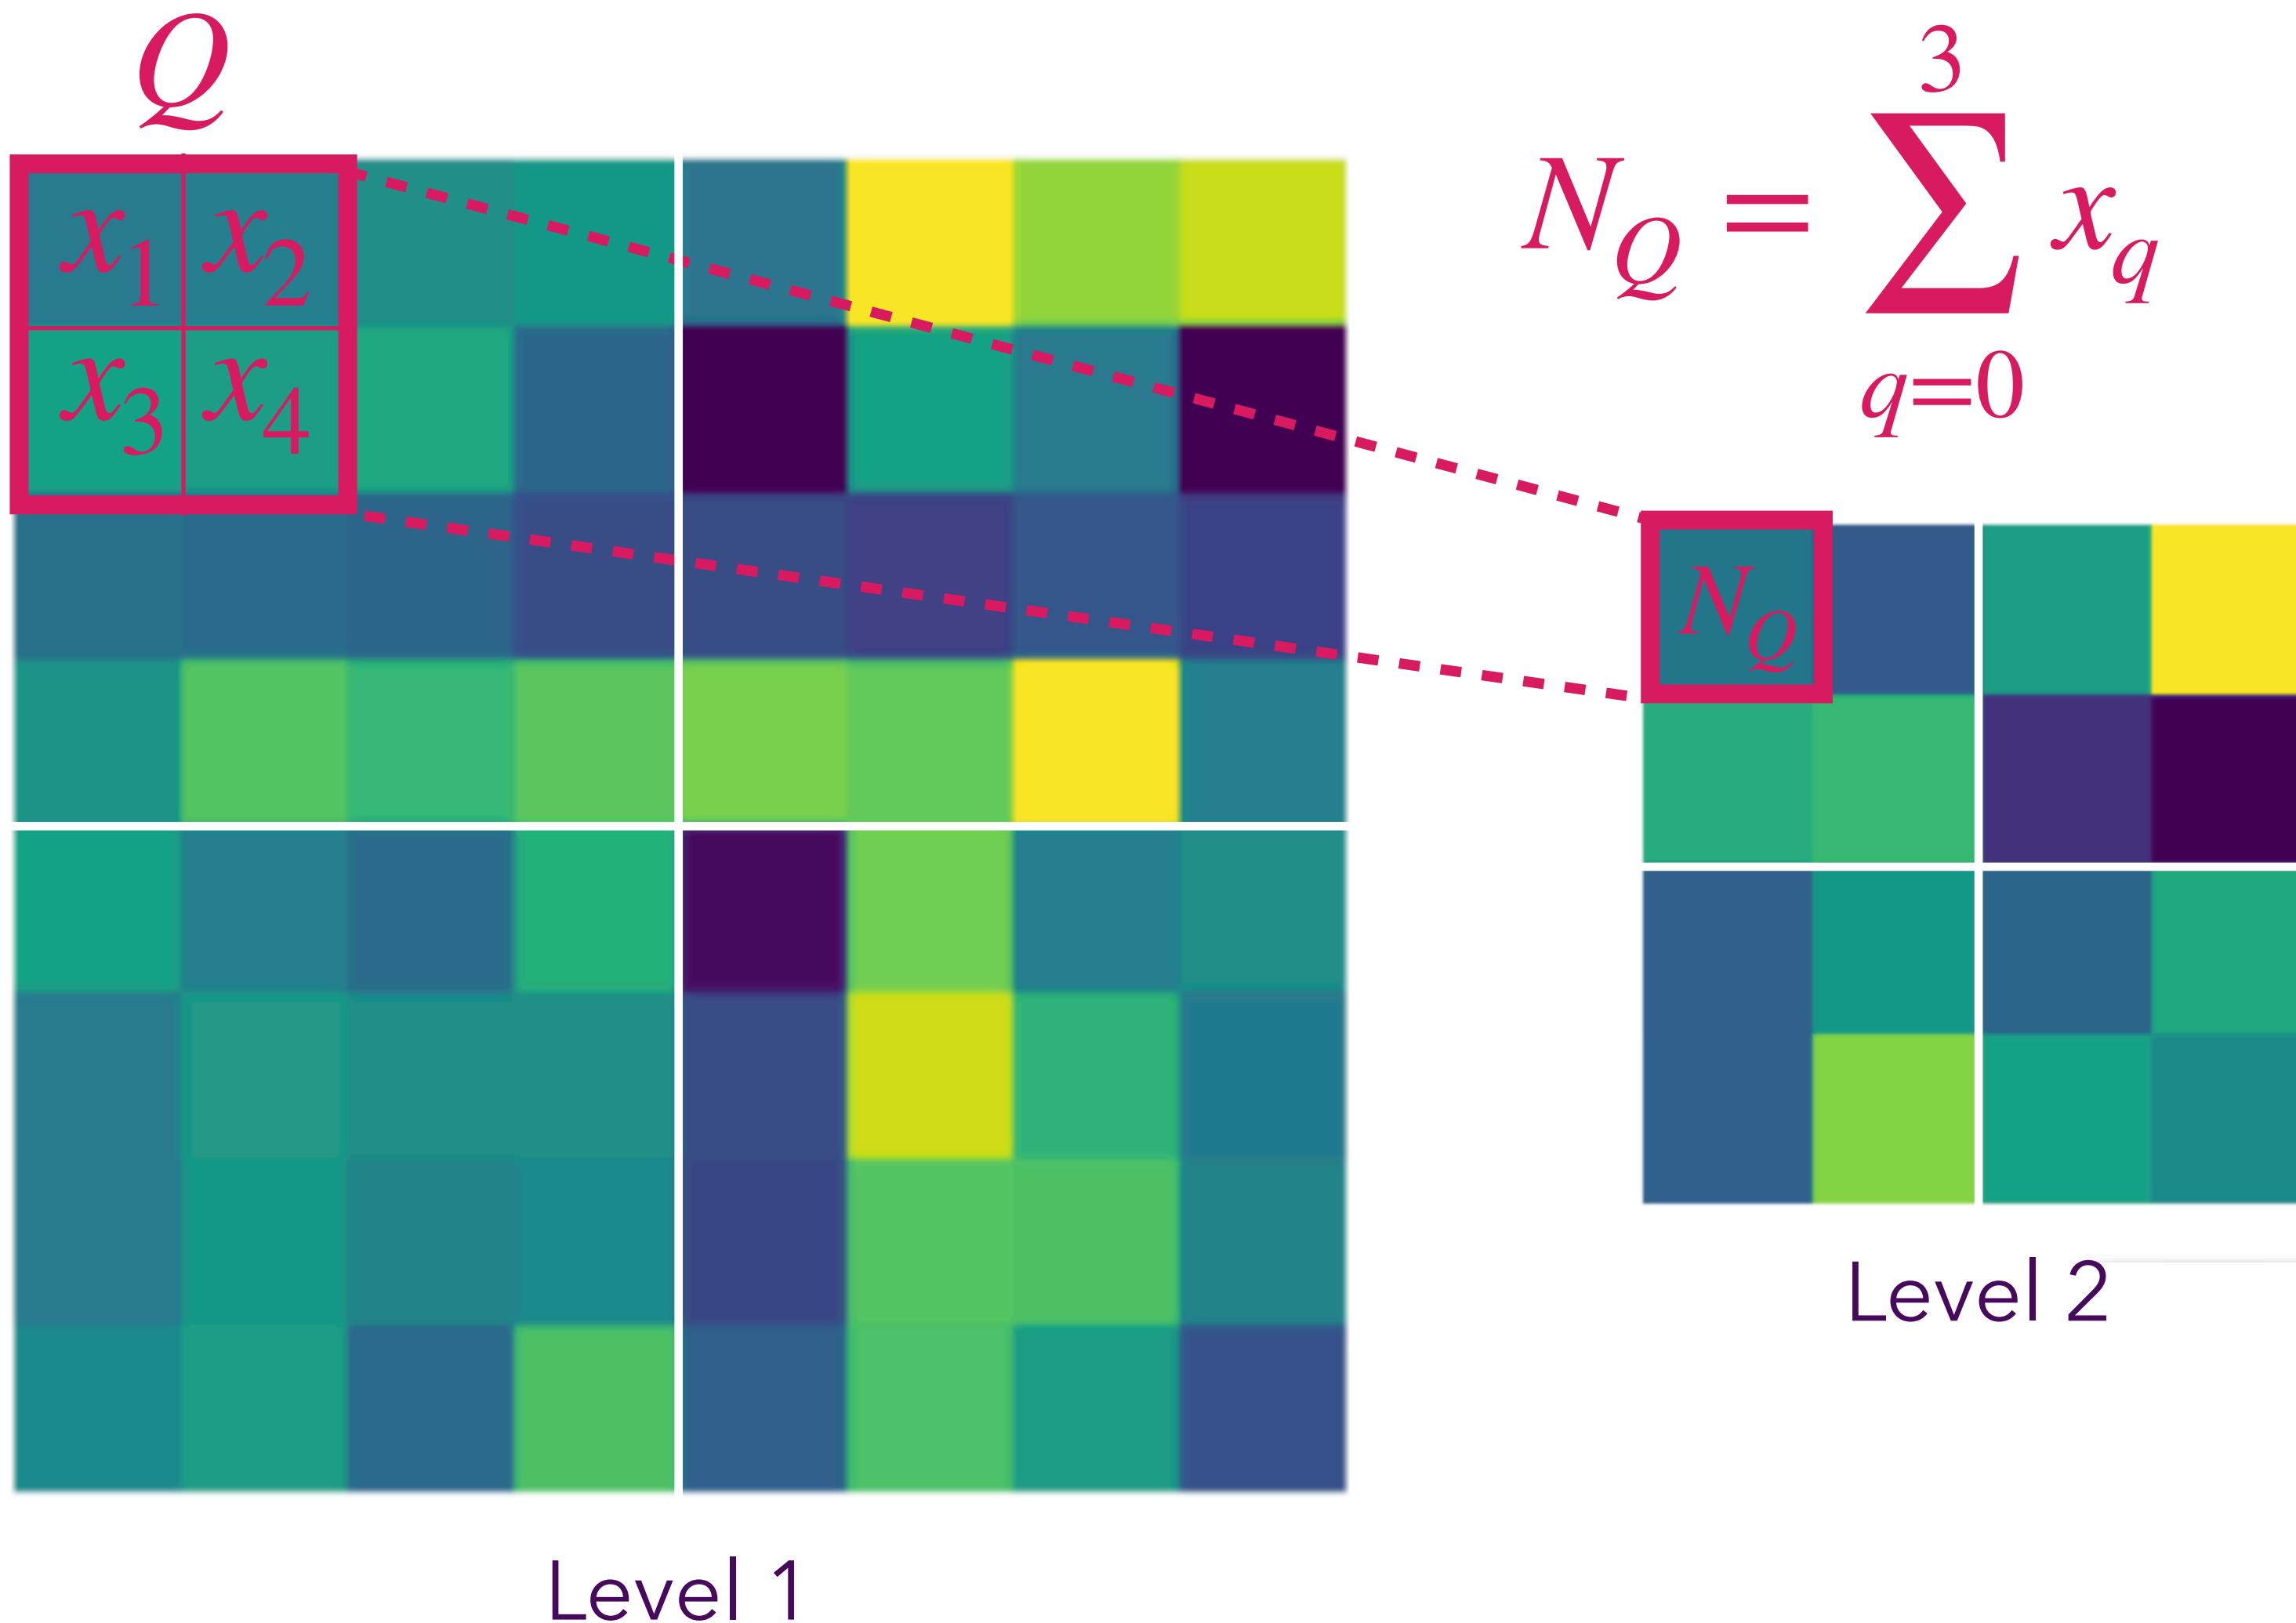


Level 1

# Multiscale decomposition

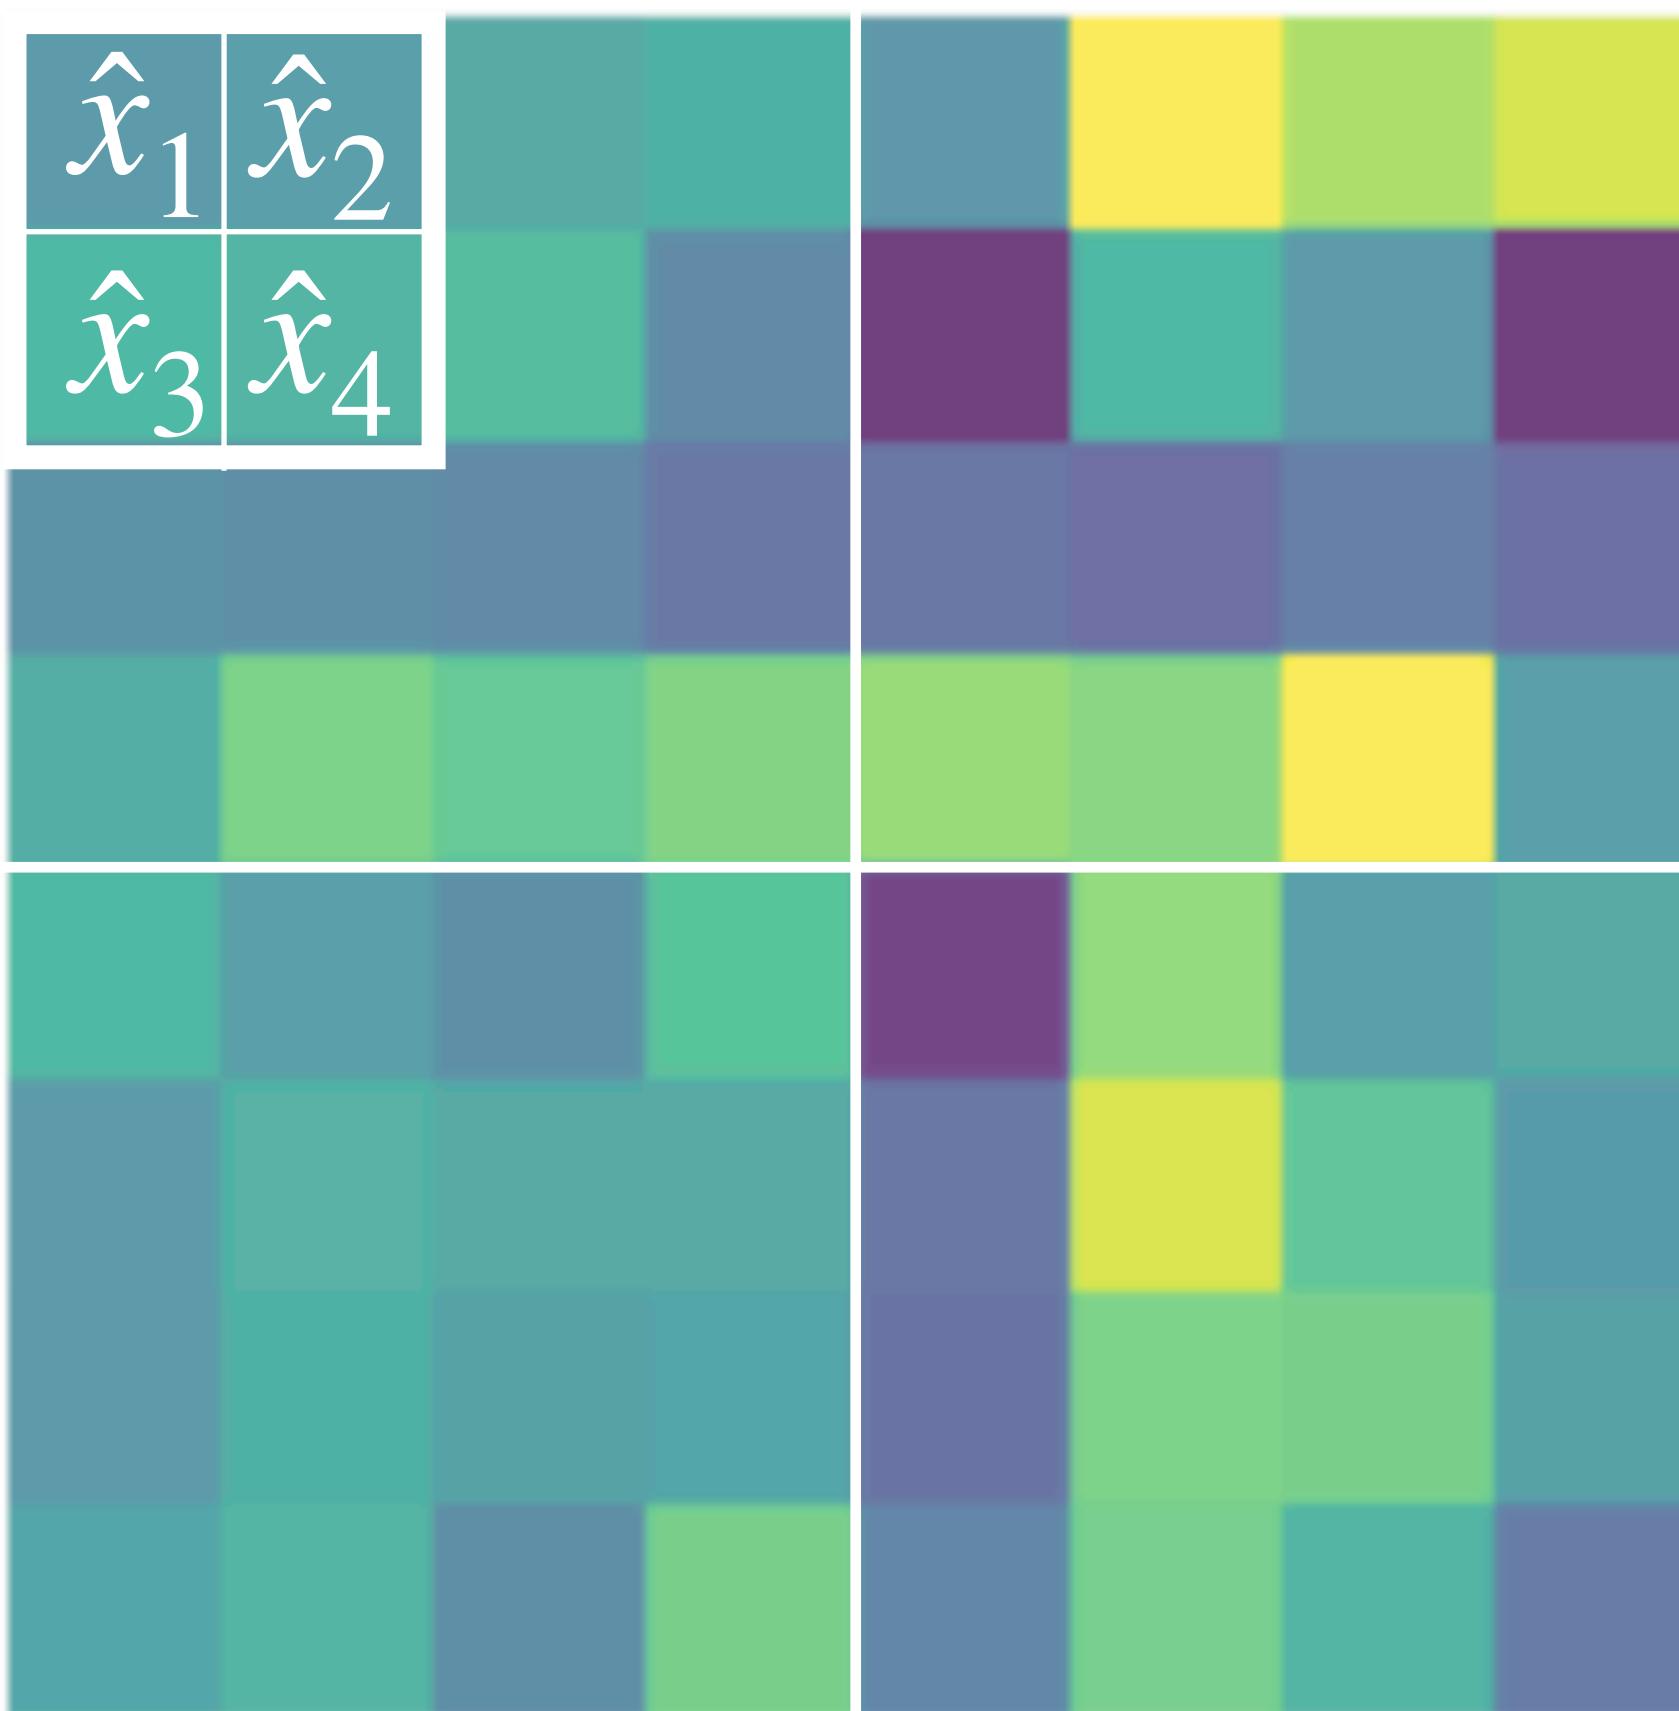


# Multiscale decomposition



# Multiscale decomposition

$$\hat{x}_q = x_q / N_Q$$

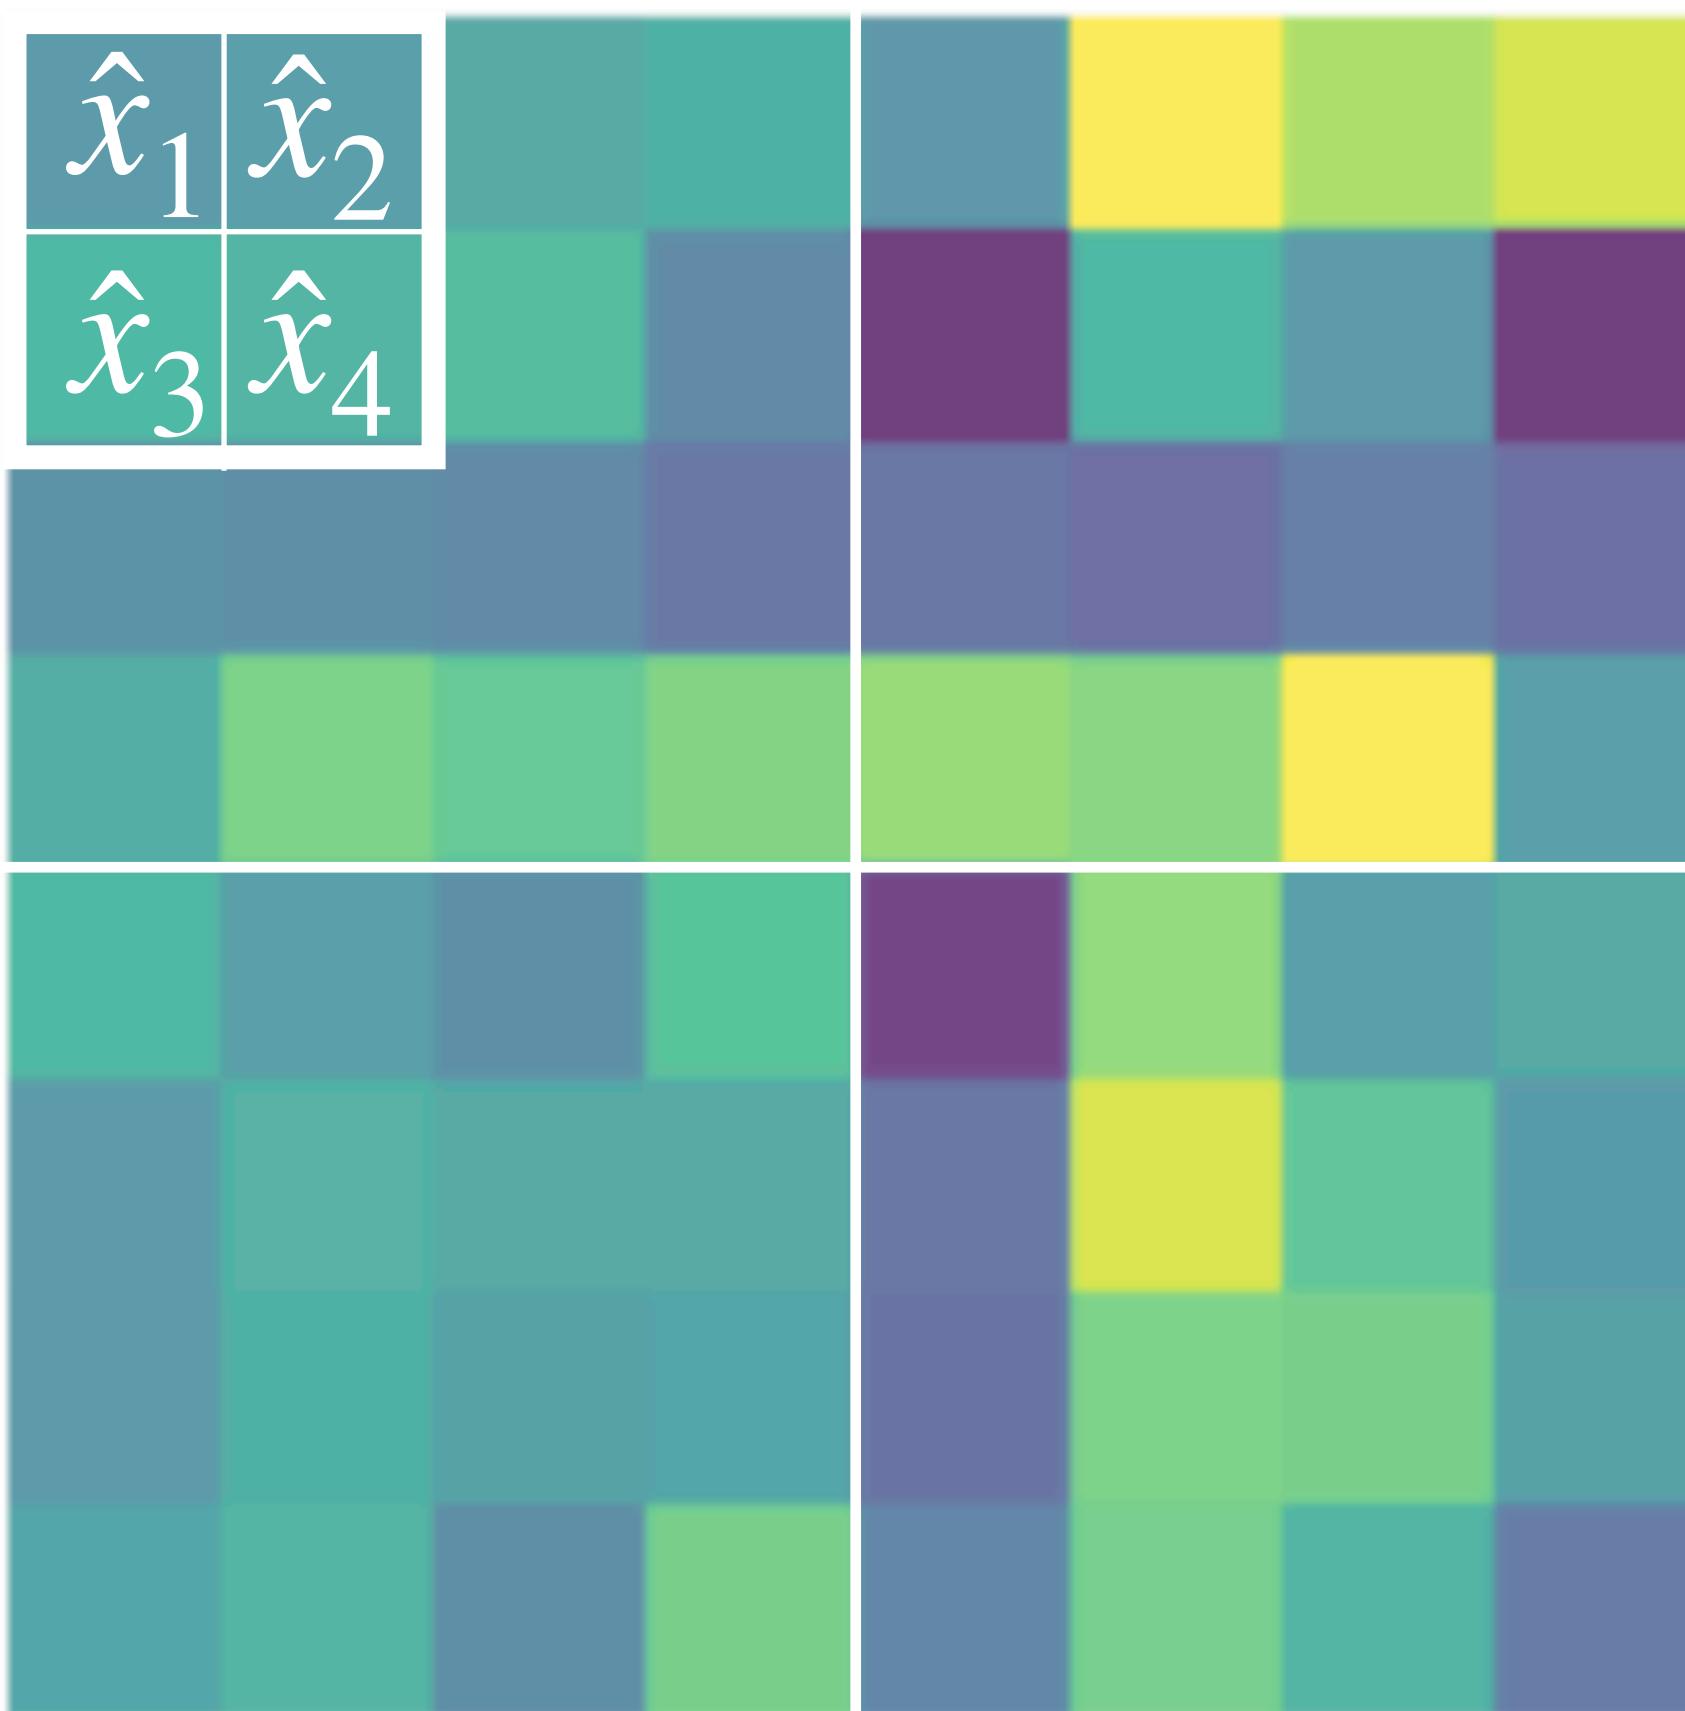


$$N_Q = \sum_{q=0}^3 x_q$$



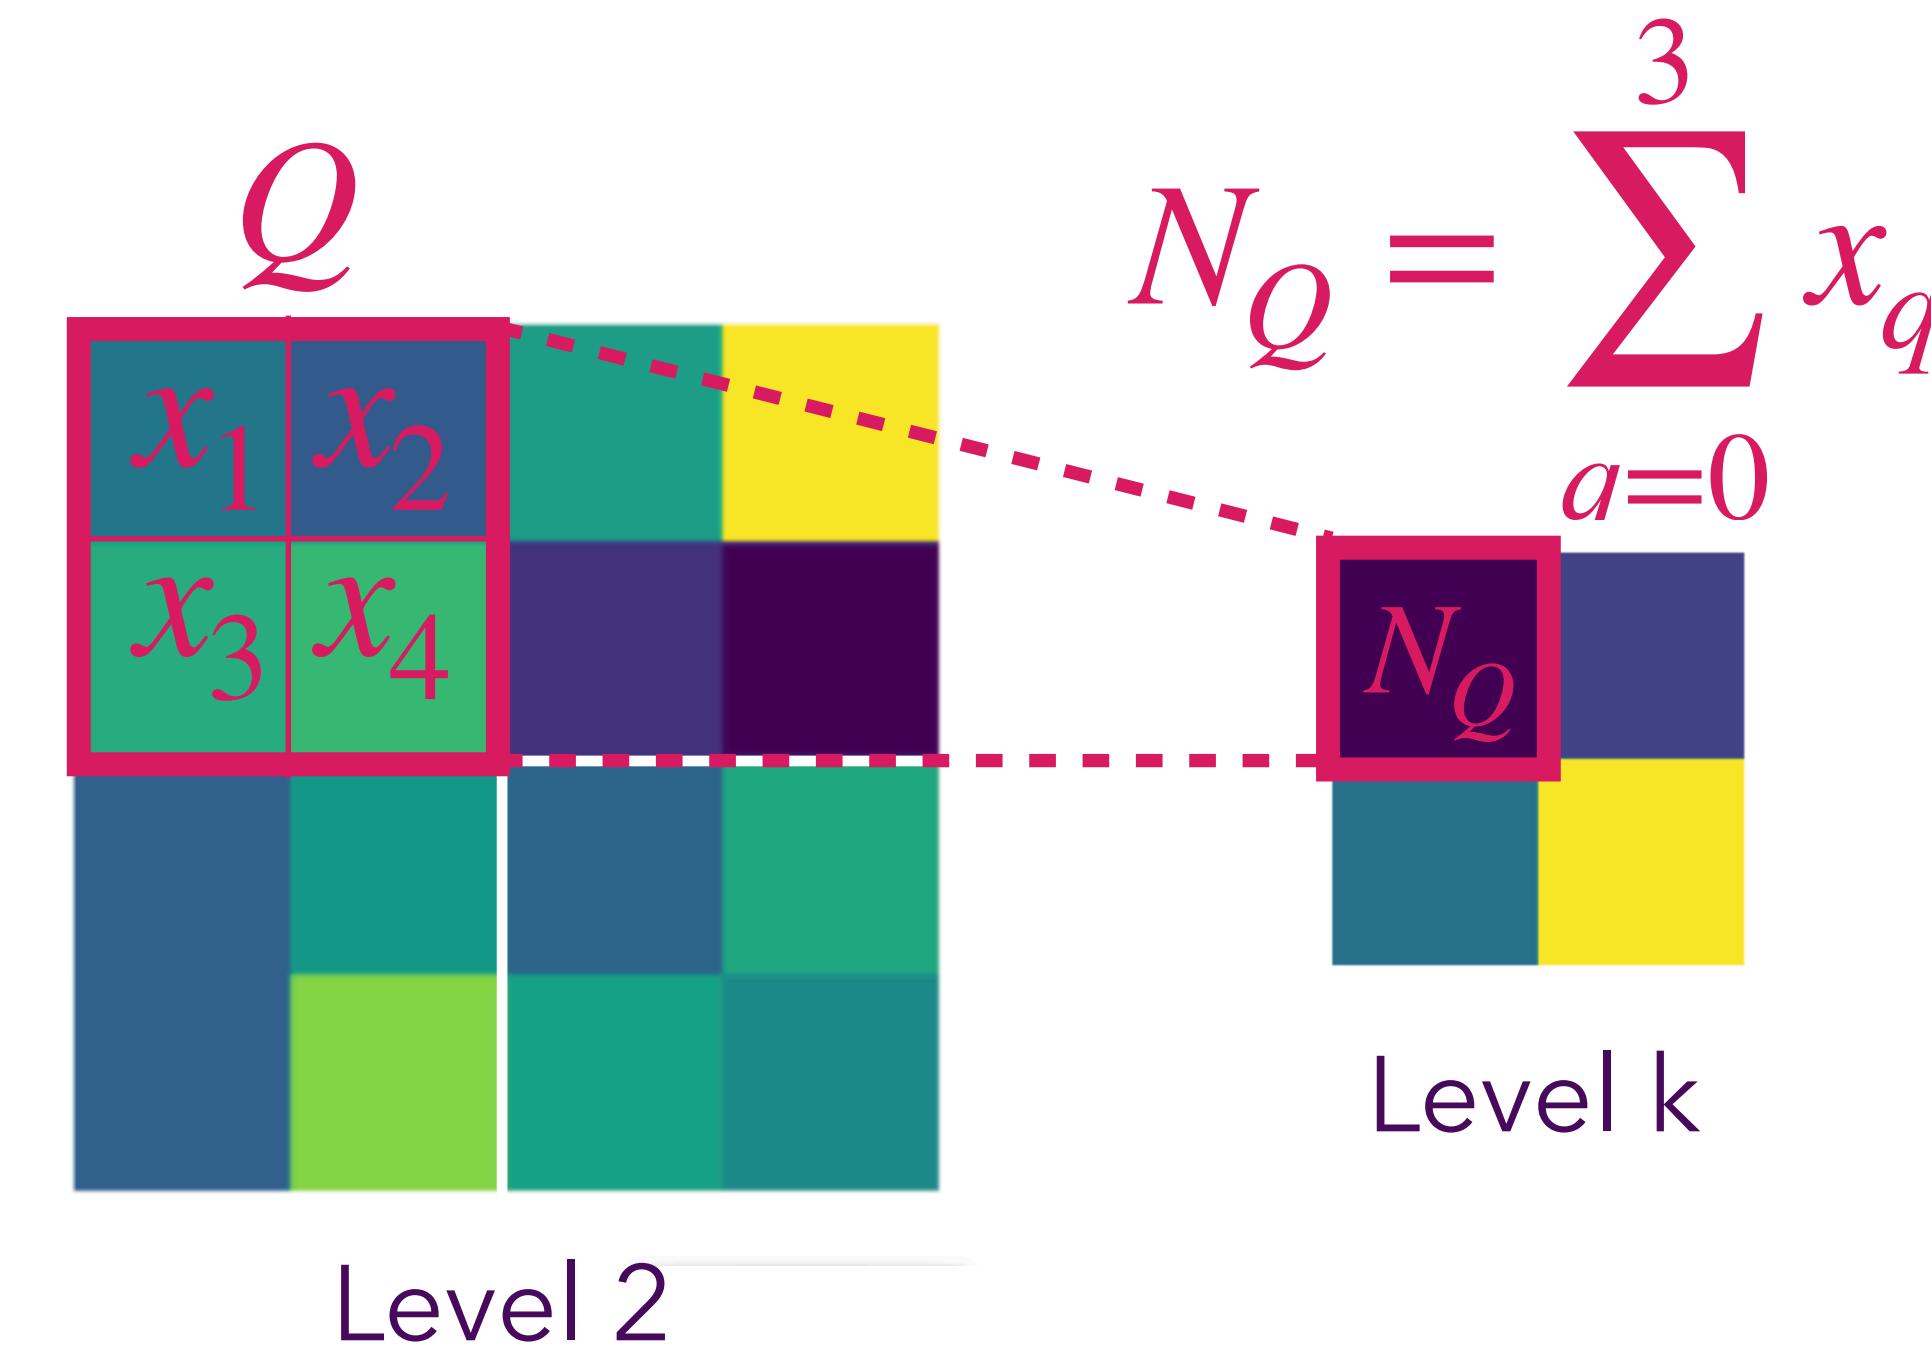
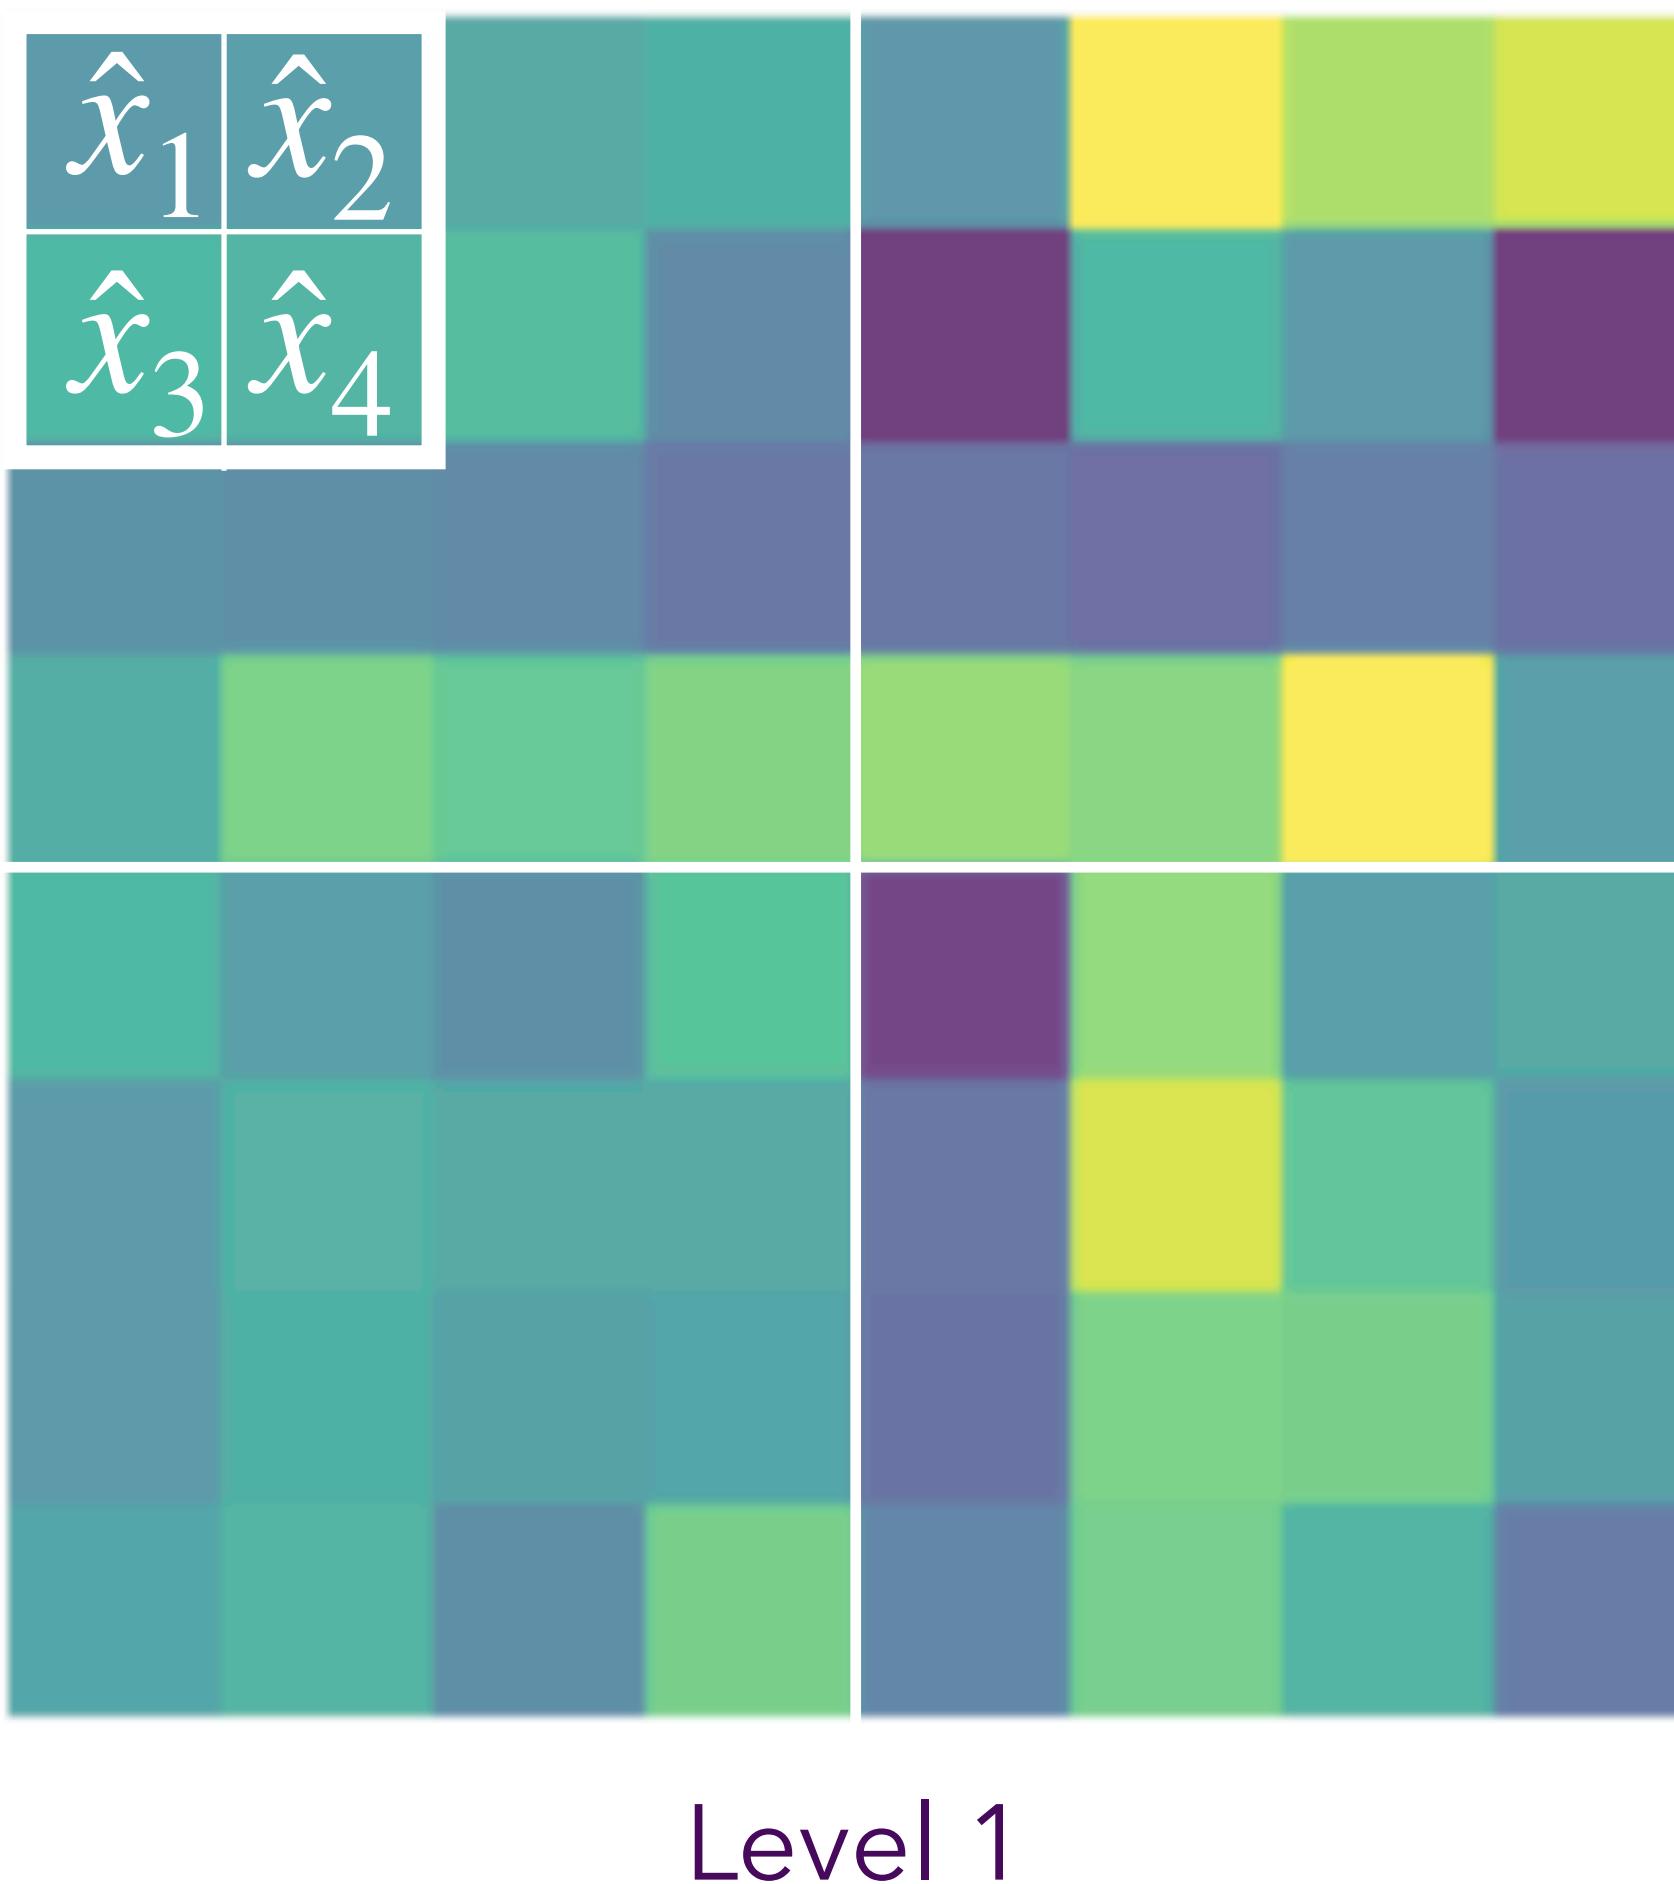
# Multiscale decomposition

$$\hat{x}_q = x_q/N_Q$$



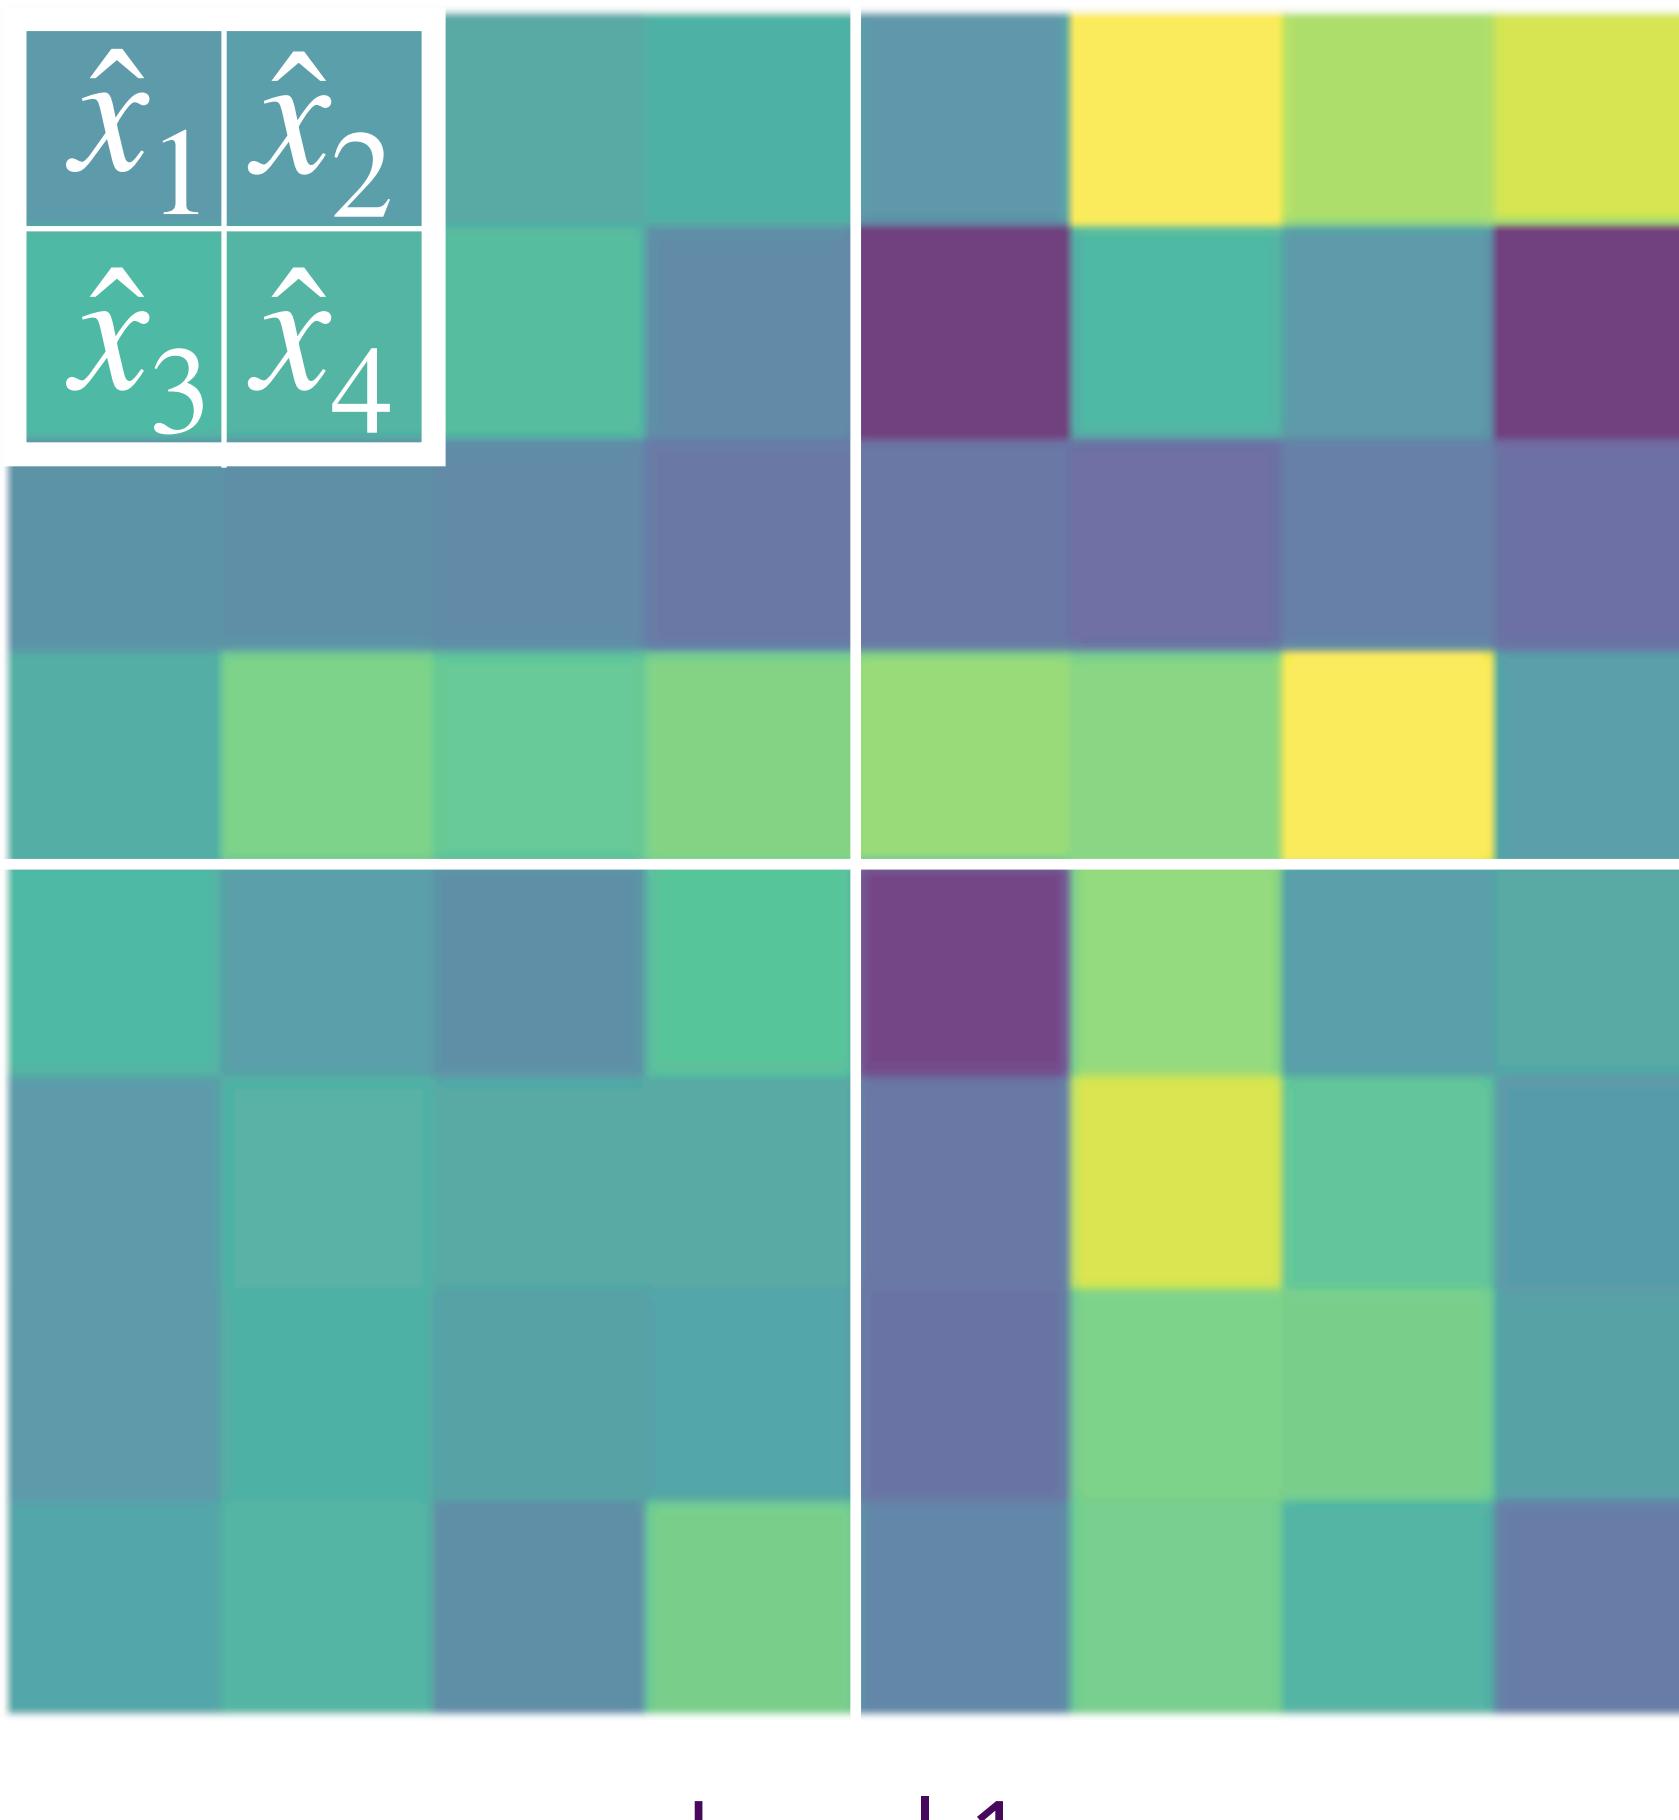
# Multiscale decomposition

$$\hat{x}_q = x_q/N_Q$$



# Multiscale decomposition

$$\hat{x}_q = x_q/N_Q$$



$$\hat{x}_q = x_q/N_Q$$



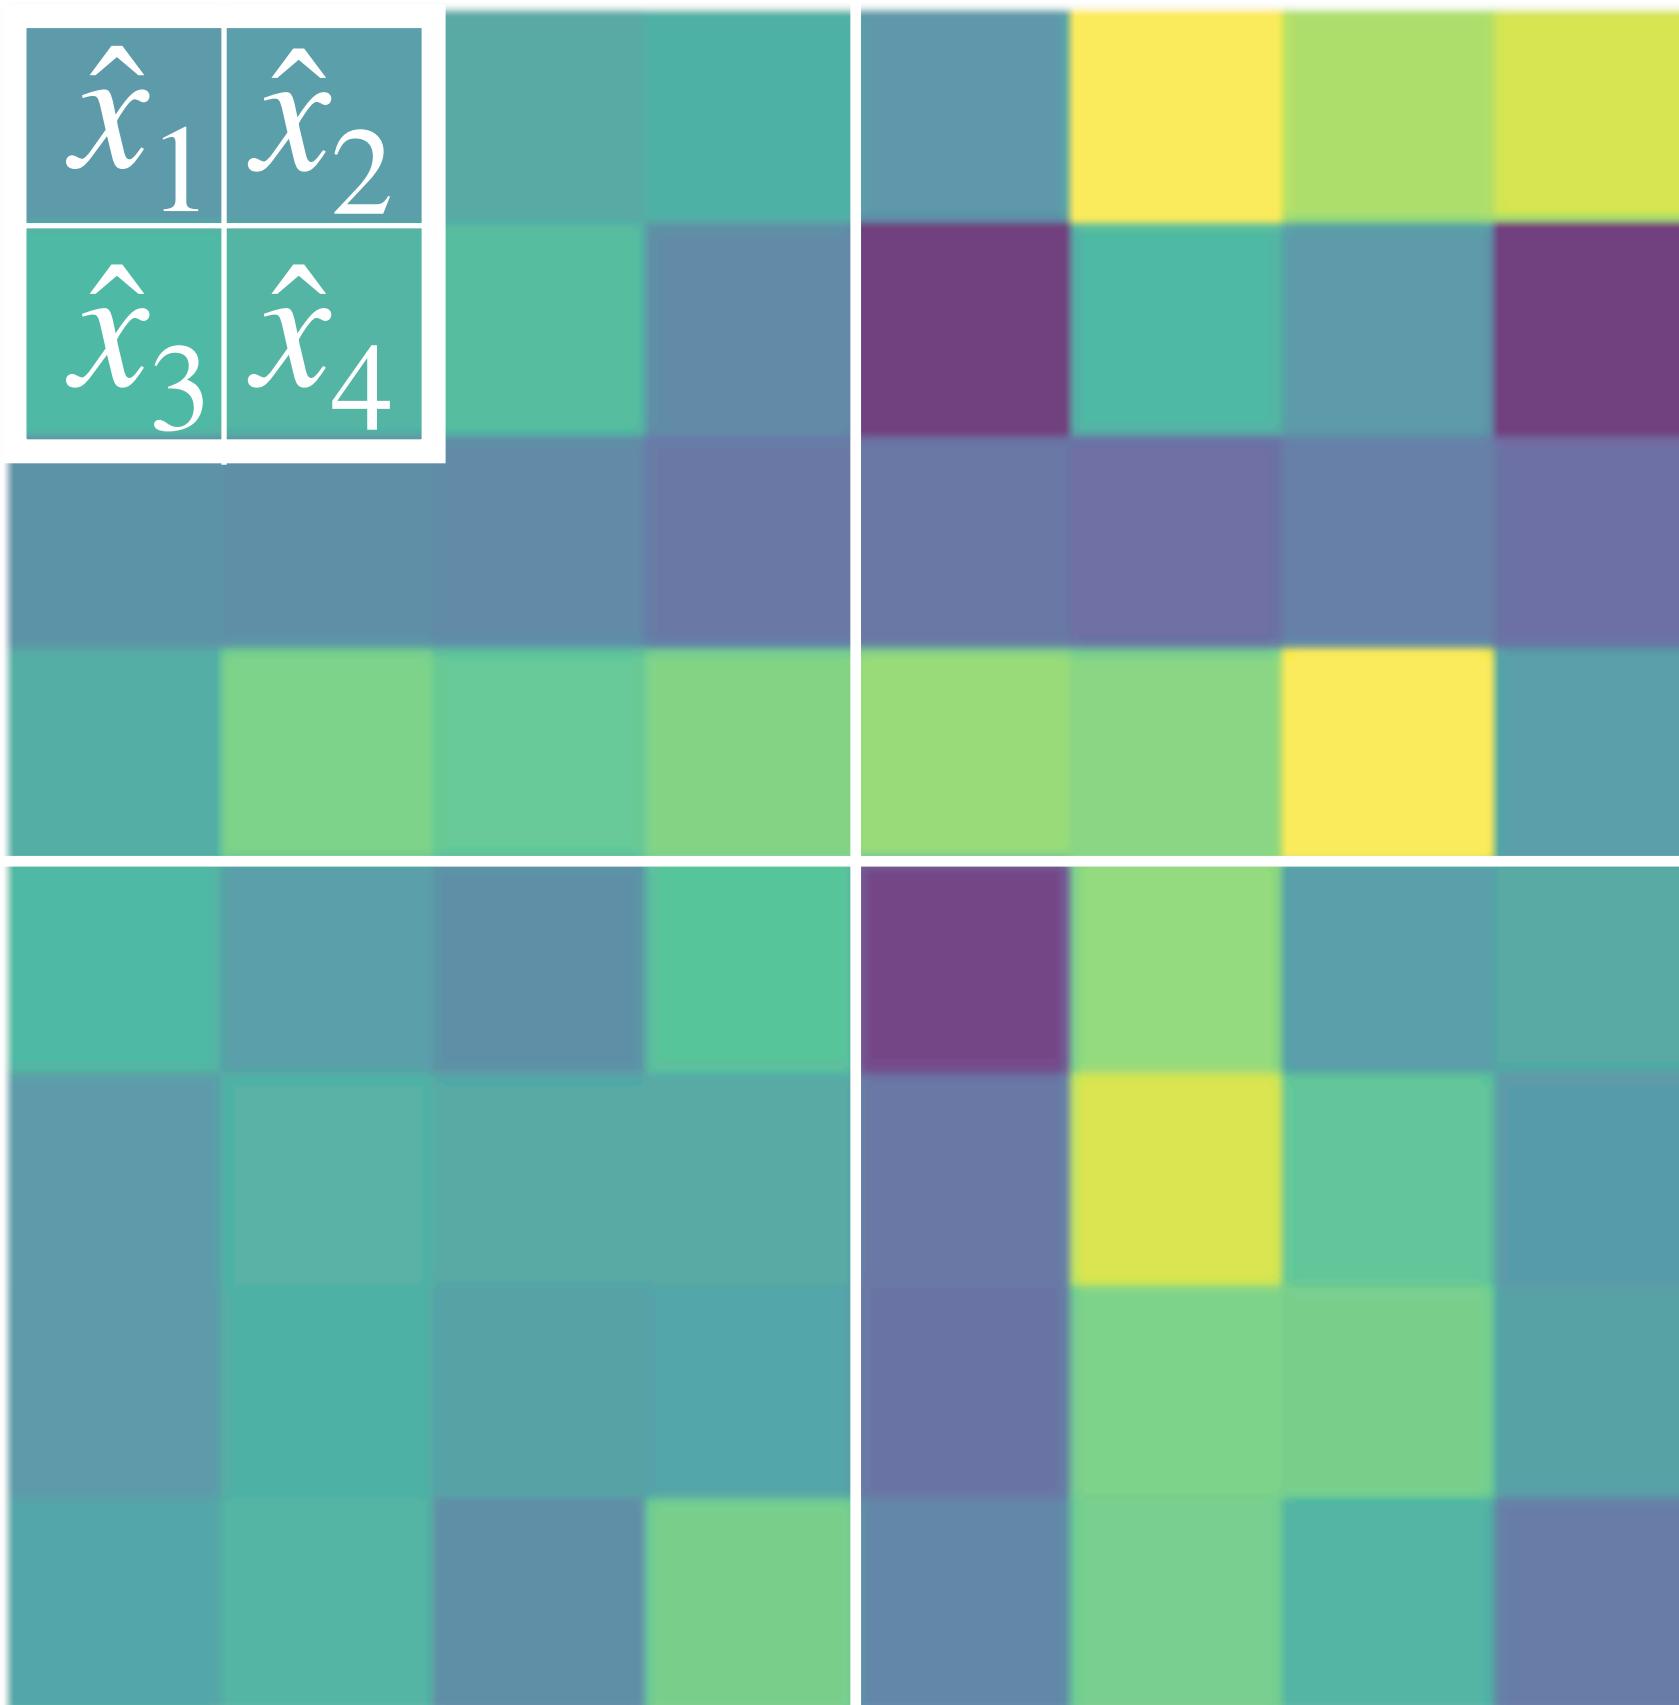
$$N_Q = \sum_{q=0}^3 x_q$$



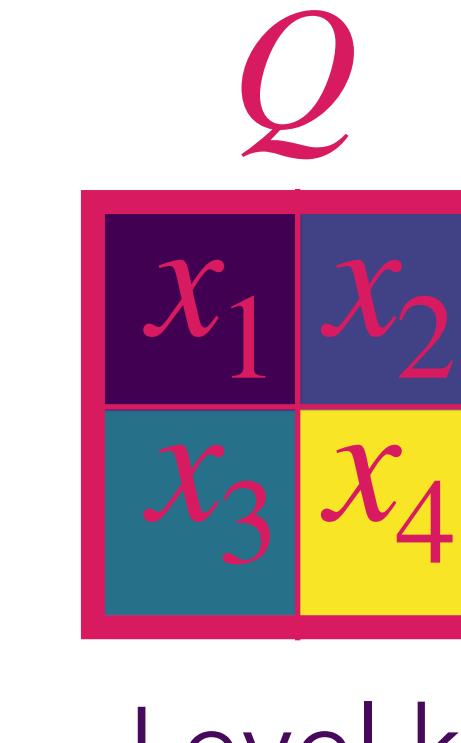
Level k

# Multiscale decomposition

$$\hat{x}_q = x_q/N_Q$$

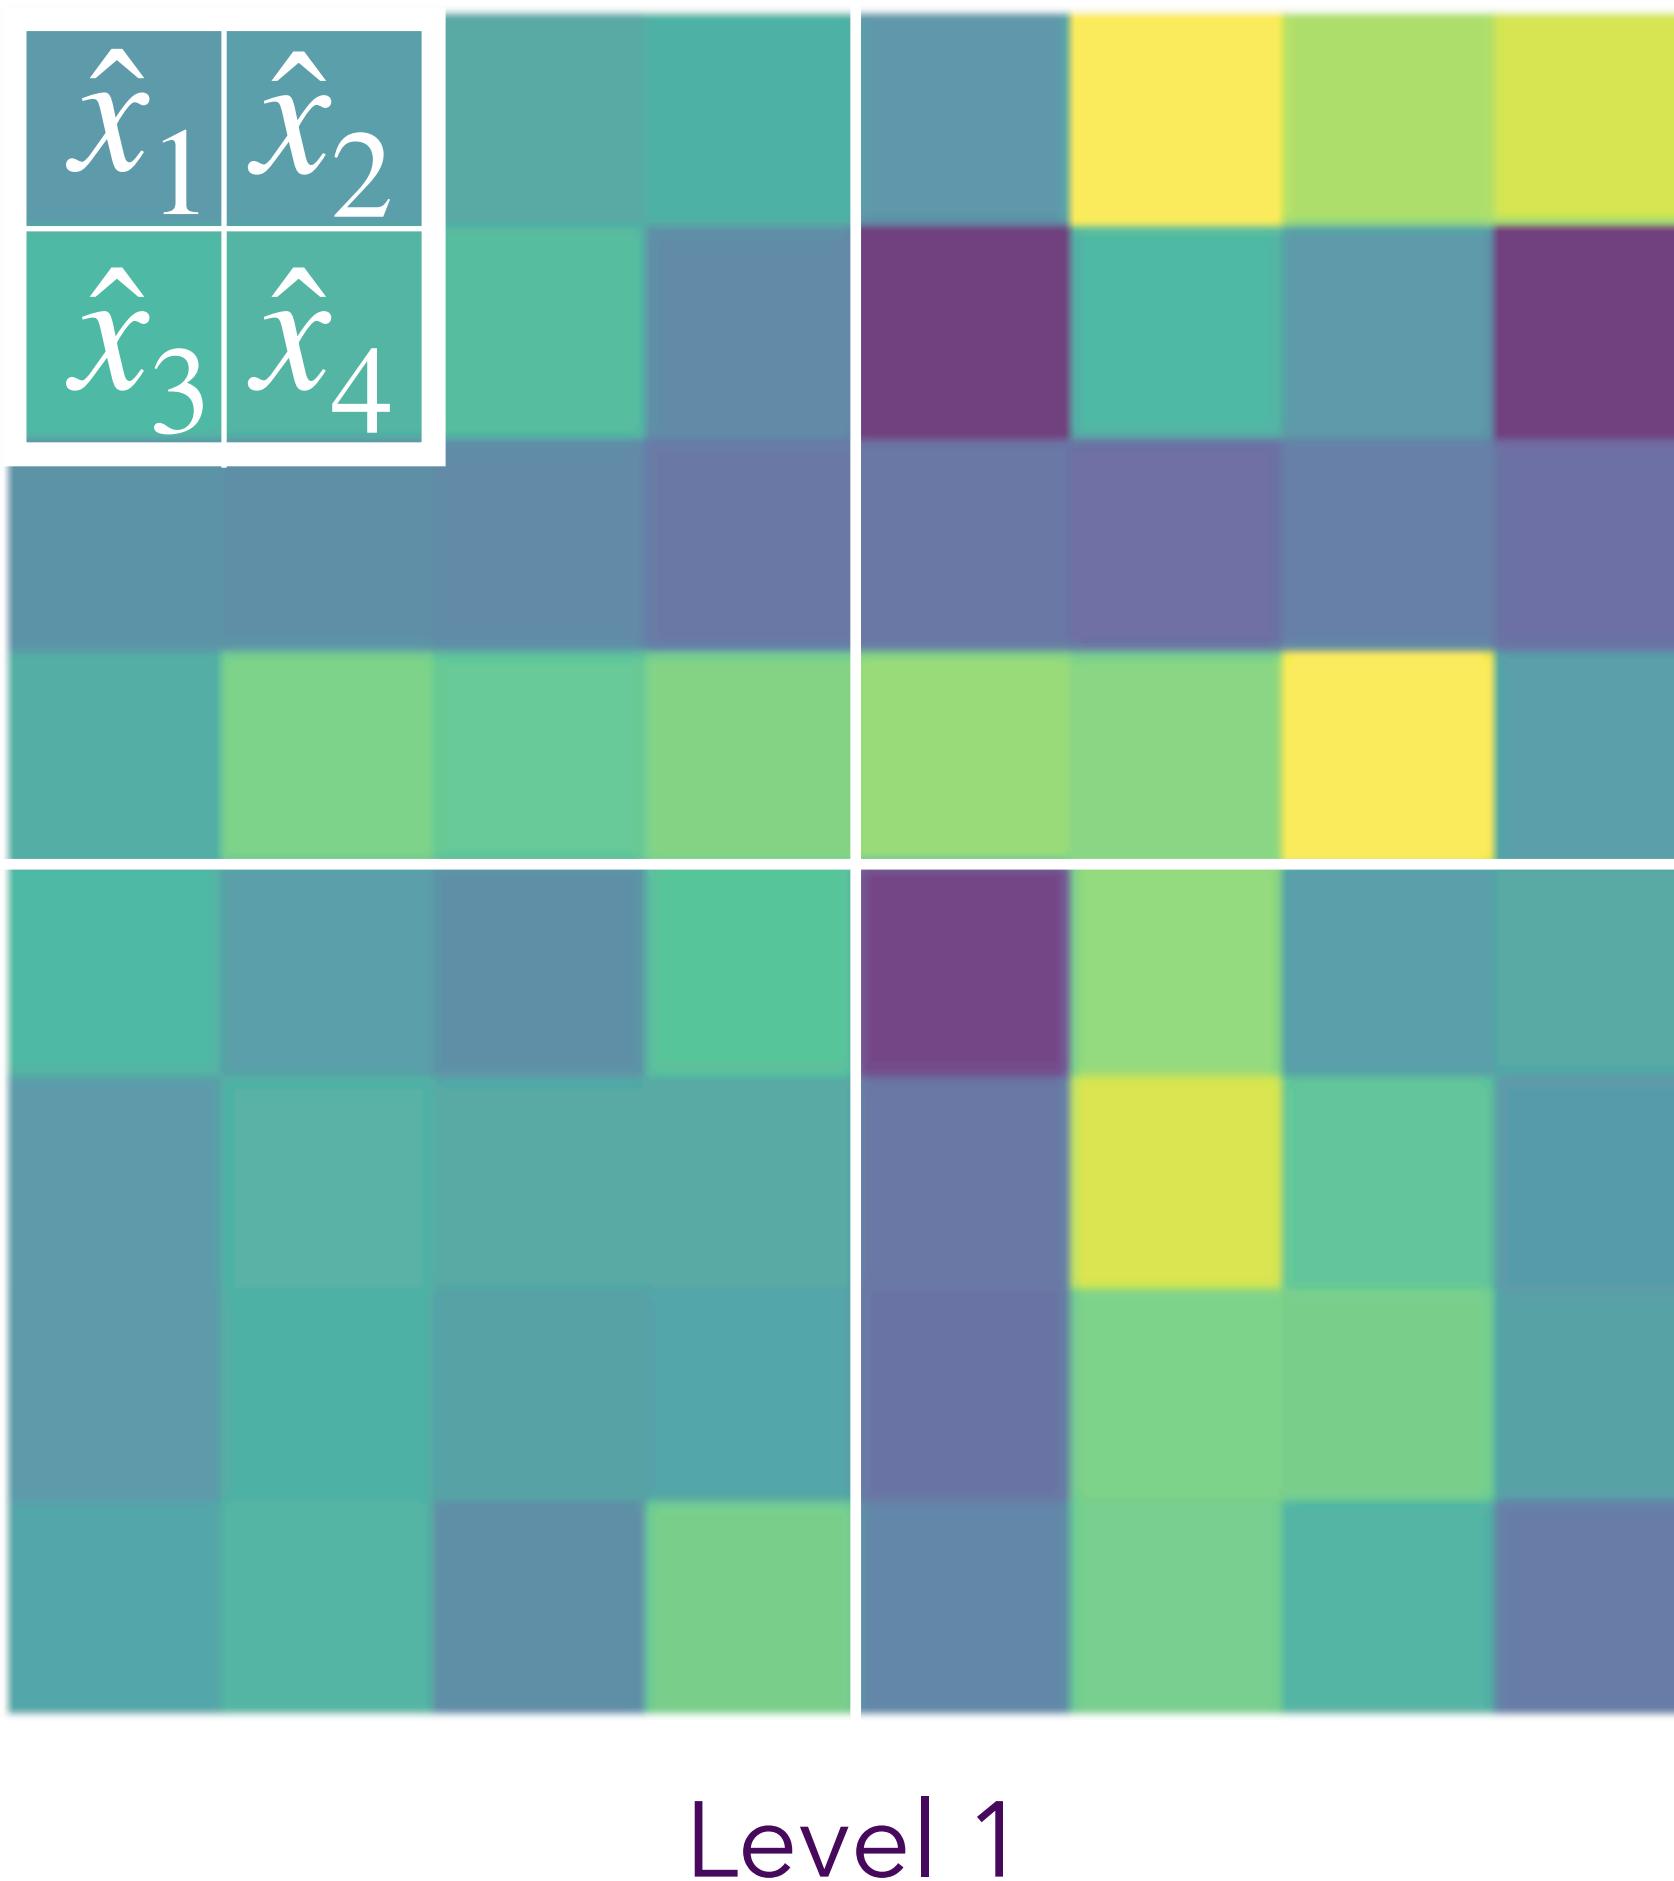


$$\hat{x}_q = x_q/N_Q$$



# Multiscale decomposition

$$\hat{x}_q = x_q/N_Q$$



$$\hat{x}_q = x_q/N_Q$$

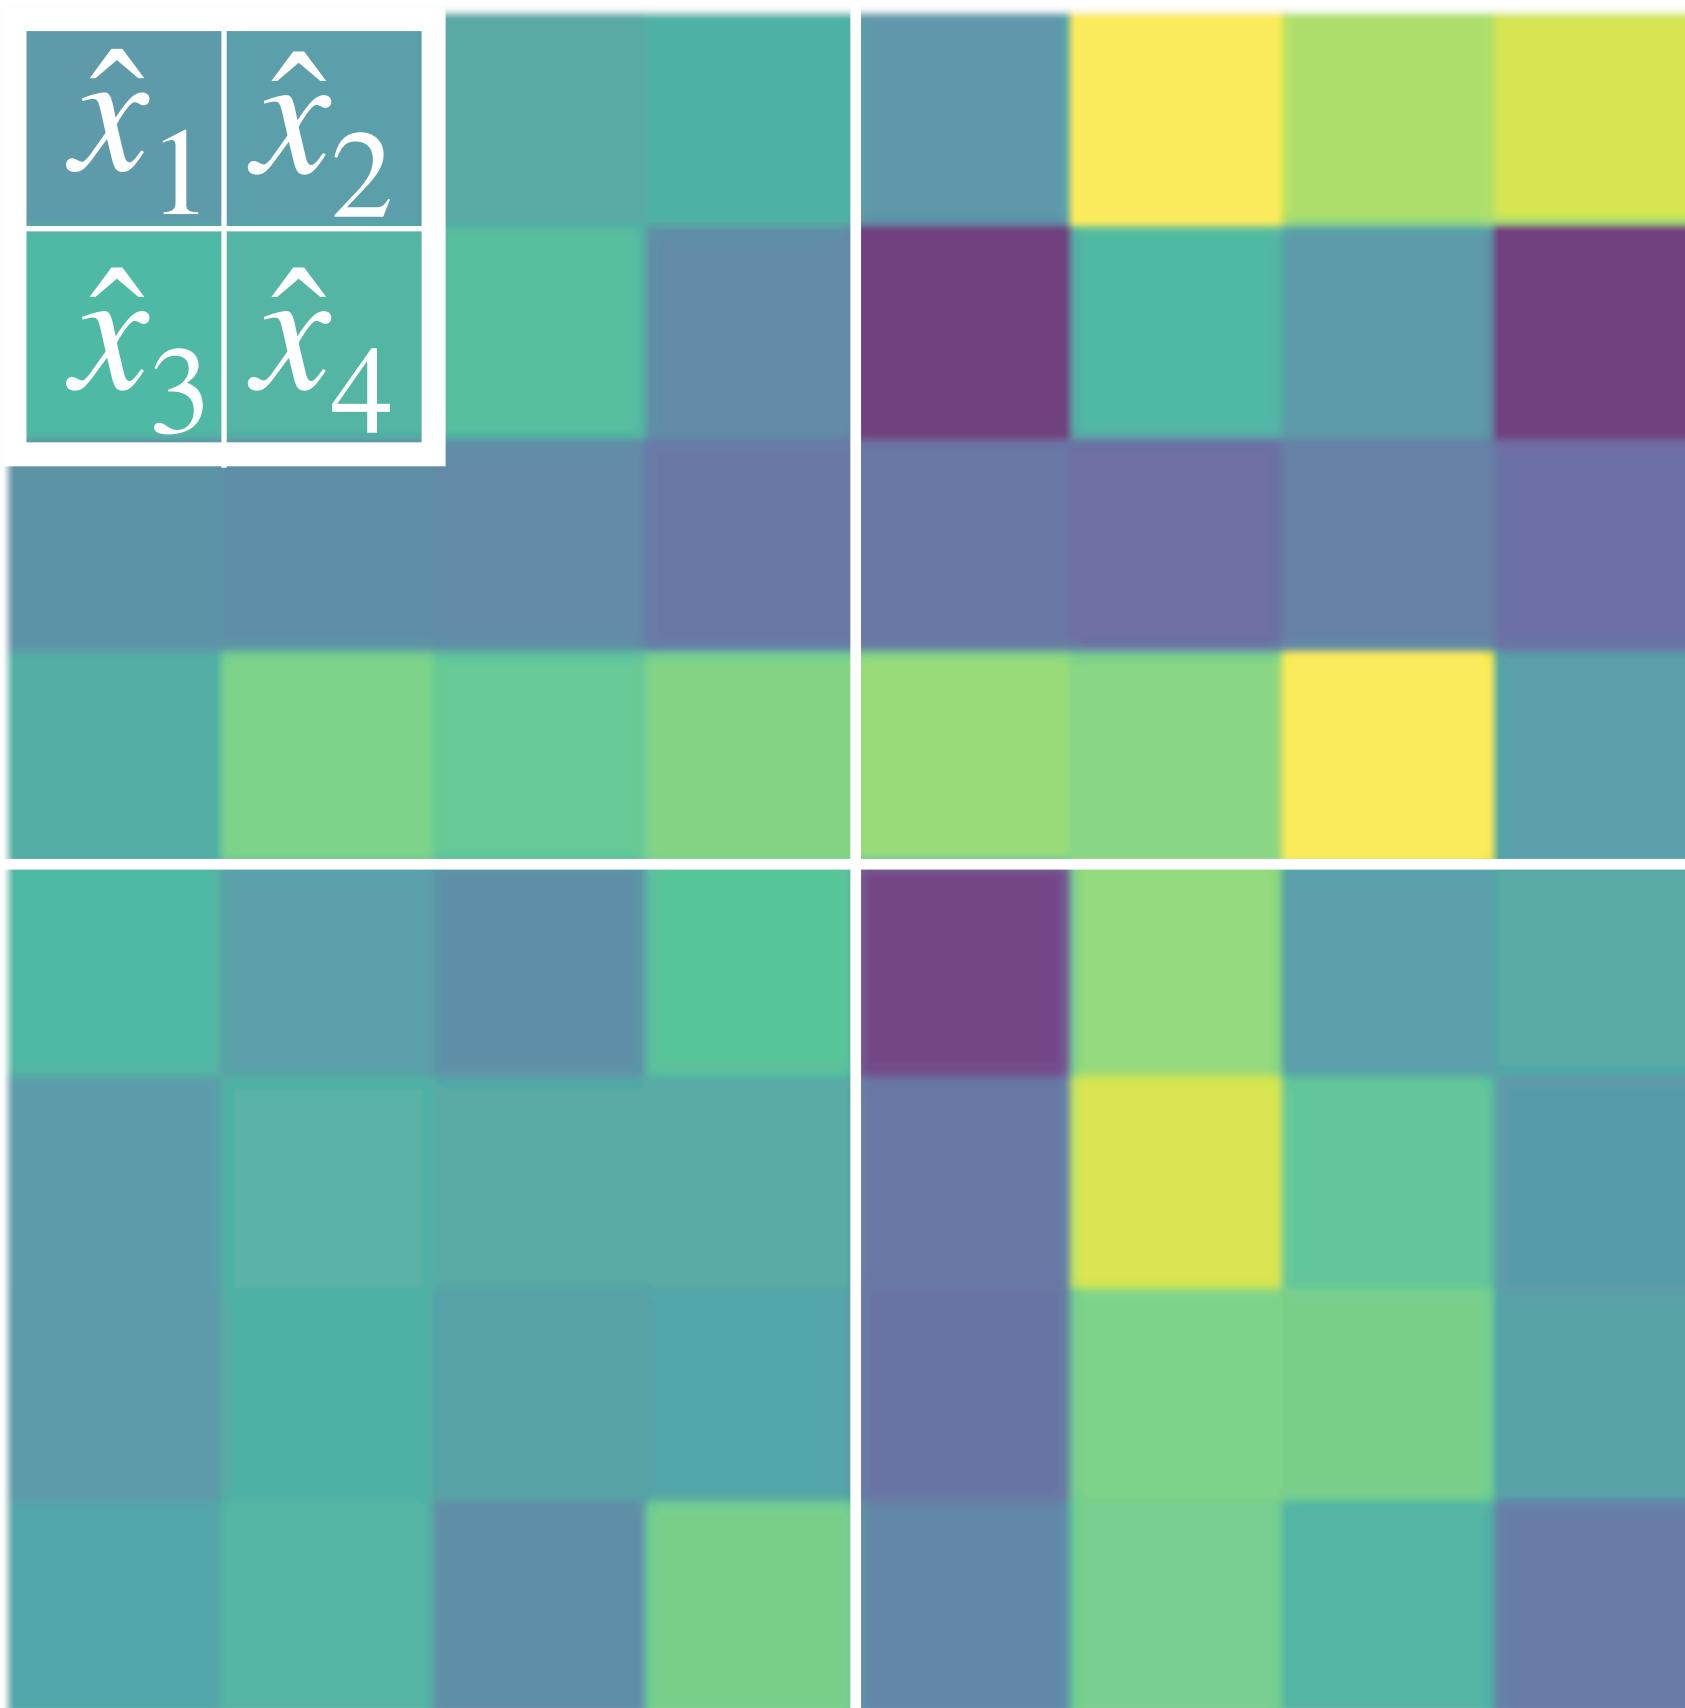


$$N_Q = \sum_{q=0}^3 x_q$$

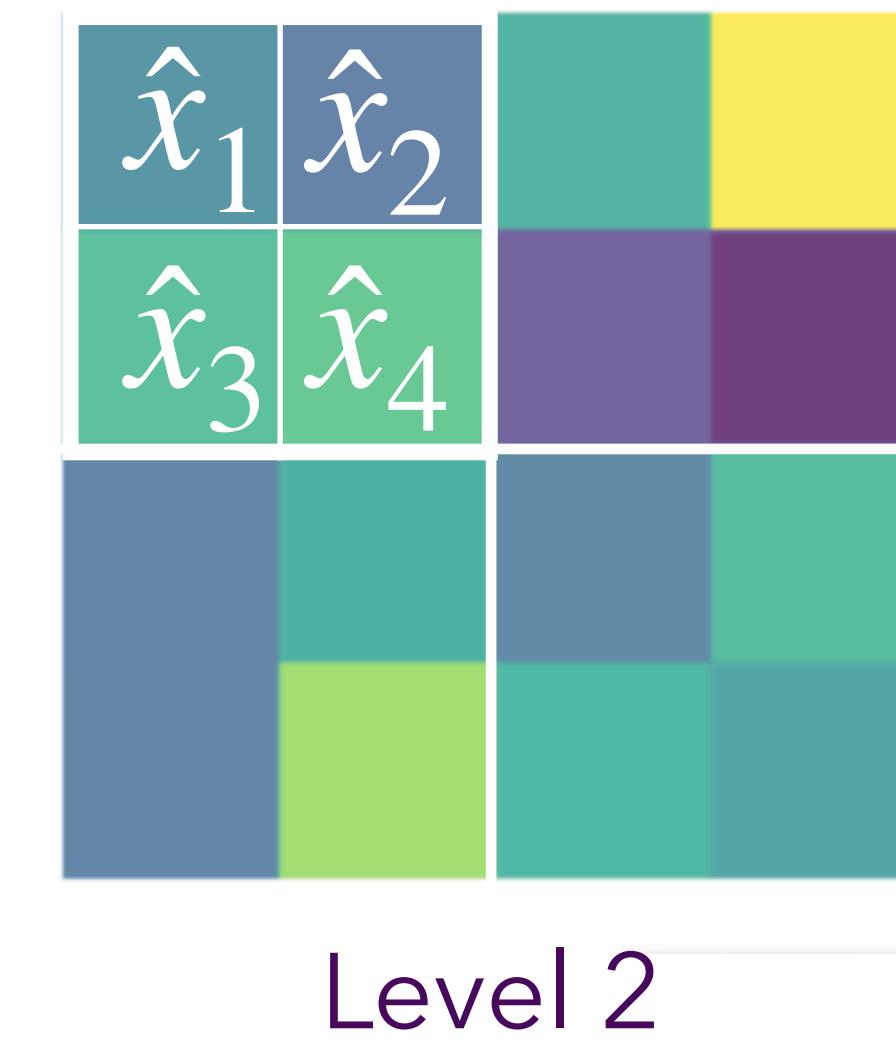
Diagram illustrating the multiscale decomposition. A 2x2 grid of squares is shown, with the top-left square labeled  $x_1$  and  $x_2$  in red, and the bottom-left square labeled  $x_3$  and  $x_4$  in red. The top-right square is teal and the bottom-right square is yellow. Dashed arrows point from the bottom-right square to a yellow box labeled  $N_Q$ , and from the entire grid to a purple box labeled  $N_Q$ . The text "Level k" is written below the grid, and "Level N" is written next to the dashed arrow pointing to  $N_Q$ .

# Multiscale decomposition

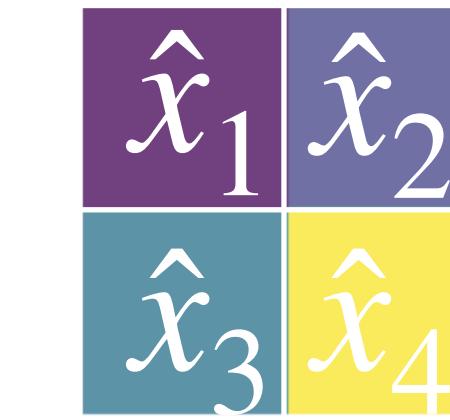
$$\hat{x}_q = x_q/N_Q$$



$$\hat{x}_q = x_q/N_Q$$



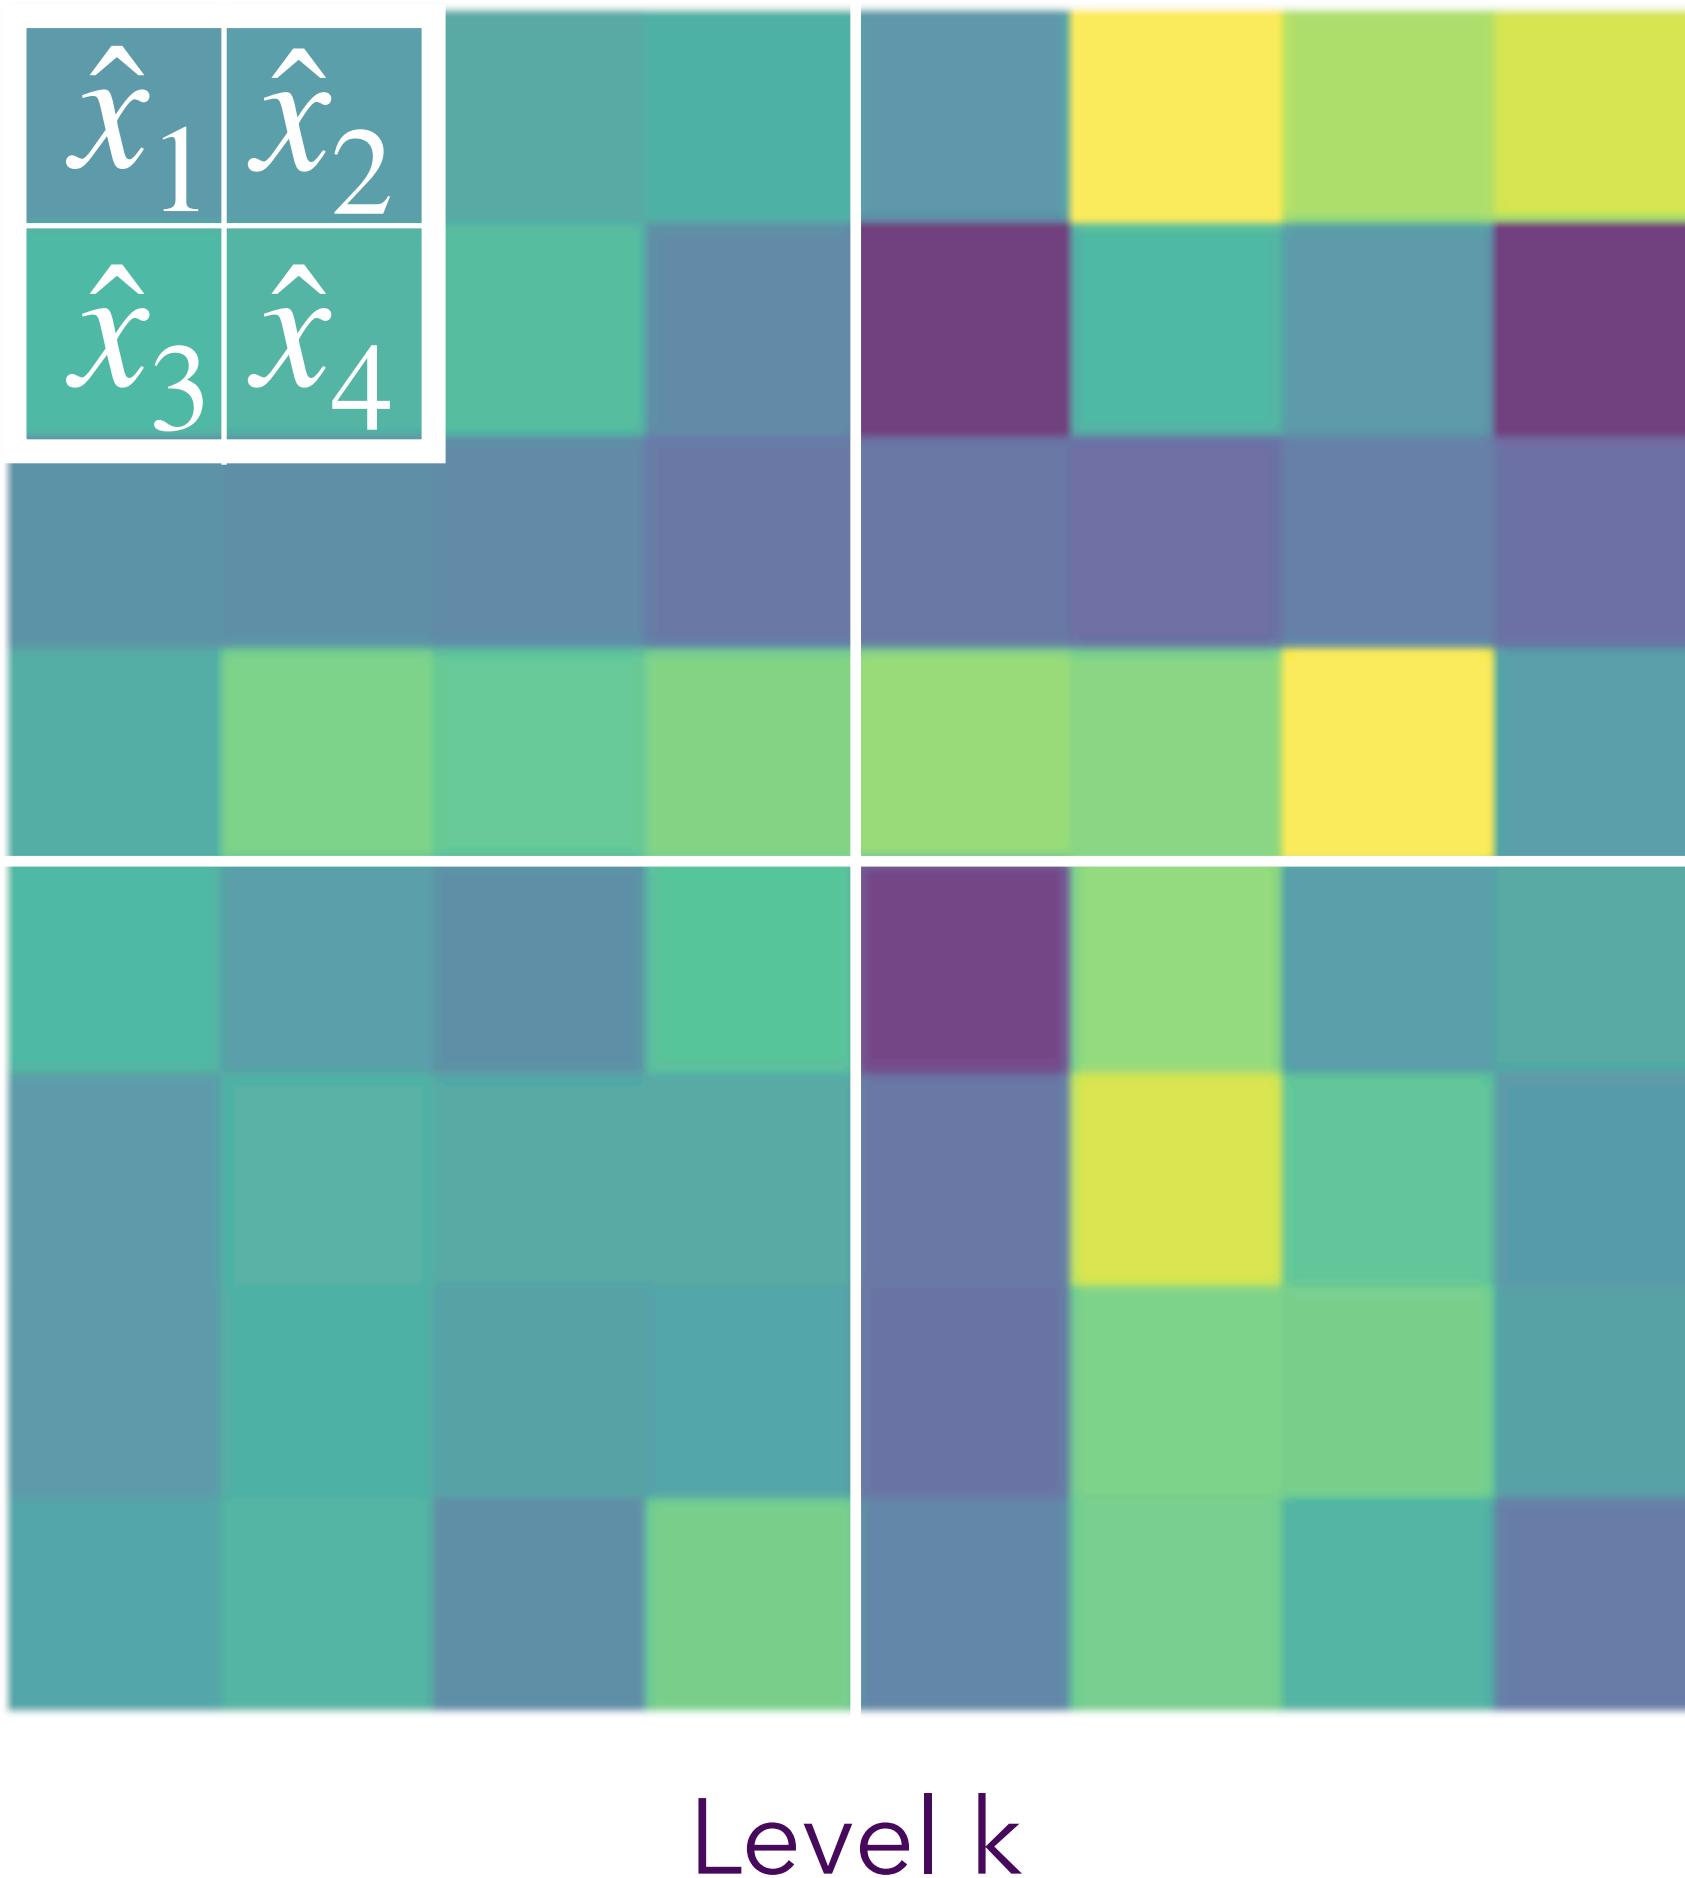
$$\hat{x}_q = x_q/N_Q$$



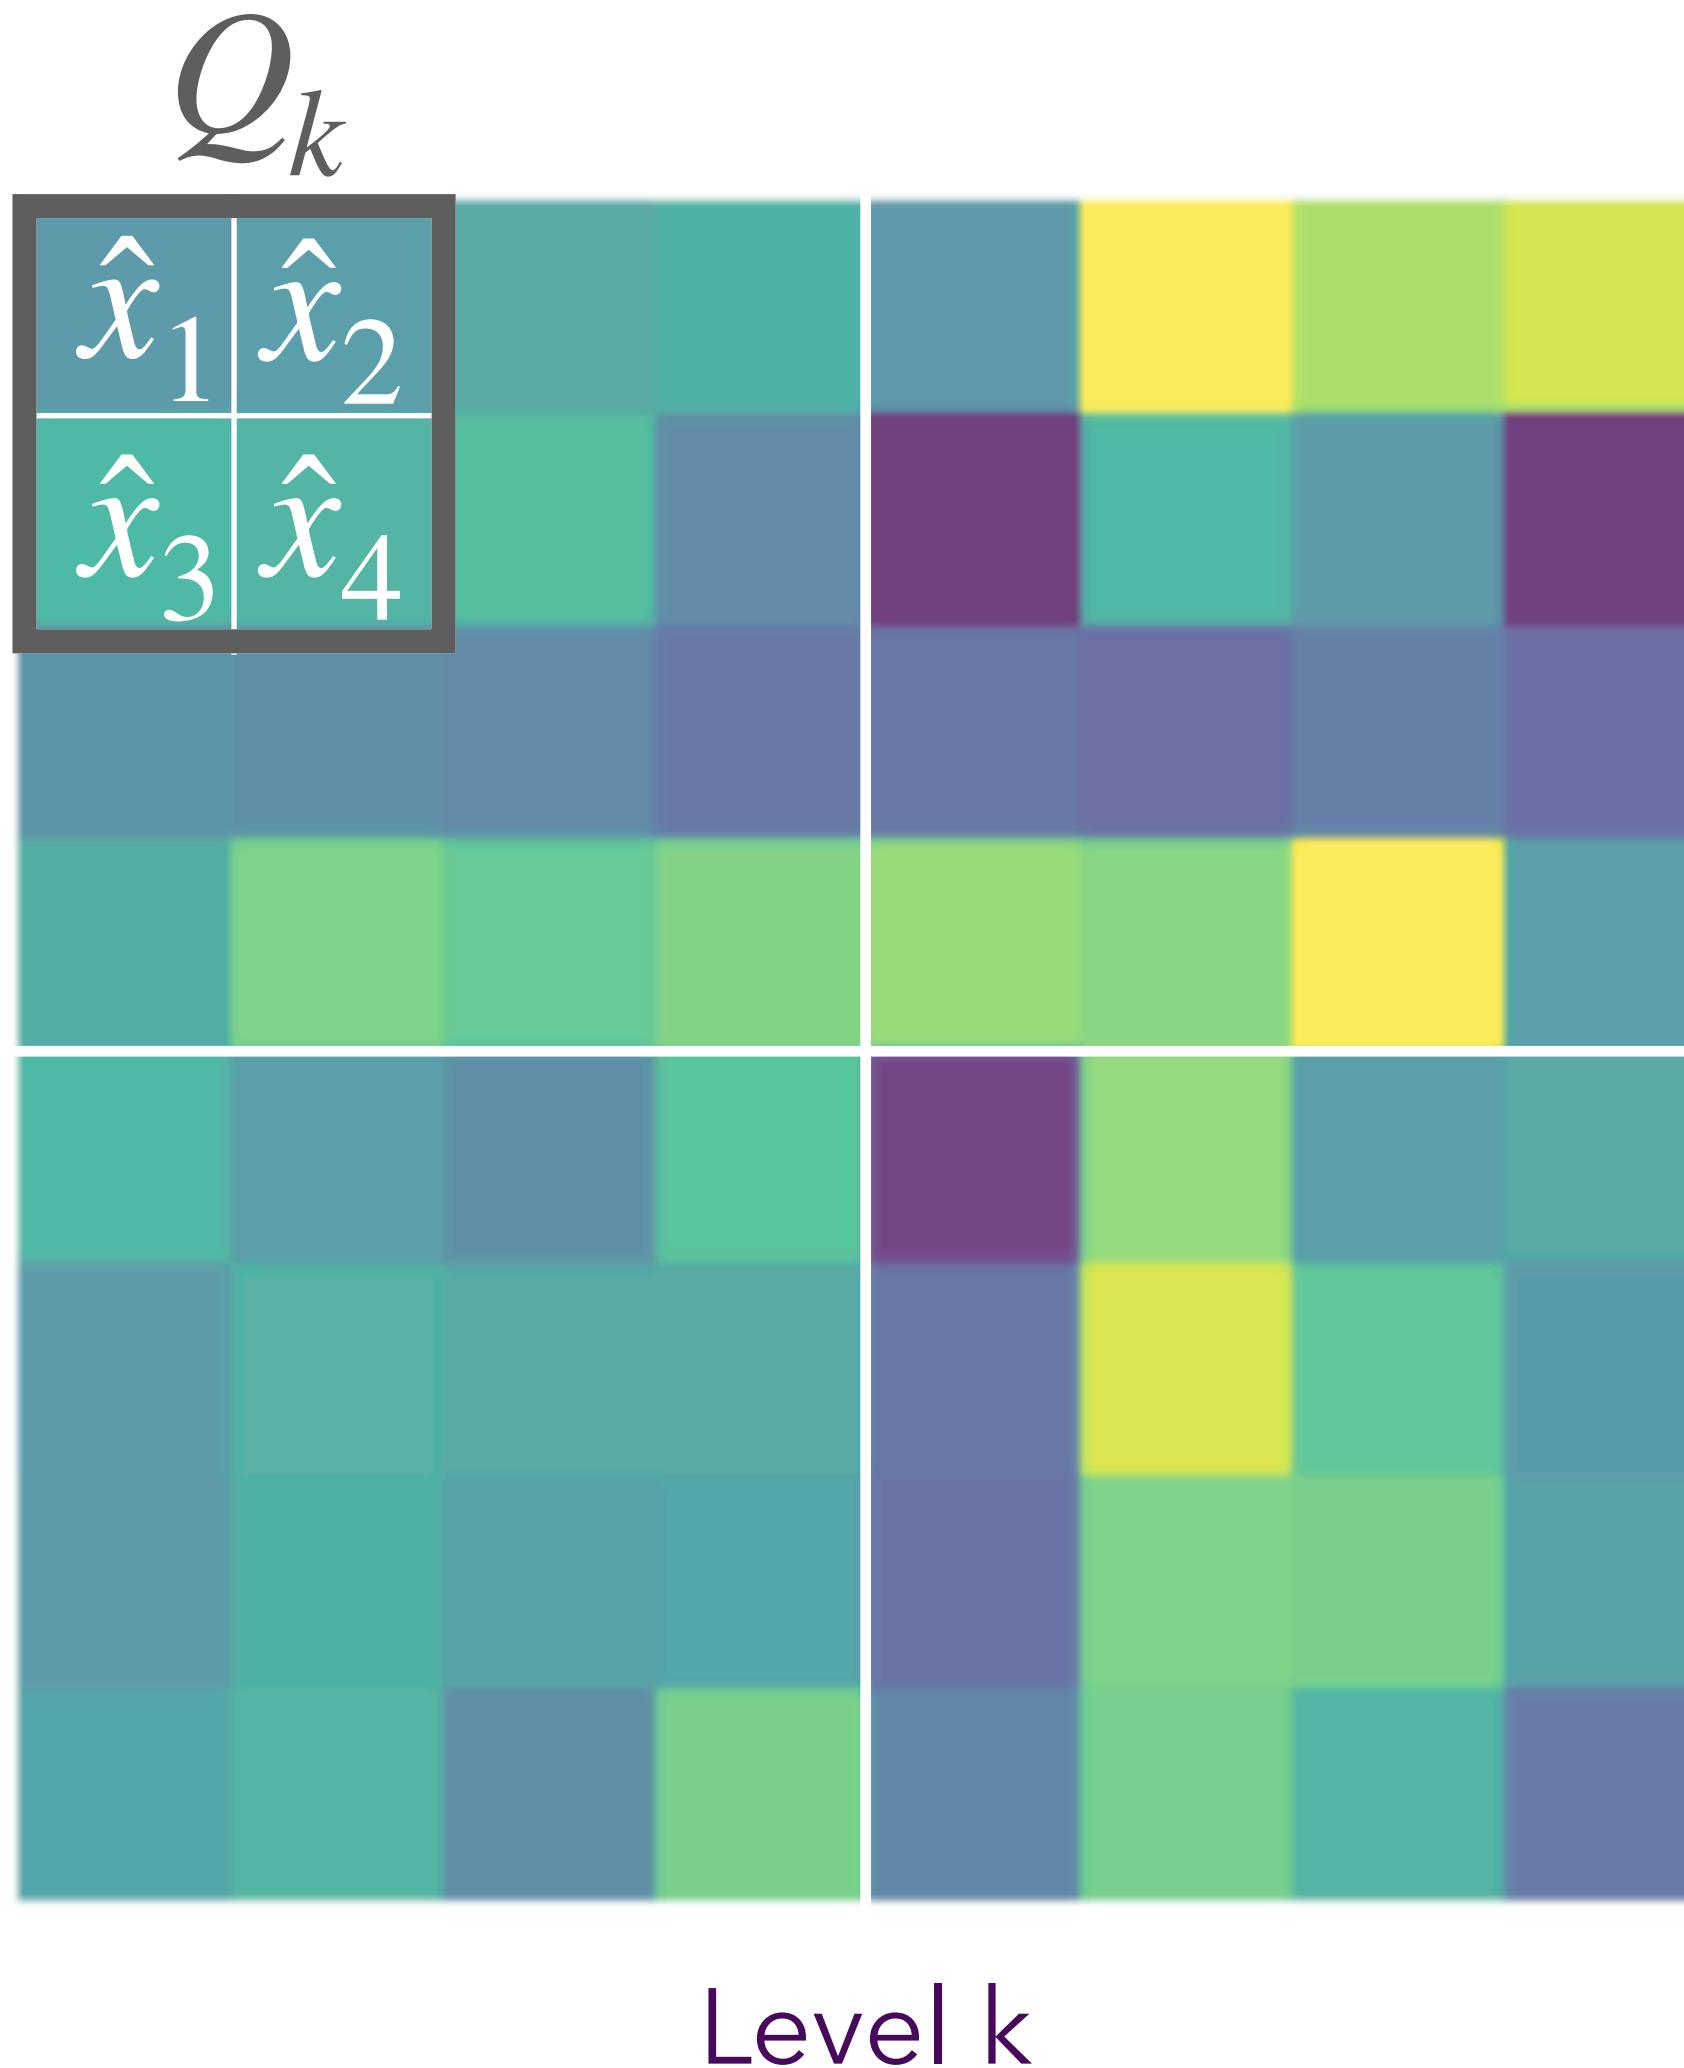
$N_Q$

Level N

# Multiscale prior



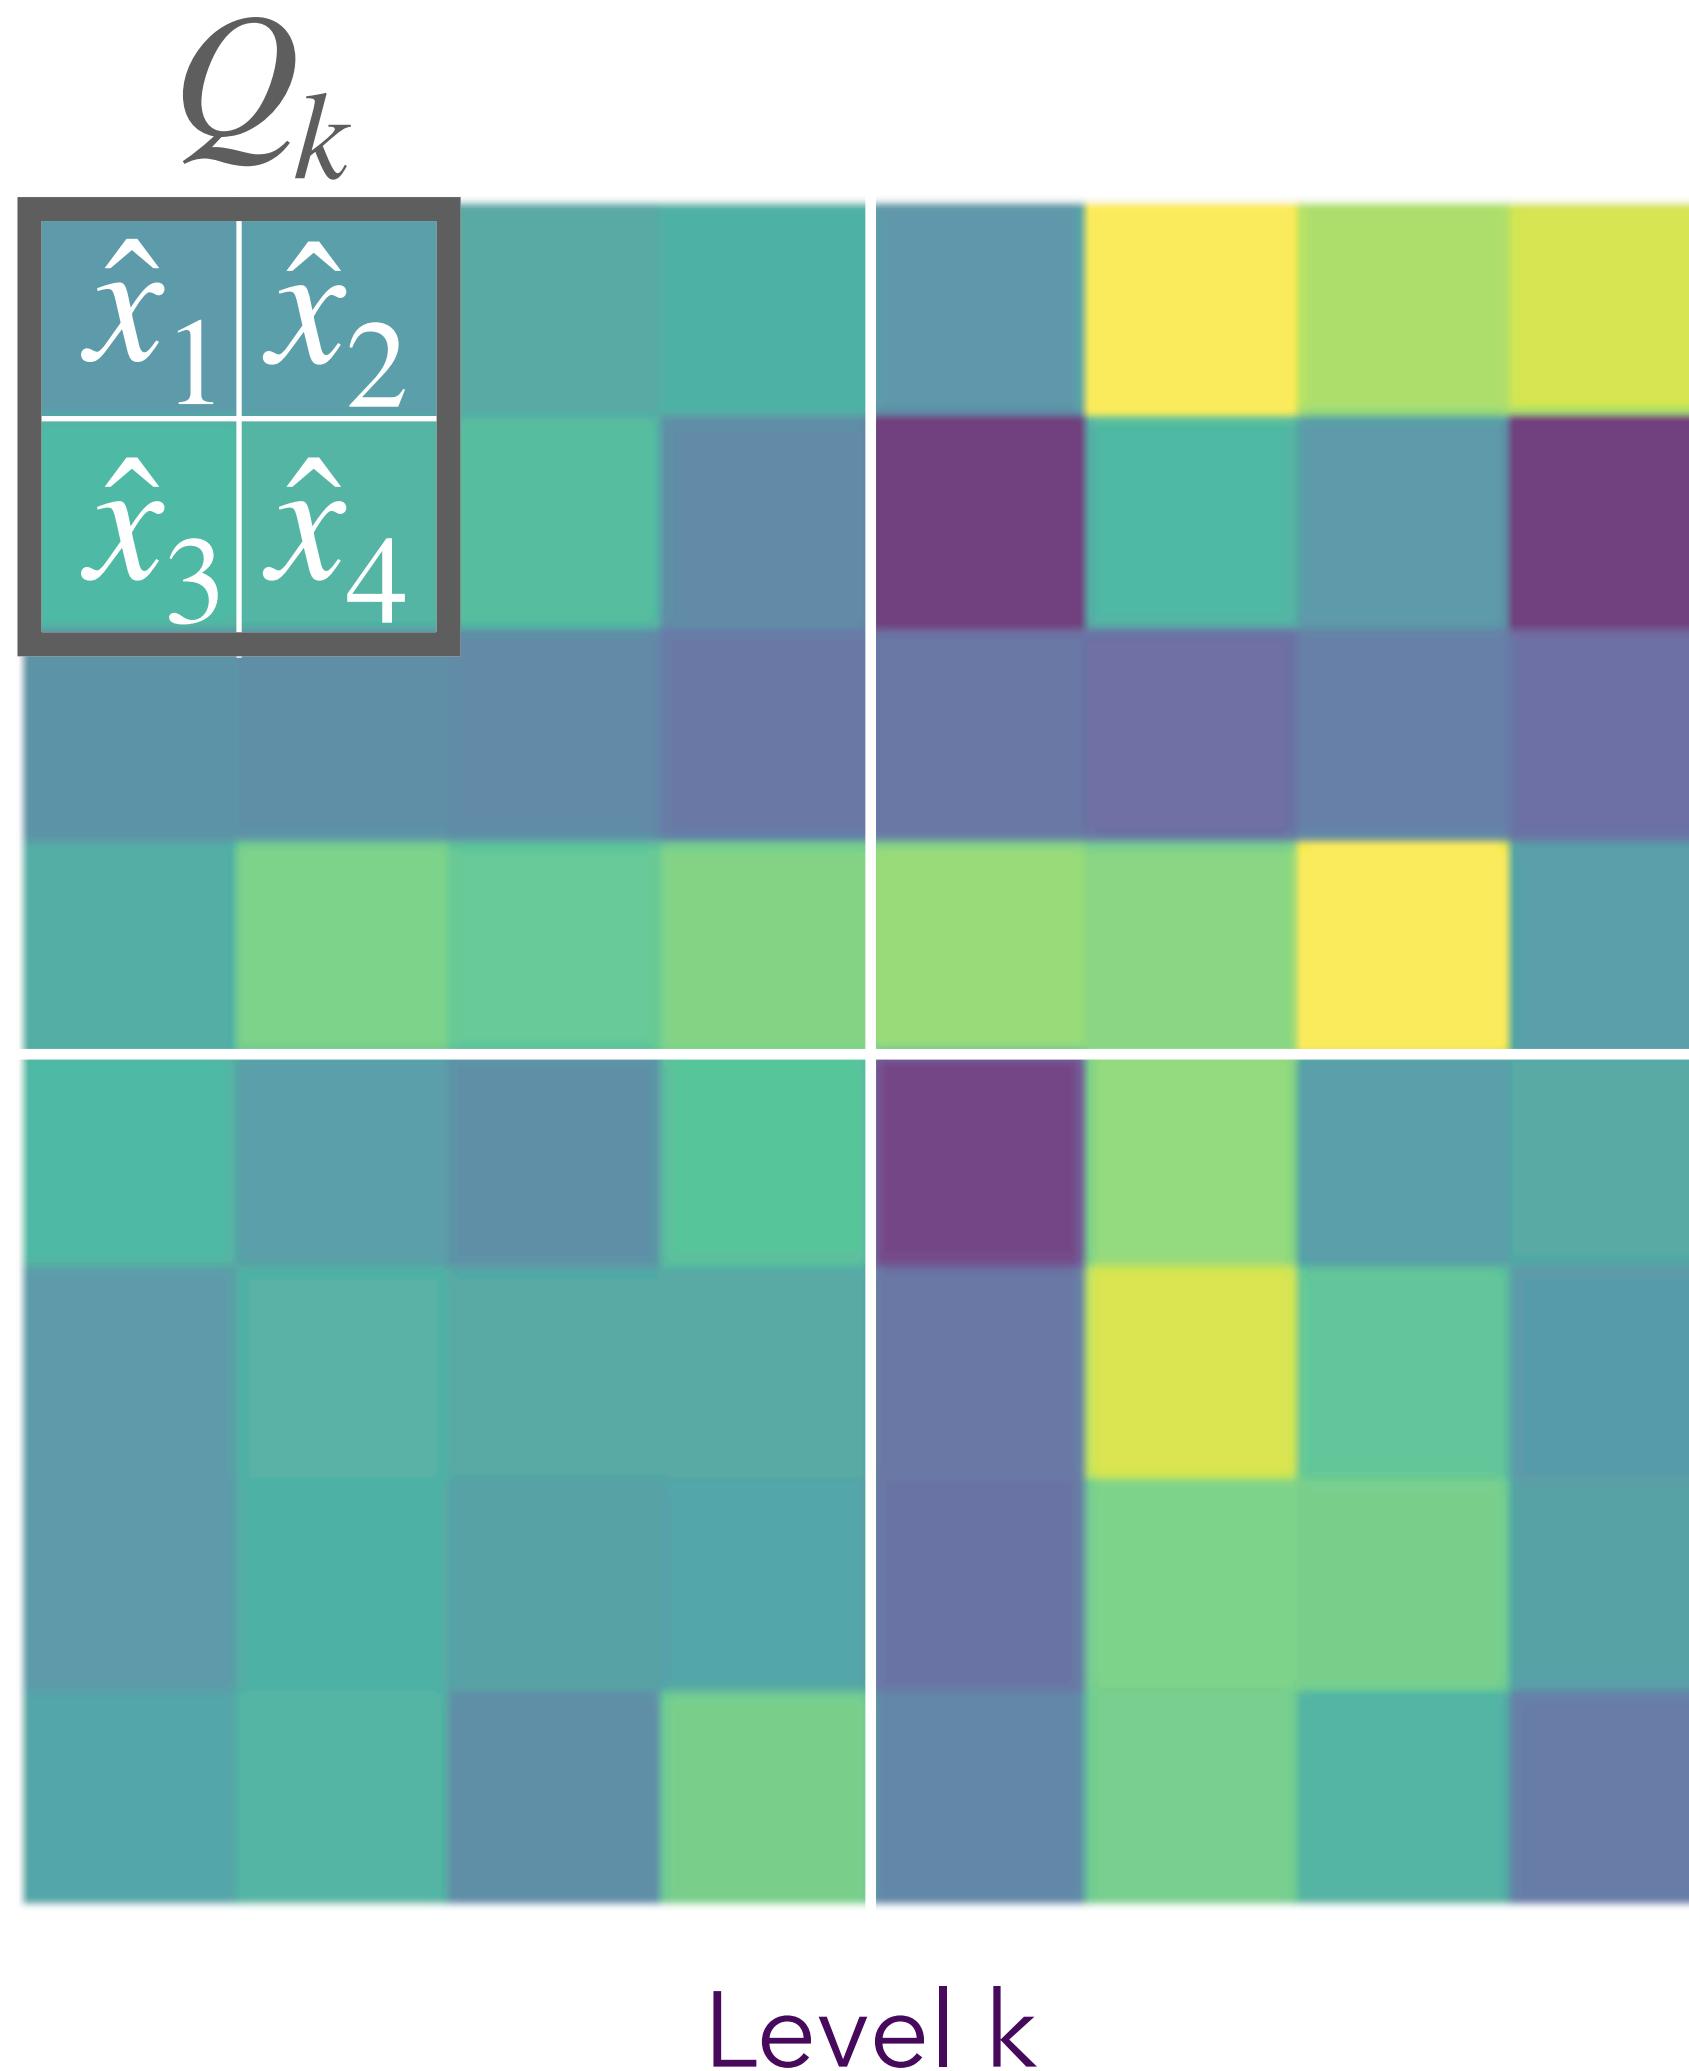
# Multiscale prior



The **Dirichlet distribution** is a multivariate distribution, where all **variables sum to unity**. And the larger the parameter  $\alpha_k$  the more similar the variables are.

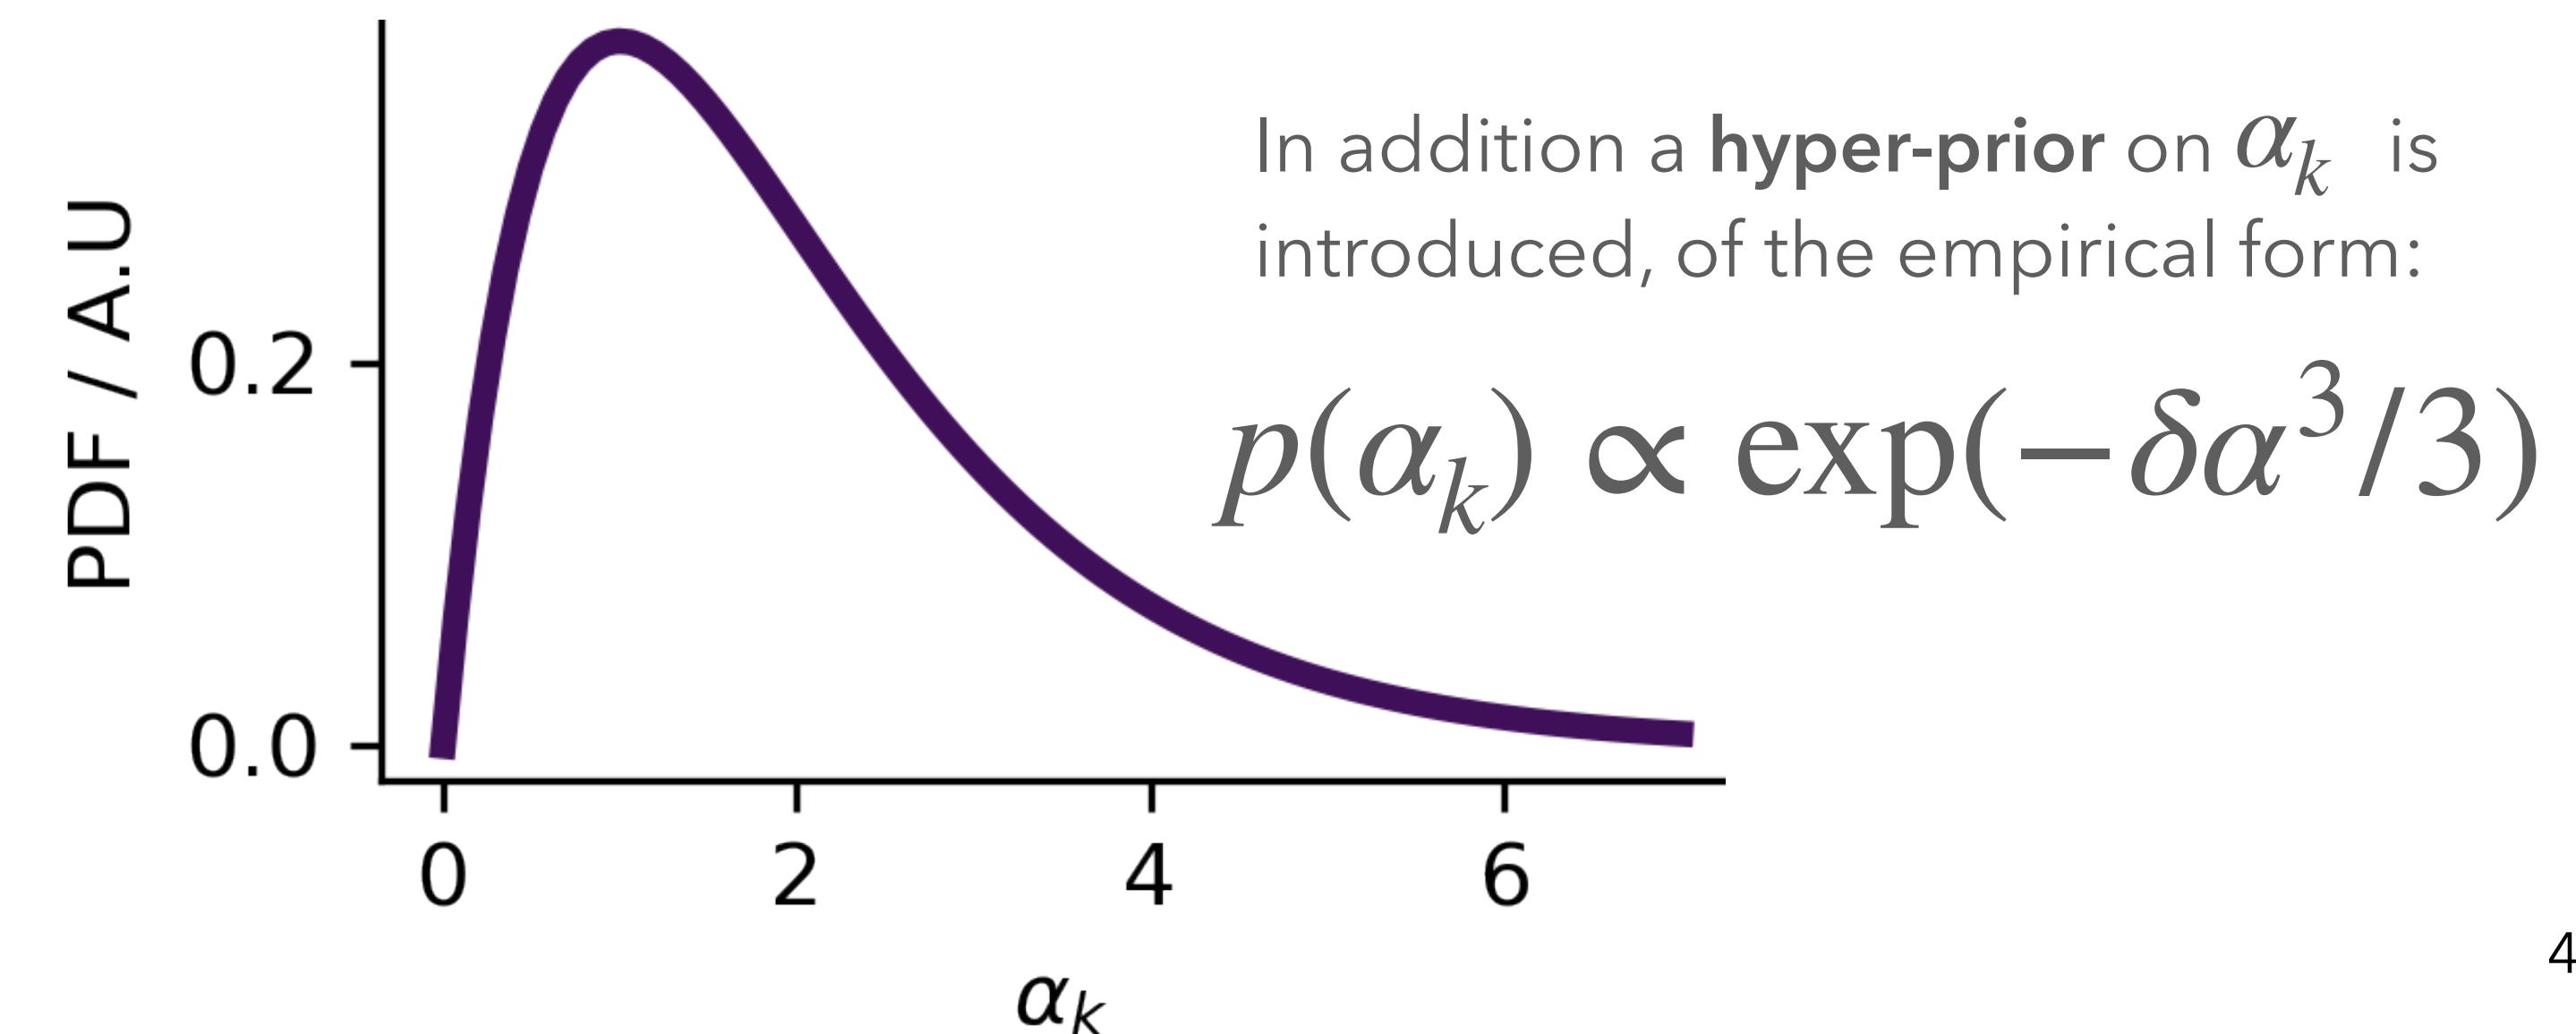
$$Q_k \propto \text{Dirichlet}(\alpha_k, \alpha_k, \alpha_k, \alpha_k)$$

# Multiscale prior

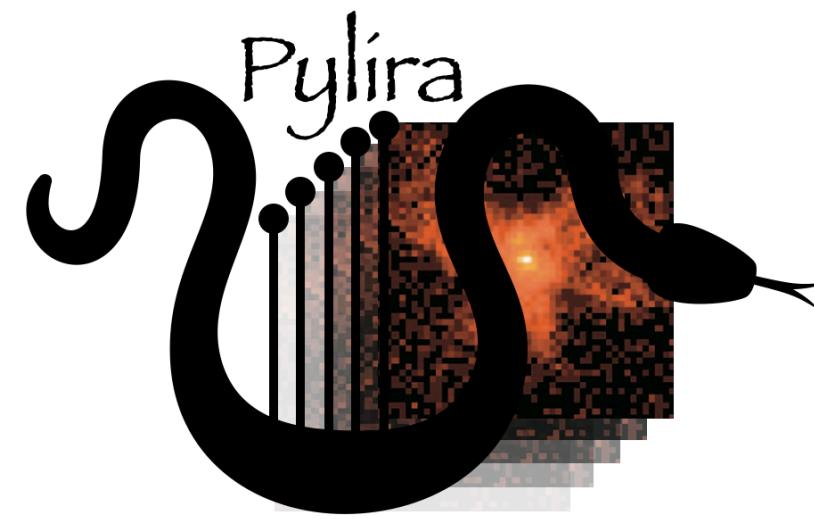


The **Dirichlet distribution** is a multivariate distribution, where all **variables sum to unity**. And the larger the parameter  $\alpha_k$  the more similar the variables are.

$$Q_k \propto \text{Dirichlet}(\alpha_k, \alpha_k, \alpha_k, \alpha_k)$$

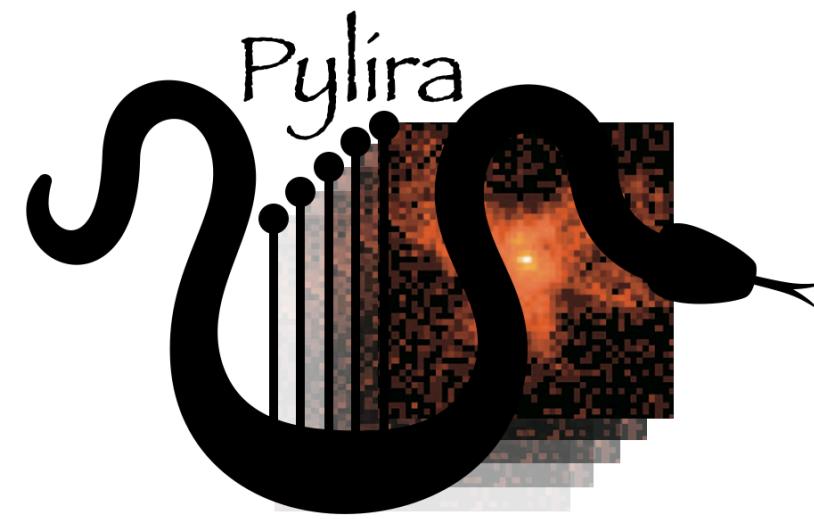


# Pylira Python Package



# Pylira Python Package

- The LIRA method was **initially implemented as an R package** by [[Connors et al. 2011](#)].  
Original code available at <https://github.com/astrostat/LIRA>
- Meanwhile Python has become the favored language of choice for astronomical data analysis.
- **Pylira is a Python wrapper** around the initial C implementation implemented via [[pybind11](#)] and based on the [[Astropy affiliated package template](#)]
- Additional dependencies are Numpy, Scipy, Astropy, Matplotlib, ... 
- Github: <https://github.com/astrostat/pylira>
- Docs: <https://pylira.readthedocs.io/en/latest/>



# Pylira Python Package

- The LIRA method was **initially implemented as an R package** by [[Connors et al. 2011](#)].  
Original code available at <https://github.com/astrostat/LIRA>
- Meanwhile Python has become the favored language of choice for astronomical data analysis.
- **Pylira is a Python wrapper** around the initial C implementation implemented via [[pybind11](#)] and based on the [[Astropy affiliated package template](#)]
- Additional dependencies are Numpy, Scipy, Astropy, Matplotlib, ...
- Github: <https://github.com/astrostat/pylira>
- Docs: <https://pylira.readthedocs.io/en/latest/>

*pybind11*



NumPy



SciPy

matplotlib



# A short code example...

```
import numpy as np
from pylira import LIRADeconvolver
from pylira.data import point_source_gauss_psf

random_state = np.RandomState(372)

# load built in test dataset
data = point_source_gauss_psf()

data["flux_init"] = random_state.gamma(10, size=(32, 32))

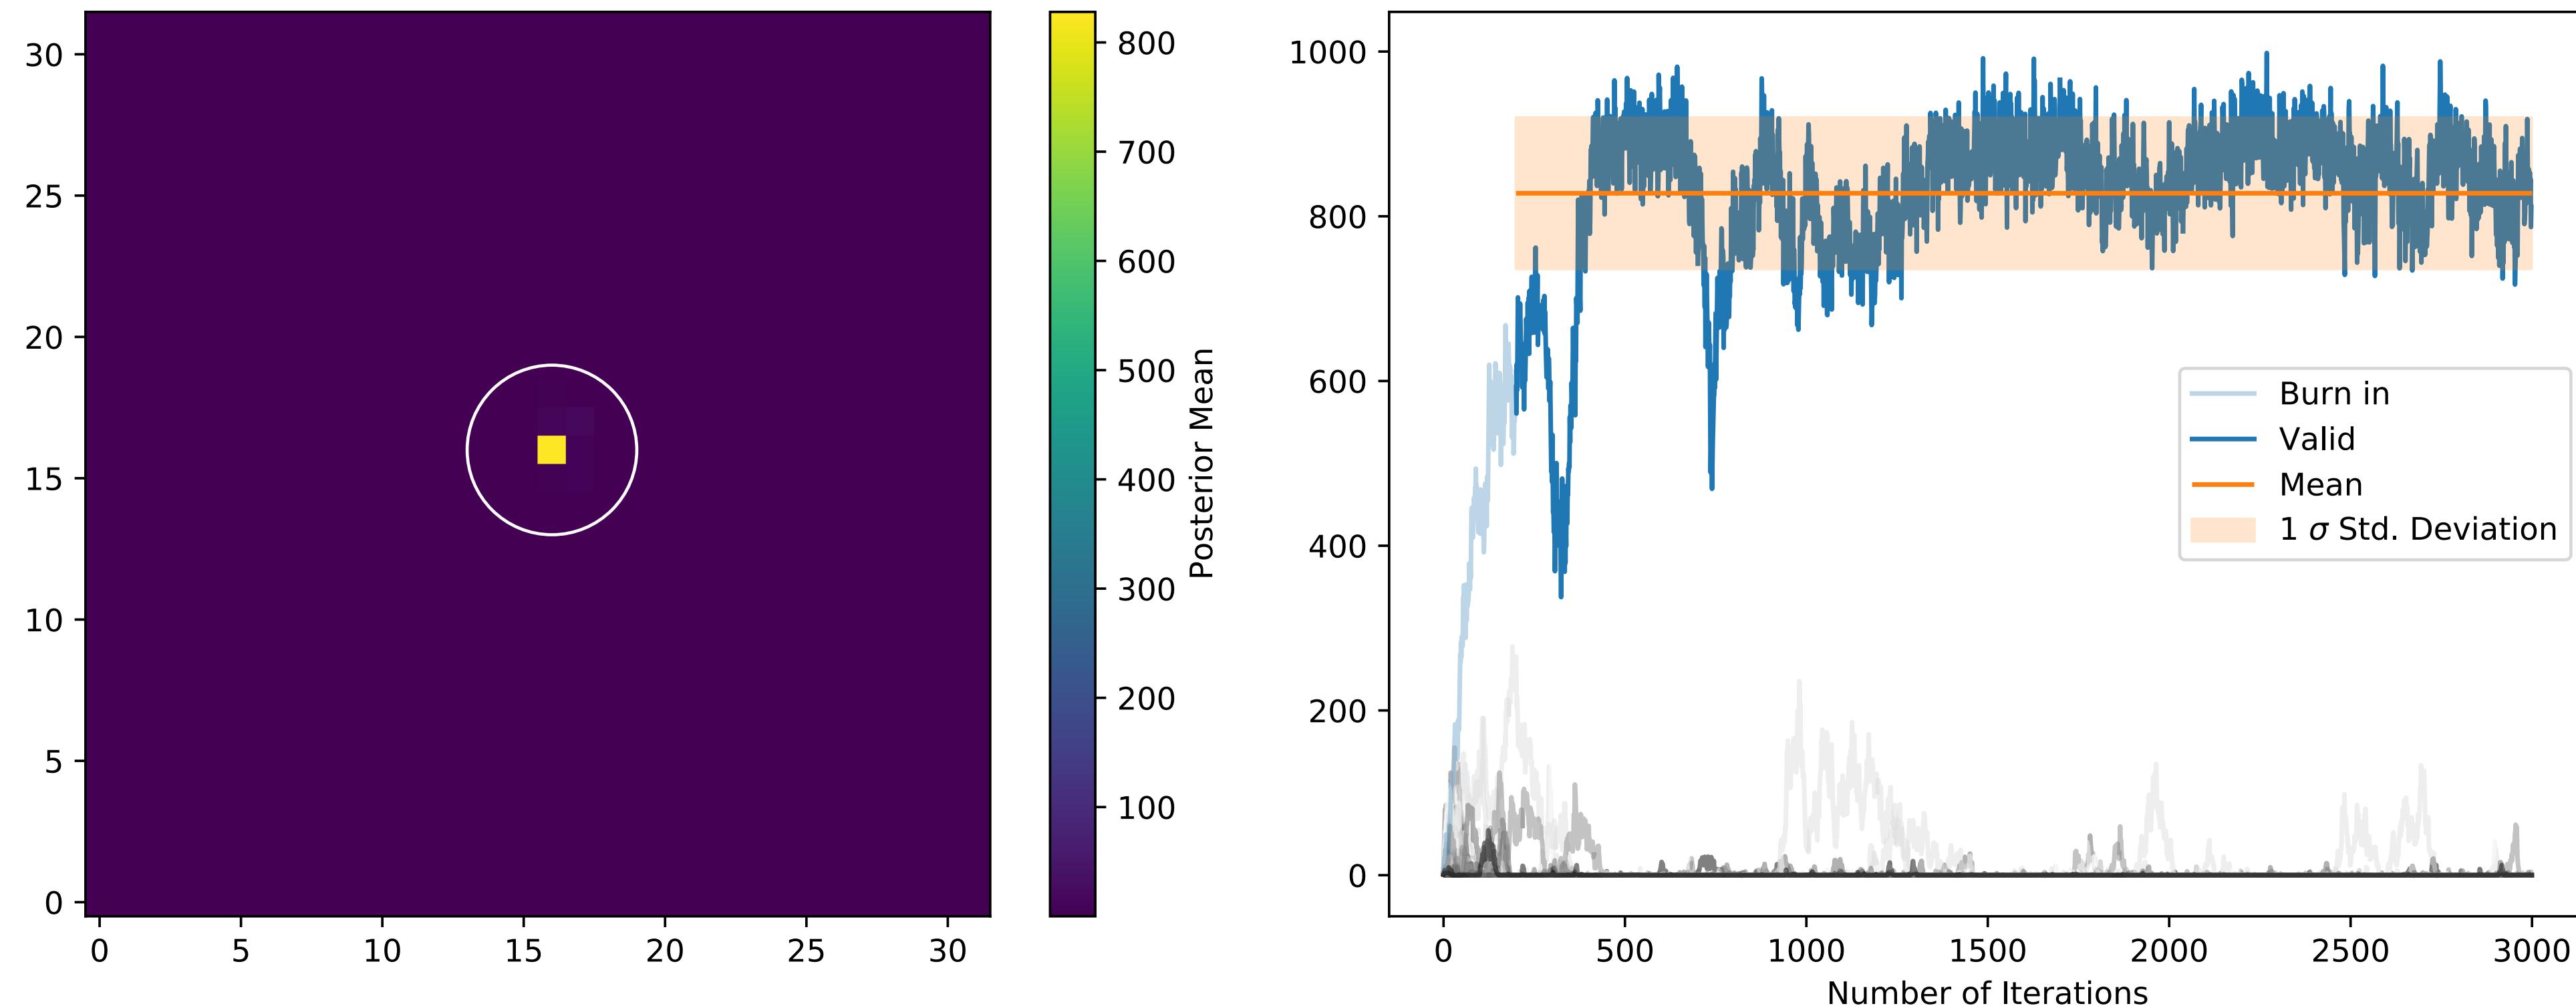
# sklearn inspired API, take config on init
deconvolve = LIRADeconvolver(
    alpha_init=np.ones(5),
    n_iter_max=5_000,
    n_burn_in=500,
    random_state=random_state,
)

# run algorithm
result = deconvolve.run(data=data)

# serialise to FITS
result.write("pylira-result.fits")
```

[[From Simple Point Source Tutorial](#)]

# Diagnostic plot I

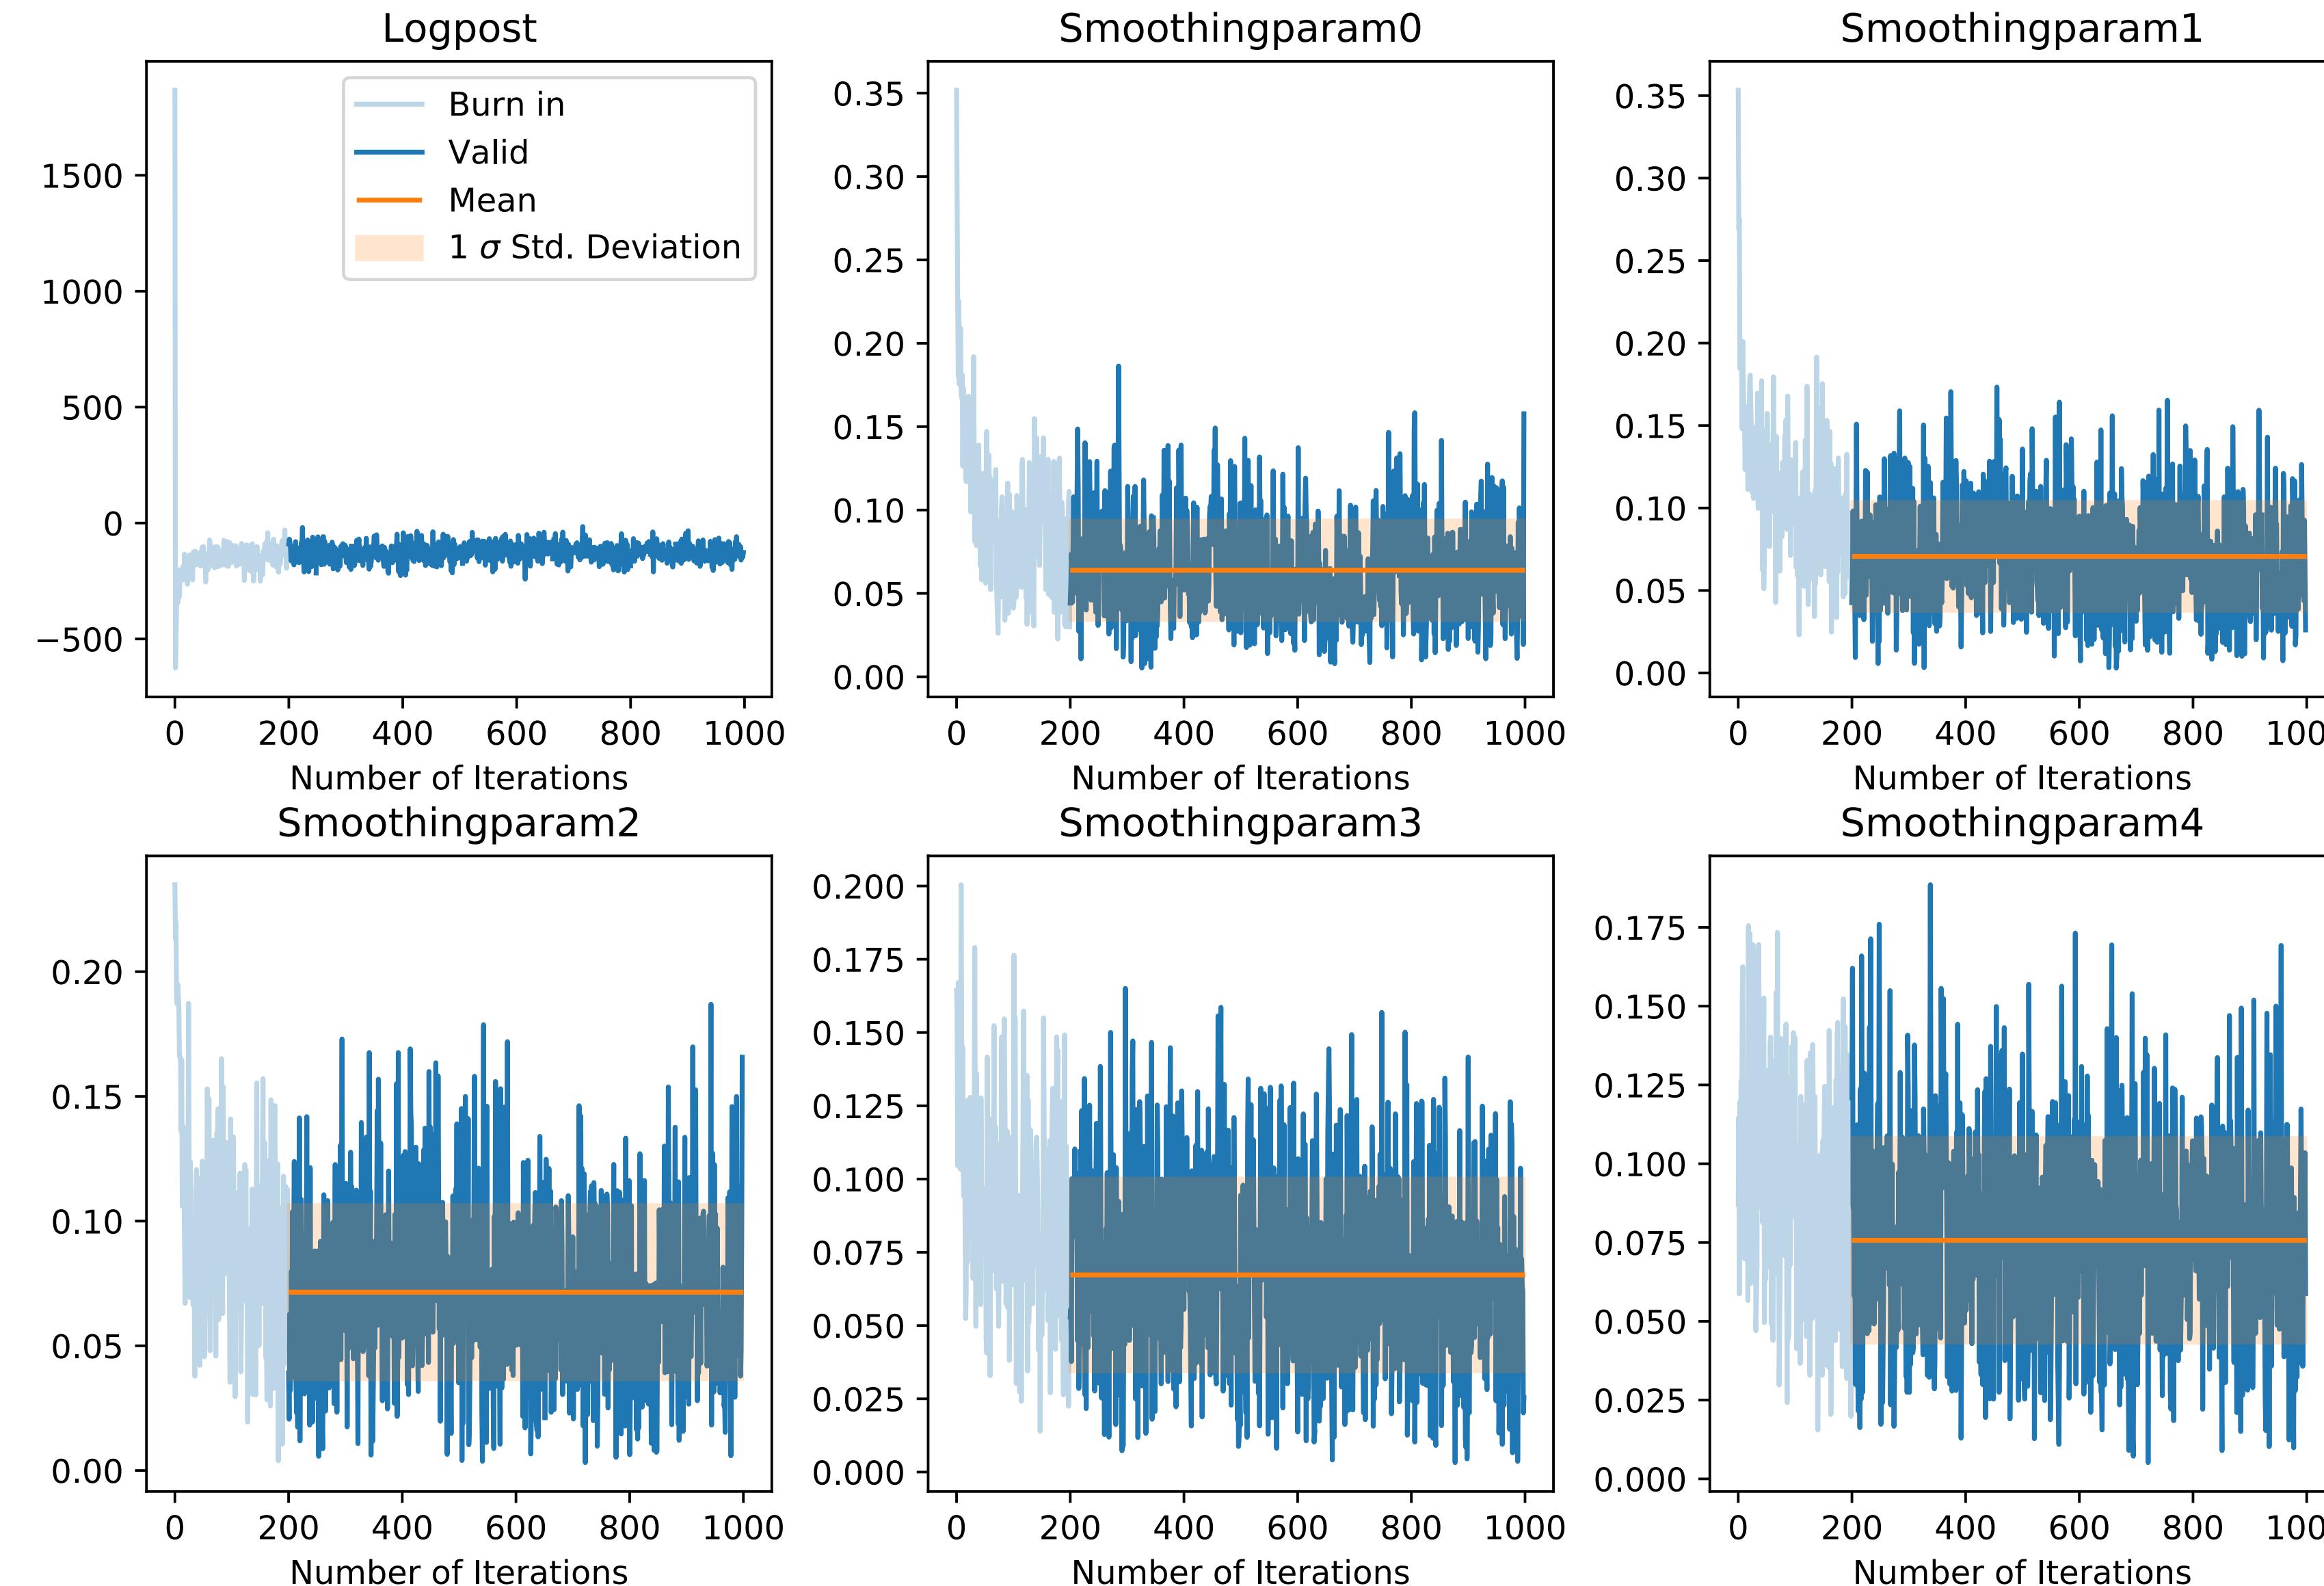


```
from pylira import LIRADeconvolverResult

# read result back from disk
result = LIRADeconvolverResult.read("pylira-result.fits")

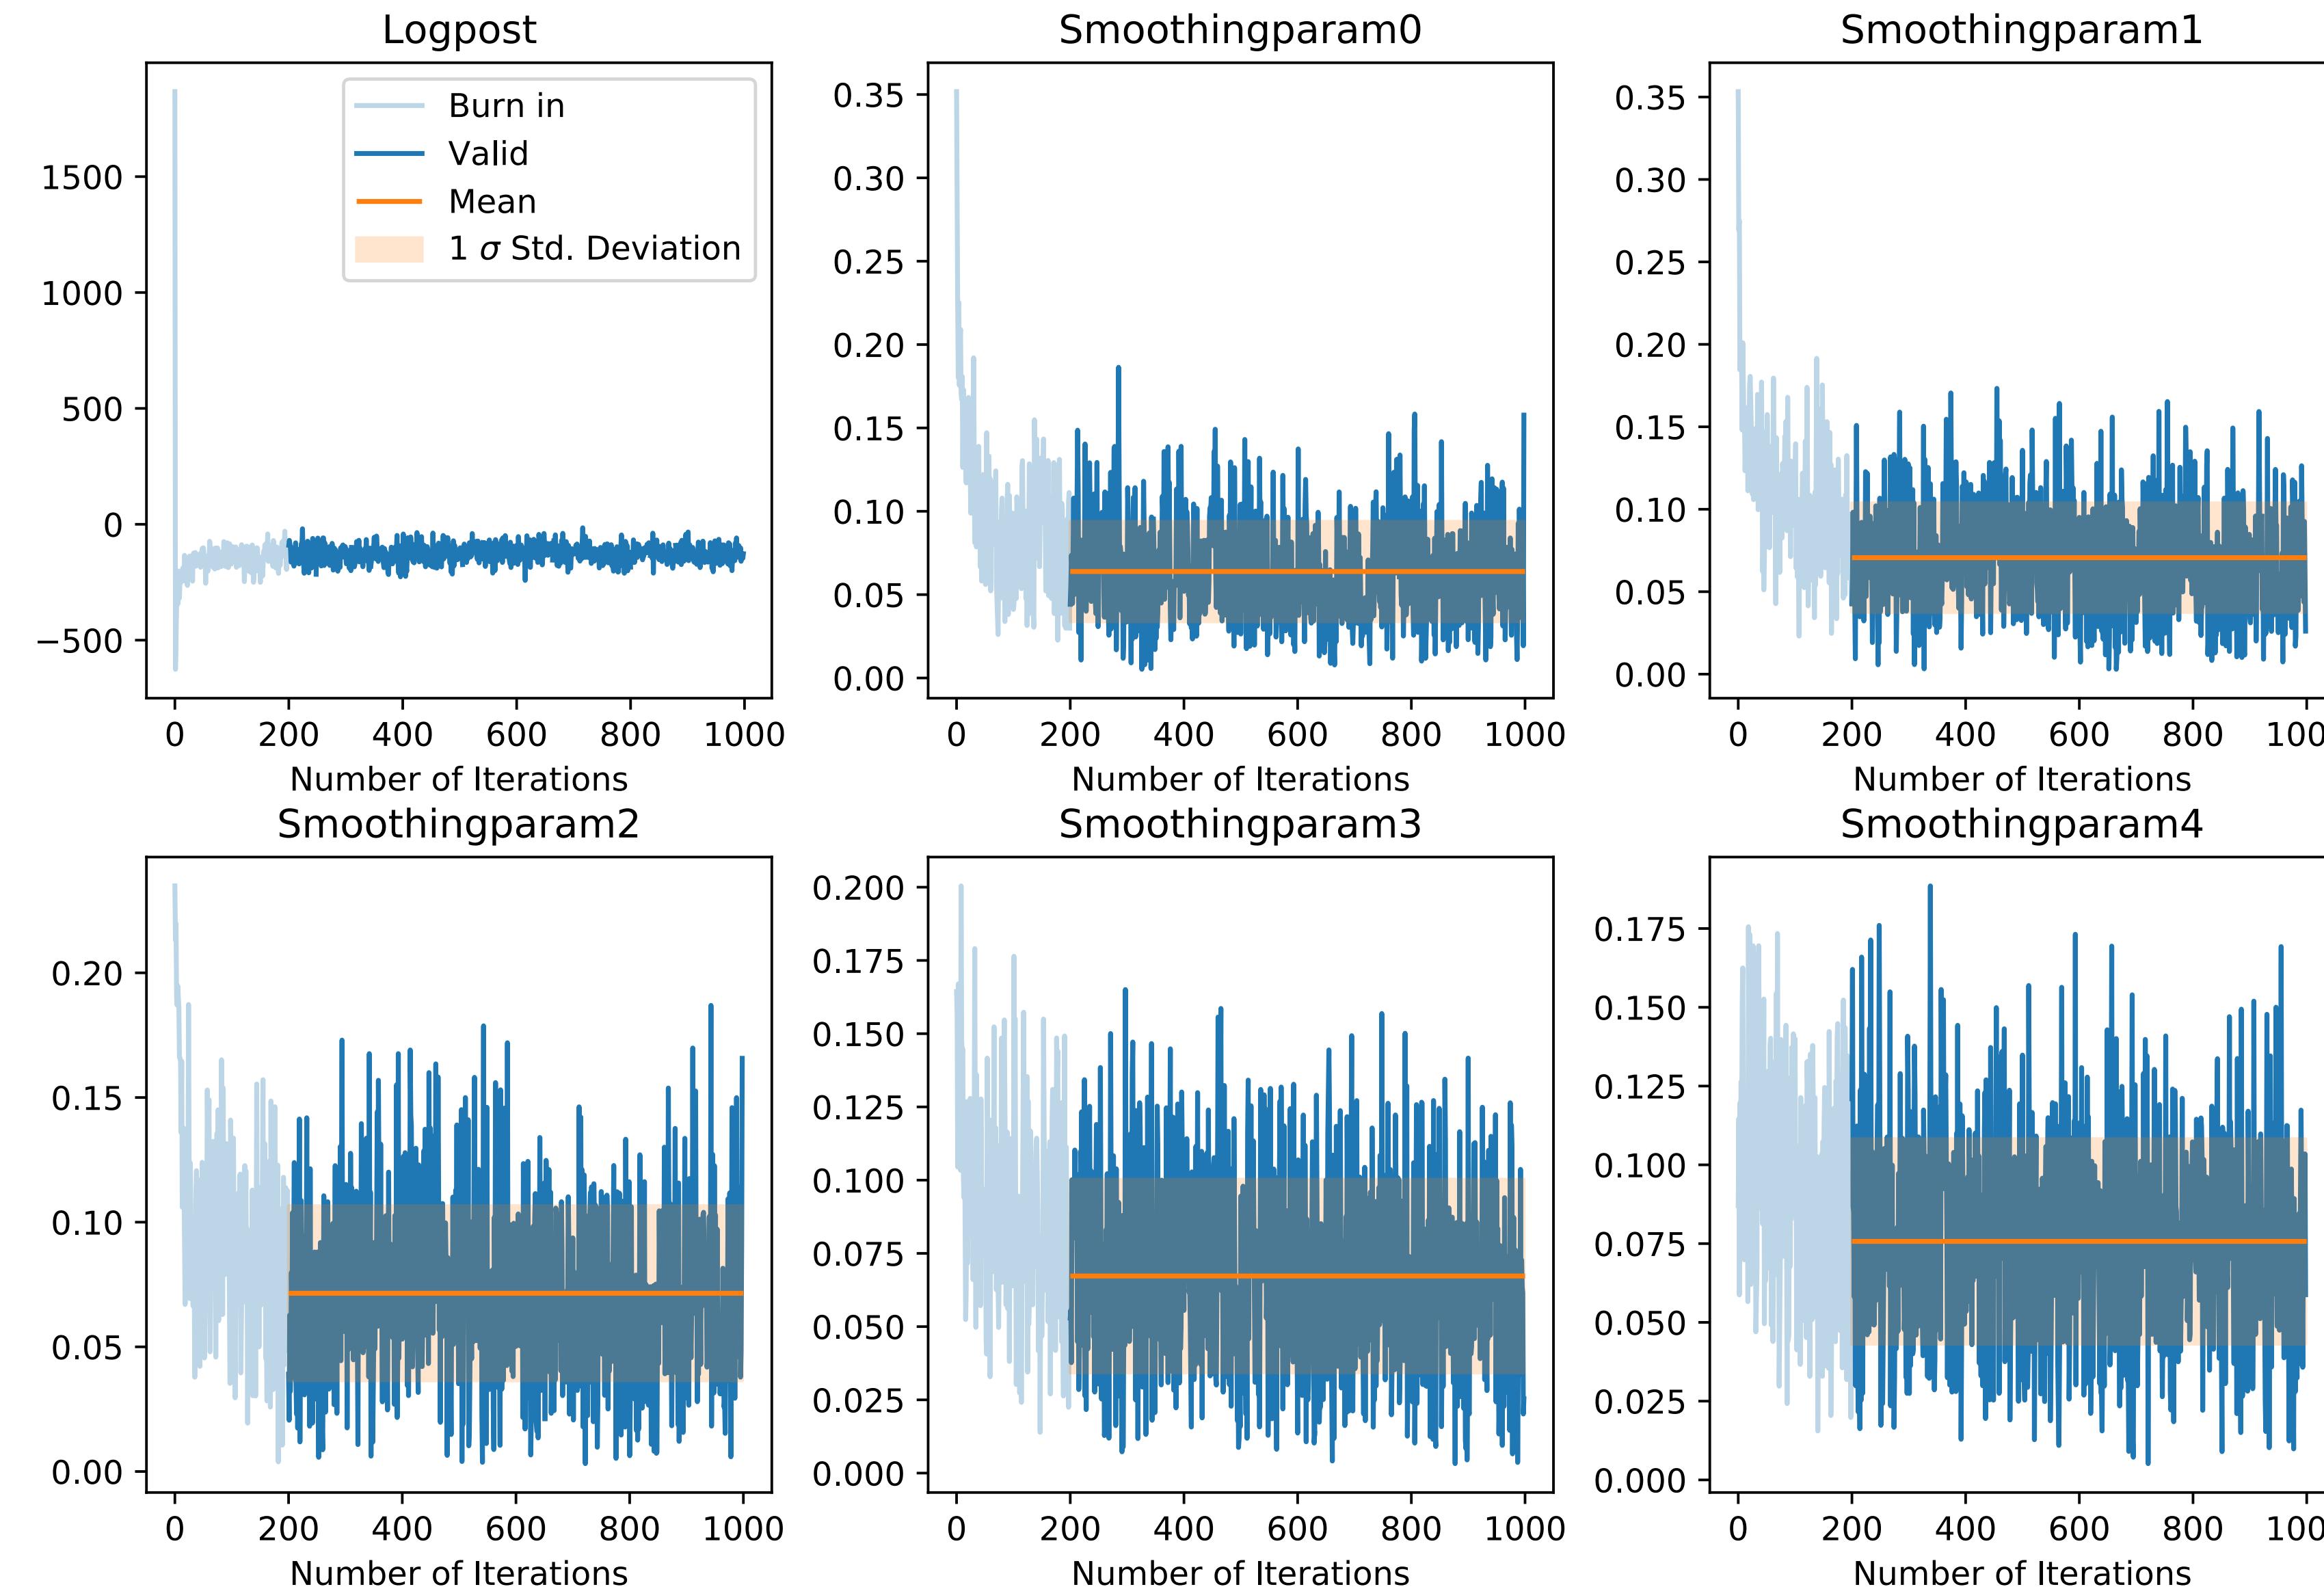
# diagnostic plot
result.plot_pixel_traces_region(center_pix=(16, 16), radius_pix=2)
```

# Diagnostic plot II



```
from pylira import LIRADeconvolverResult  
# read result back from disk  
result = LIRADeconvolverResult.read(  
    "pylira-result.fits"  
)  
  
# diagnostic plot  
result.plot_parameter_traces()
```

# Diagnostic plot II

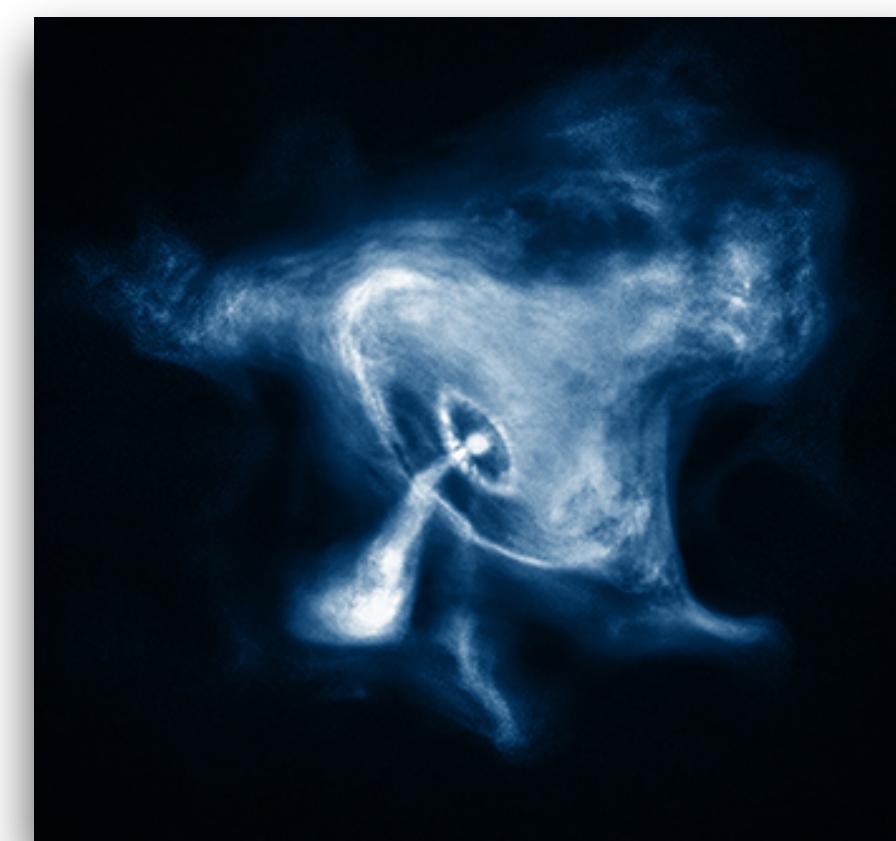
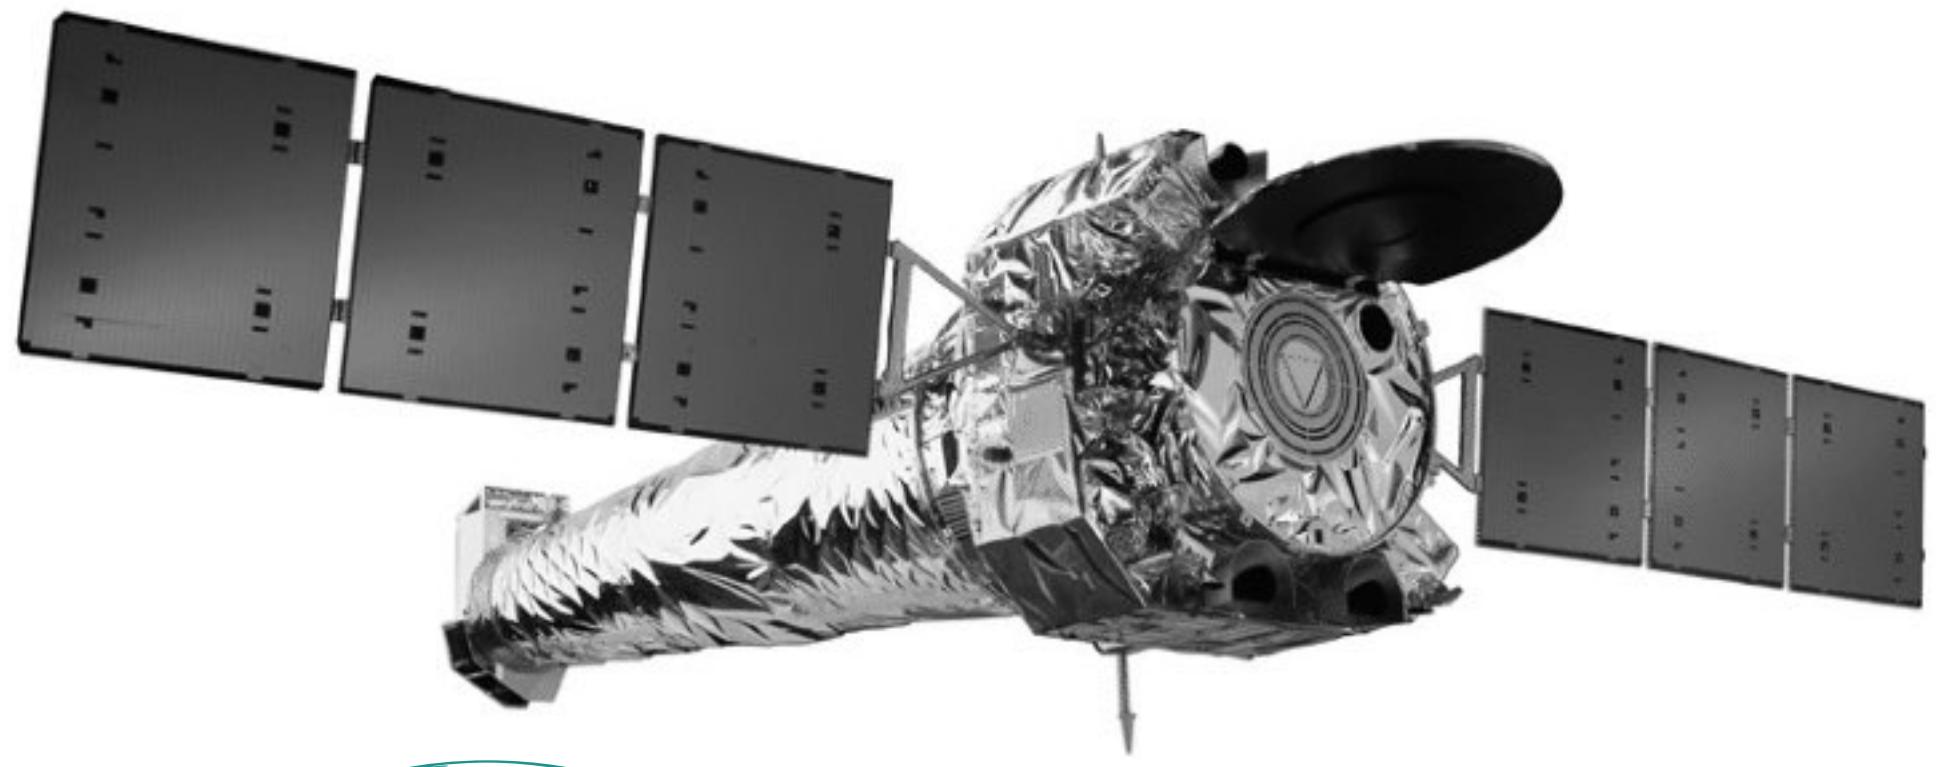


```
from pylira import LIRADeconvolverResult  
  
# read result back from disk  
result = LIRADeconvolverResult.read(  
    "pylira-result.fits"  
)  
  
# diagnostic plot  
result.plot_parameter_traces()
```

**Goal:** provide as much additional diagnostic information as possible!  
It is important to check and trust sampling results!

# Application Example

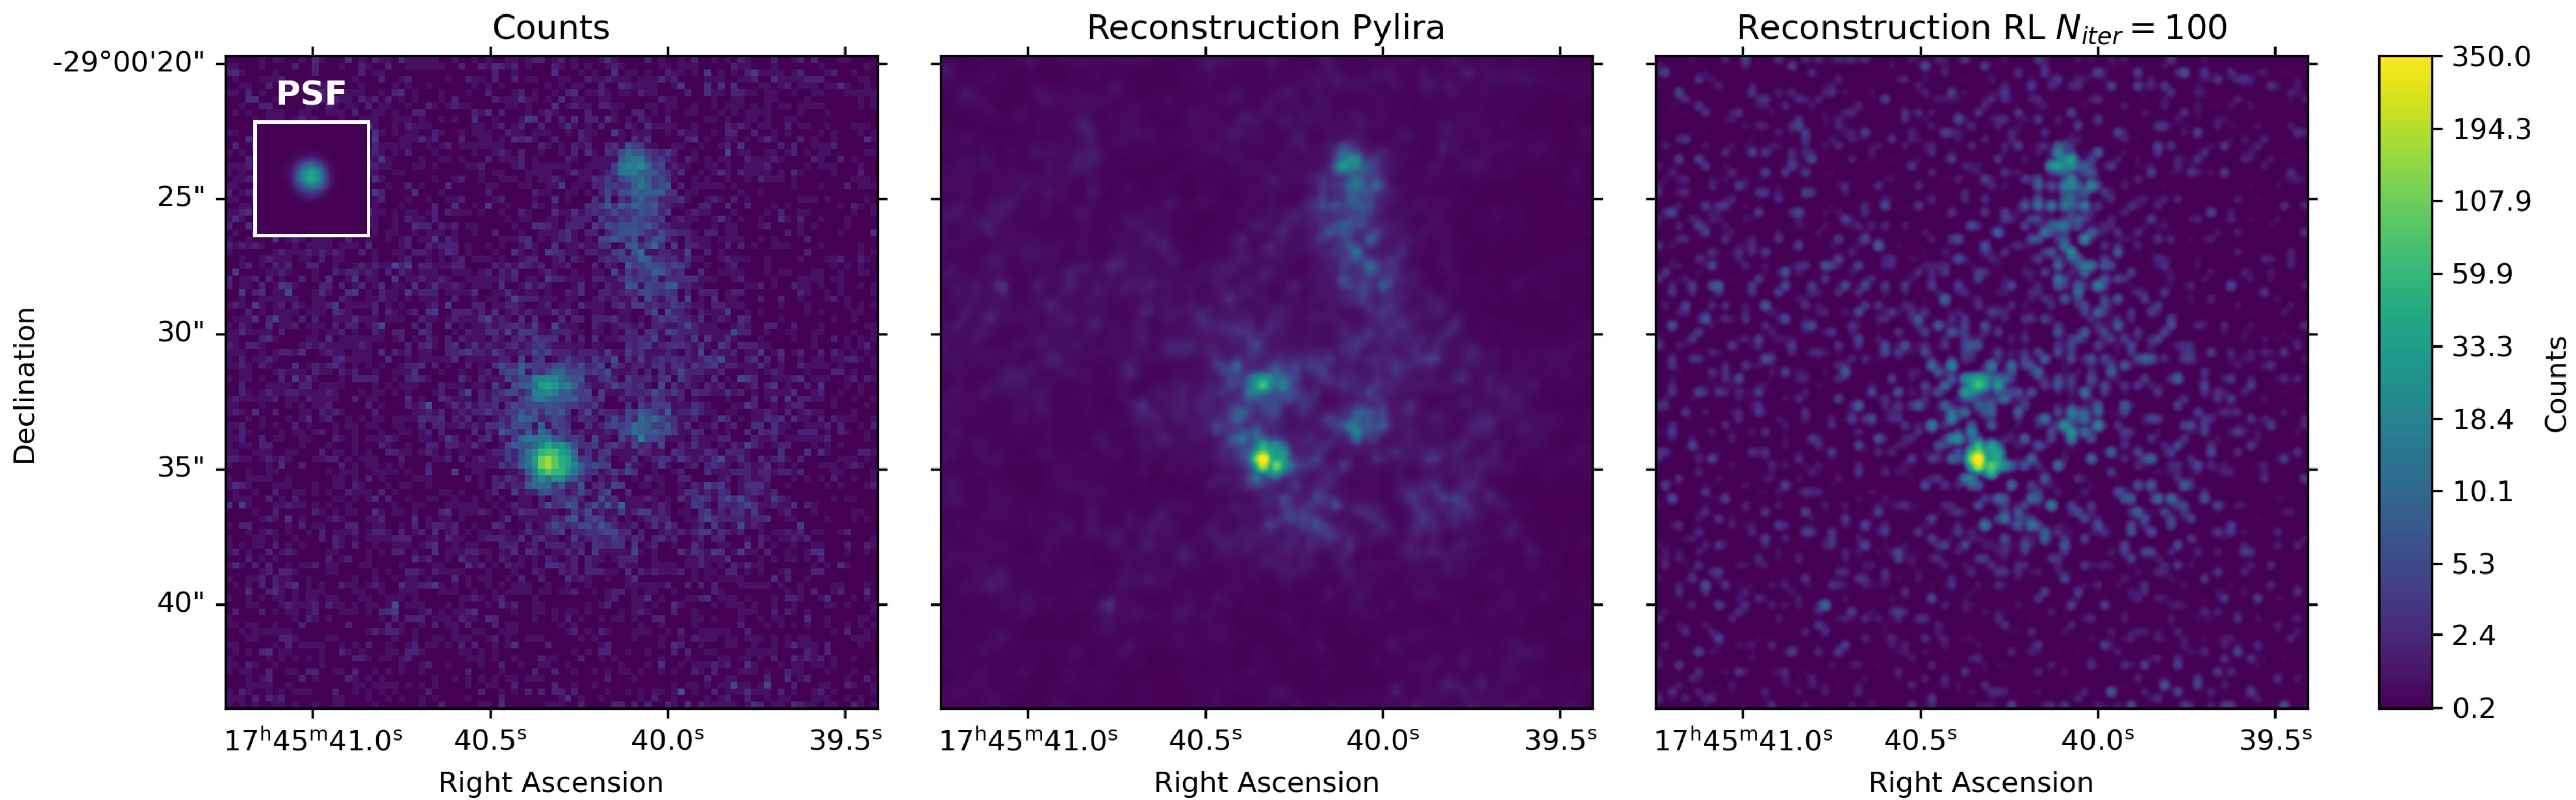
# CHANDRA Observatory



Crab Nebula seen with CHANDRA...

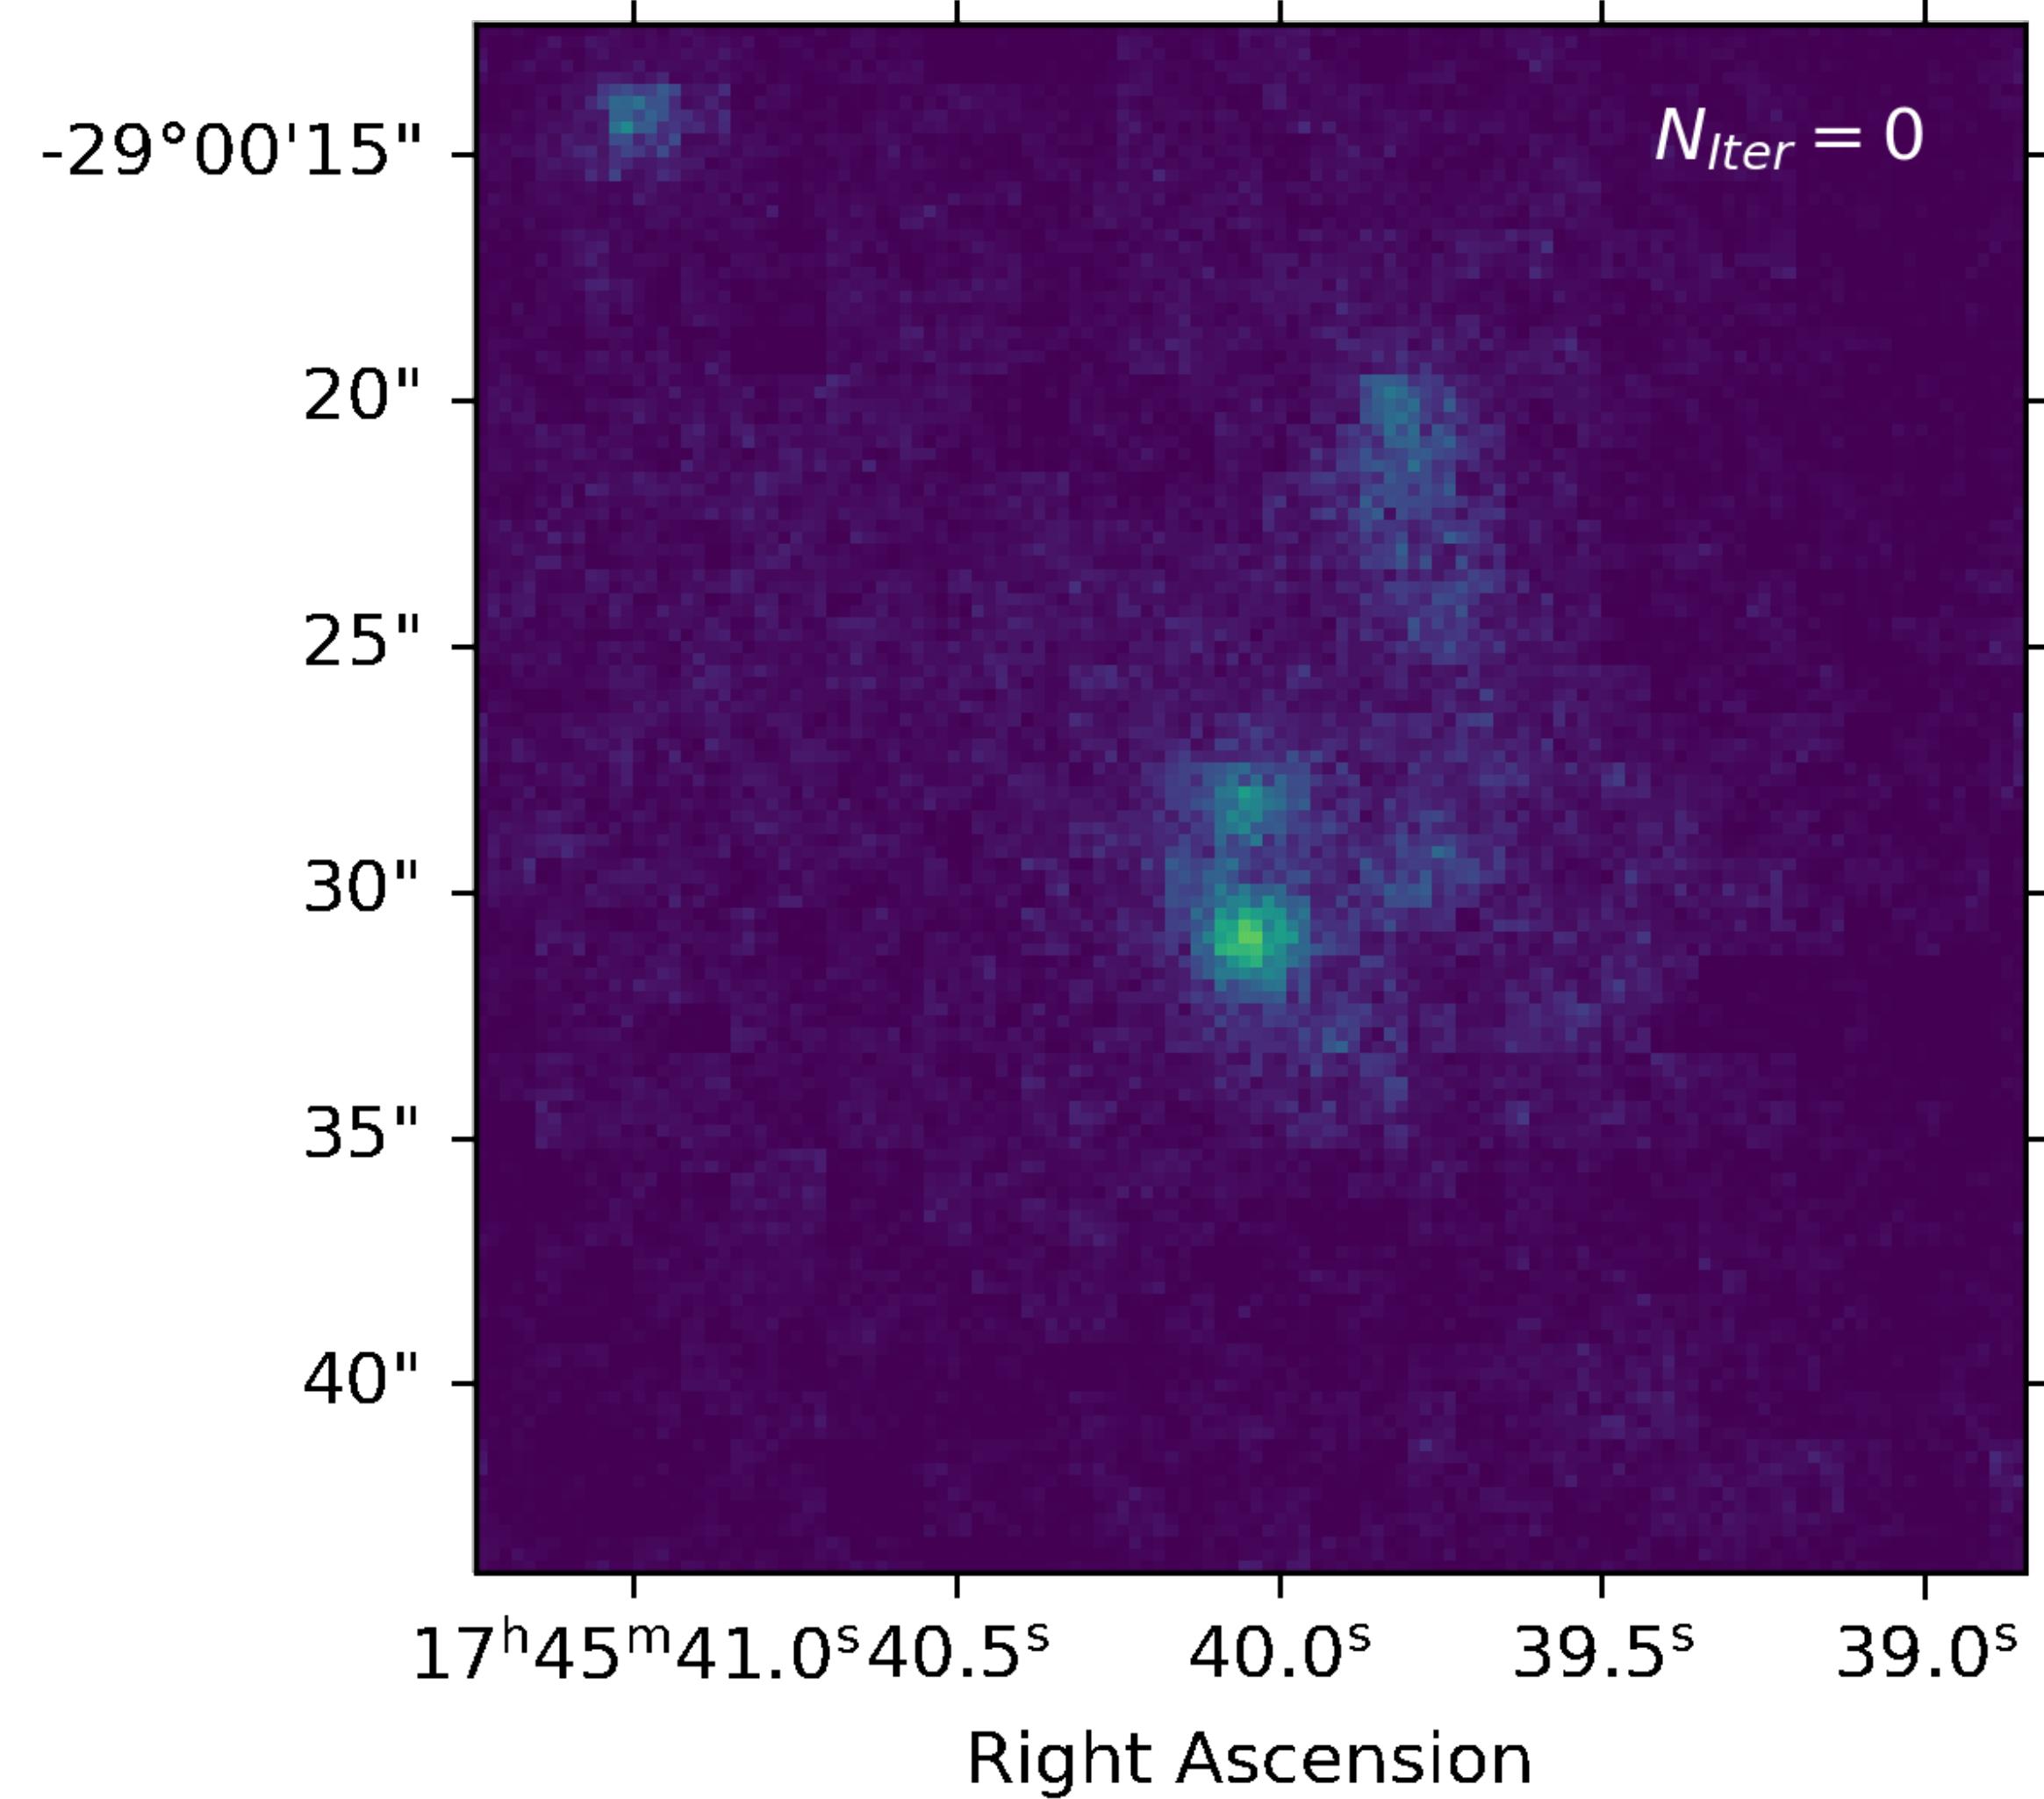
- Chandra is a **space-based X-ray observatory**, which has been in operation since 1999
- It consists of nested cylindrical paraboloid and hyperboloid surfaces, which form an **imaging optical system for X-rays**. In the focal plane, it has multiple instruments for different scientific purposes. This includes a high-resolution camera (HRC) and an Advanced CCD Imaging Spectrometer (ACIS)
- The typical **angular resolution is 0.5 arcsecond** and the covered **energy ranges from 0.1 - 10 keV**.

# CHANDRA Galactic Center

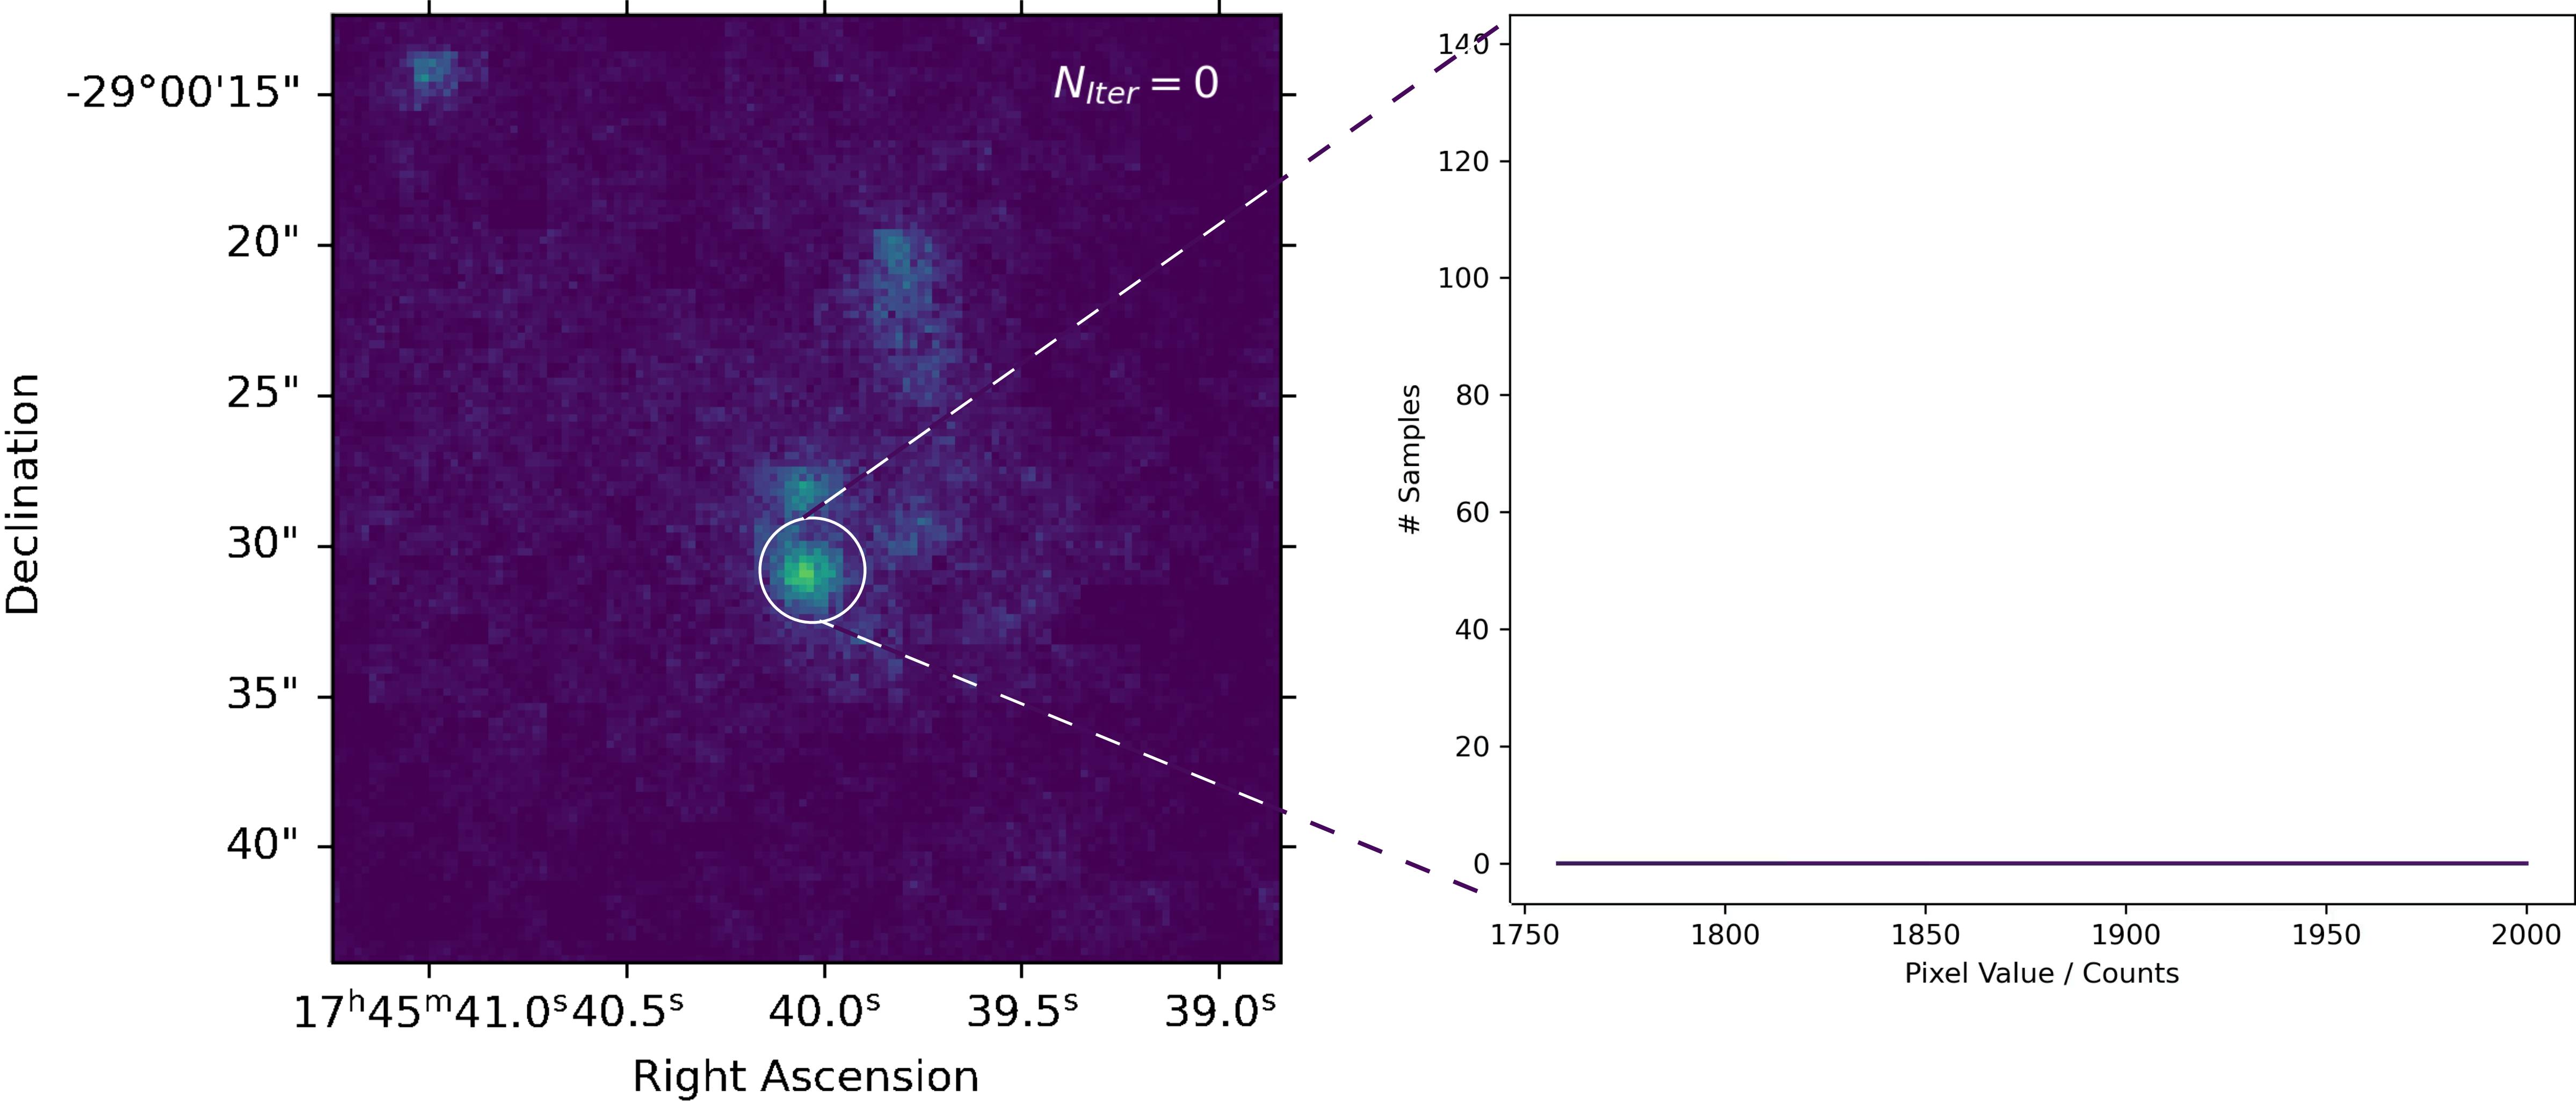


Pylira applied to Chandra ACIS data of the Galactic Center region, using the observation IDs 4684 and 4684.

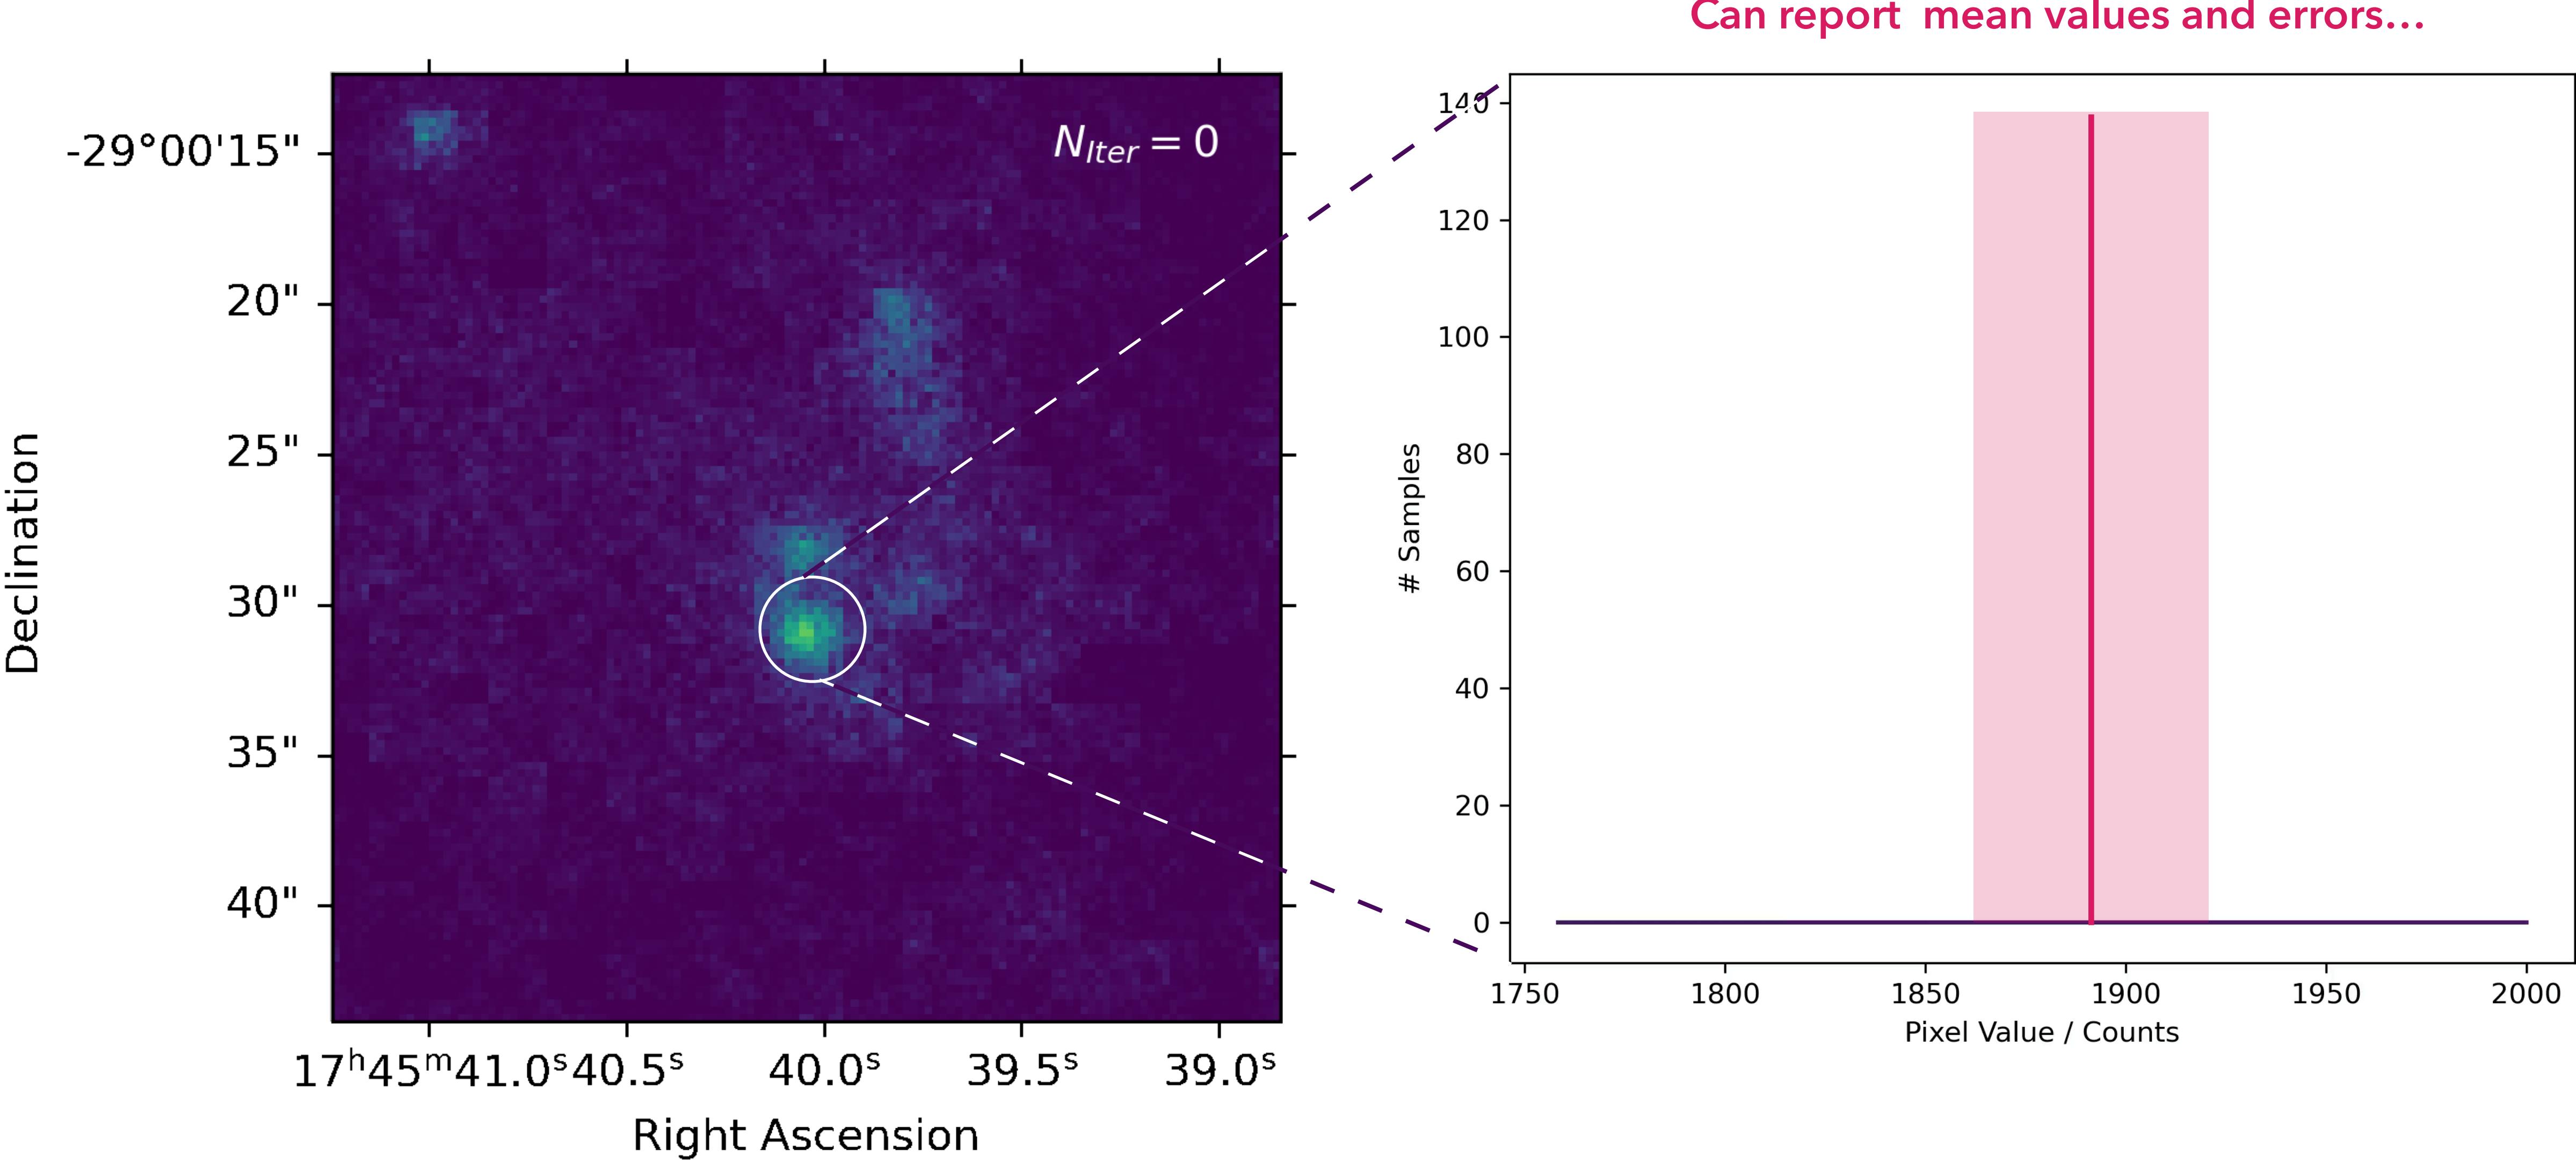
# CHANDRA Galactic Center



# CHANDRA Galactic Center



# CHANDRA Galactic Center



# Summary & Outlook

- Pylira is a **Python package for deconvolution of images** in the presence of Poisson noise. It acts as a wrapper around an original C (RMath) implementation.
- The LIRA method is a nice example for a “**pre ML era**” method, which is statistically fully “traceable”. So it delivers **error estimates and even per-pixel distributions**.
- The current implementation based on the initial C code is very inflexible in terms of prior definitions / modifications. Plans to **extend for more flexible prior definitions**.
- Could also **consider a re-implementation based on more modern technology**, such as [[pytorch](#)] / [[jax](#)] to take advantage from GPU acceleration
- Possibly combine with ML based methods e.g. by **learning patch based image priors**